

Stockholm Game Developer Forum 2010

Parallel Futures of a Game Engine (v2.0)

Johan Andersson
Rendering Architect, DICE



Background

- DICE
 - Stockholm, Sweden
 - ~250 employees
 - Part of *Electronic Arts*
 - *Battlefield & Mirror's Edge* game series

- Frostbite
 - Proprietary game engine used at DICE & EA
 - Developed by DICE over the last 5 years



BATTLEFIELD

BAD COMPANY 2



<http://badcompany2.ea.com/>

BATTLEFIELD

BAD COMPANY 2



<http://badcompany2.ea.com/>

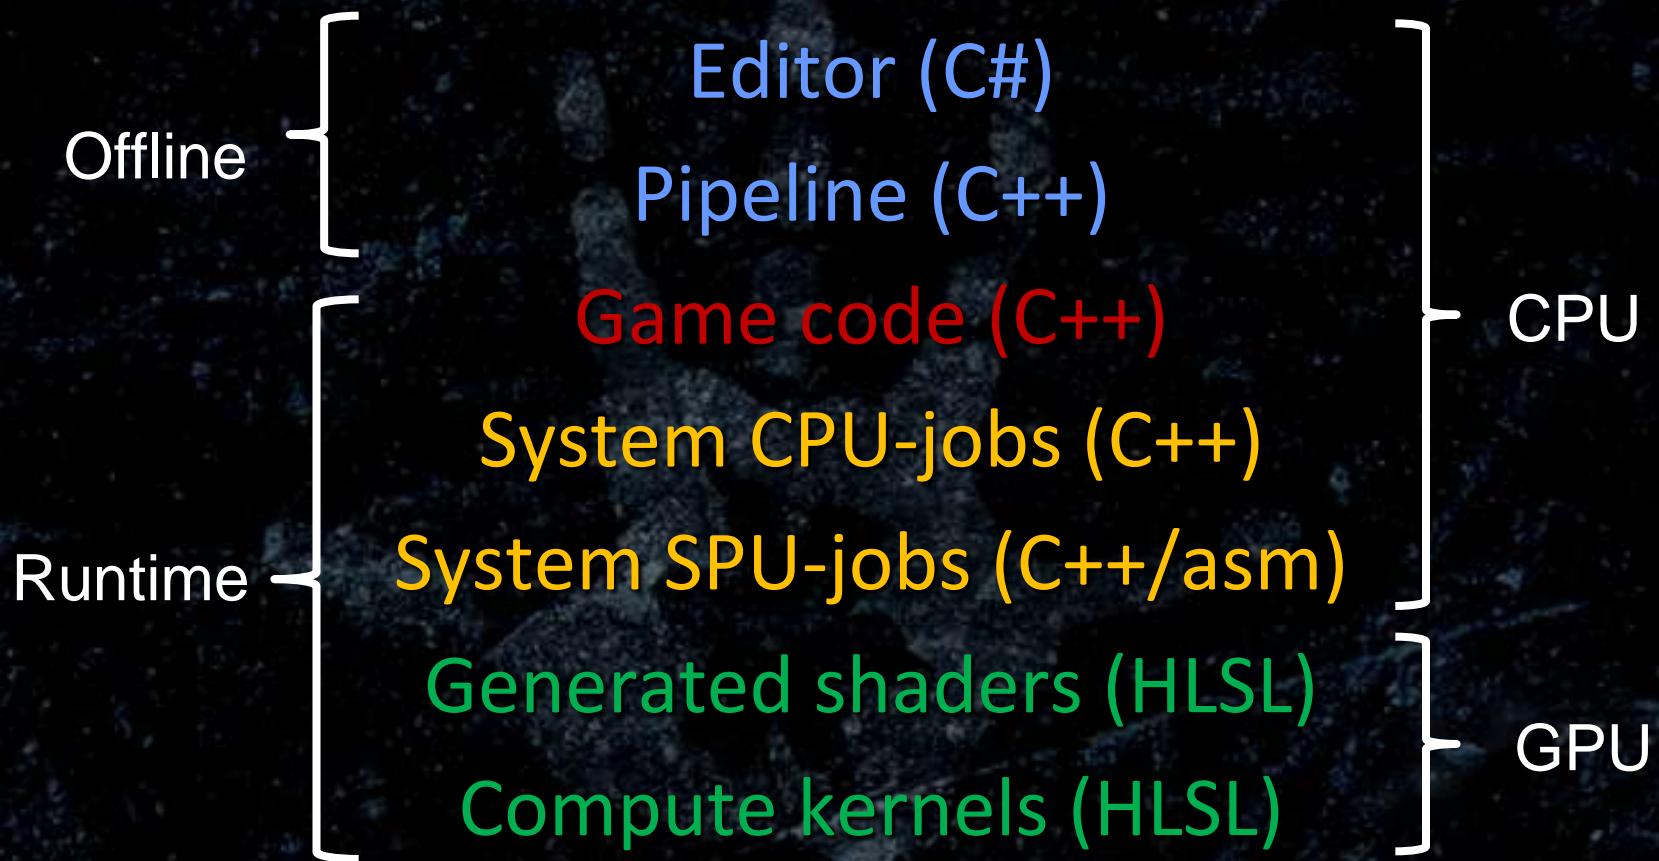
Outline

- Current parallelism
- Futures
- Q&A
- Slides will be available online later today
 - <http://publications.dice.se> and <http://repi.se>

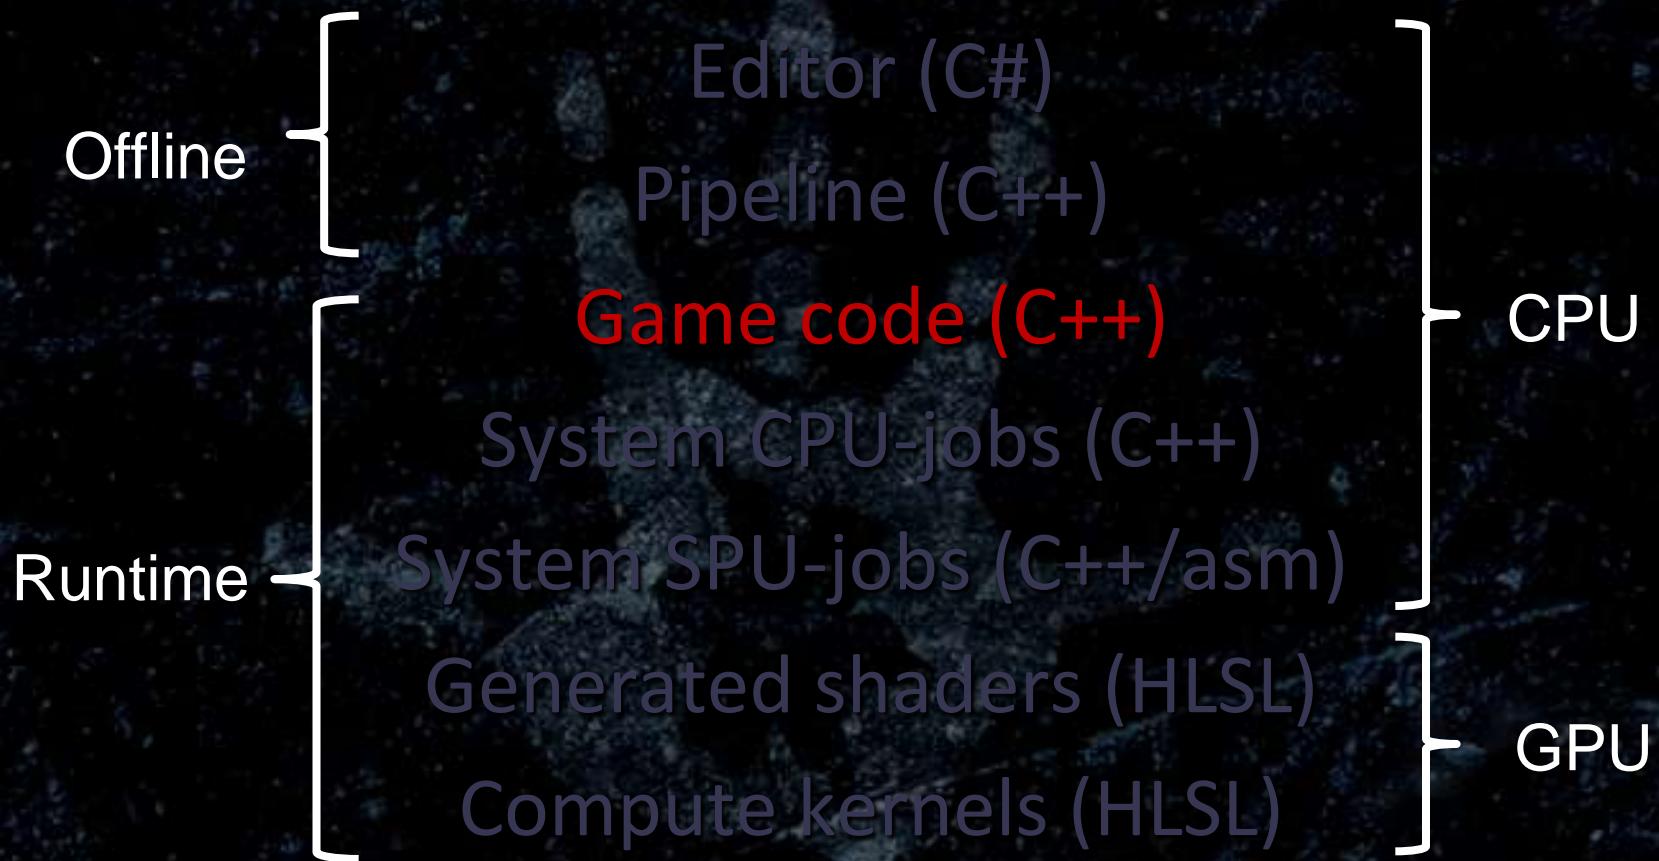
Current parallelism



Levels of code in Frostbite



Levels of code in Frostbite

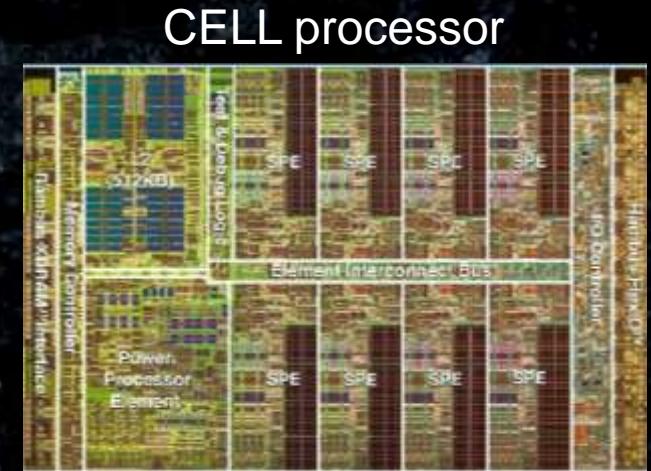


General “game code” (1/2)

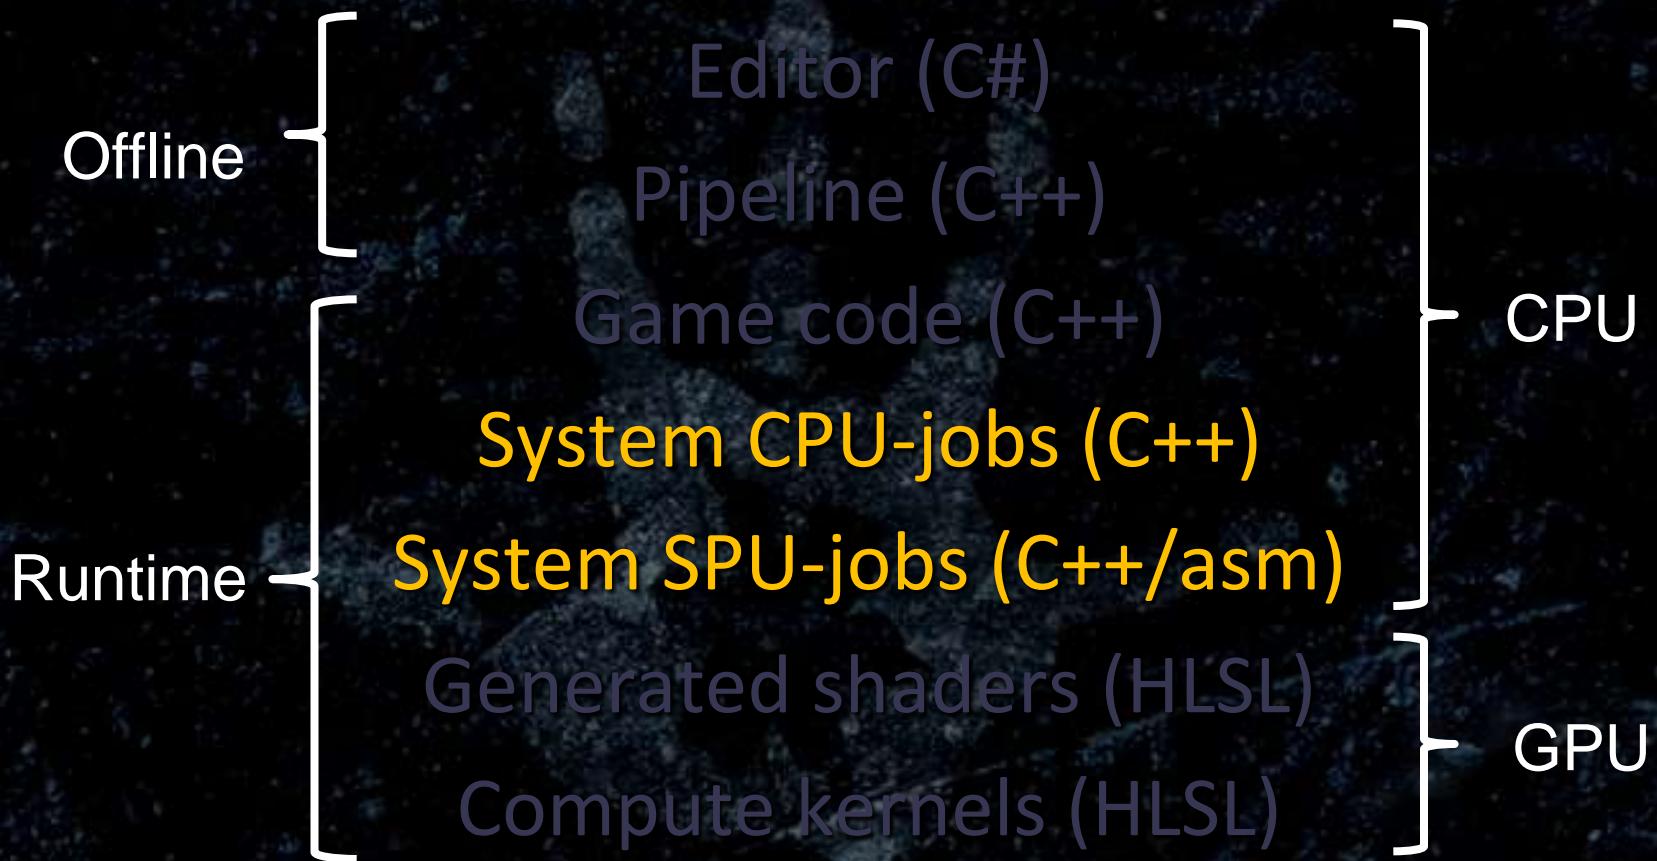
- This is the majority of our **1.5 million** lines of C++
 - Runs on Win32, Win64, Xbox 360 and PS3
 - We do not use any scripting language
- Similar to general application code
 - Huge amount of code & logic to maintain + continue to develop
 - Low compute density
 - “Glue code”
 - Scattered in memory (pointer chasing)
 - Difficult to efficiently parallelize
 - Out-of-order execution is a big help, but consoles are in-order 😞
- Key to be able to quickly iterate & change
 - This is the actual game logic & glue that **builds the game**
 - C++ not ideal, but has the invested infrastructure

General “game code” (2/2)

- PS3 is one of the main challenges
 - Standard CPU parallelization doesn't help (much)
 - CELL only has **2 HW threads on the PPU**
- Split the code in 2: **game code & system code**
 - Game logic, policy and glue code only on CPU
 - *“If it runs well on the PS3 PPU, it runs well everywhere”*
 - Lower-level systems on **PS3 SPUs**
- Main goals going forward:
 - Simplify & structure code base
 - Reduce coupling with lower-level systems
 - Increase in task parallelism for PC



Levels of code in Frostbite



Job-based parallelism

- *Essential* to utilize the cores on our target platforms
 - Xbox 360: 6 HW threads
 - PlayStation 3: 2 HW threads + 6 powerful SPUs
 - PC: 2-16 HW threads
- Divide up system work into *Jobs* (a.k.a. *Tasks*)
 - *15-200k* C++ code each. *25k* is common
 - Can depend on each other (if needed)
 - Dependencies create job graph
- All HW threads **consume** jobs
 - ~200-300 / frame

What is a Job for us?

An asynchronous function call

- Function ptr + 4 *uintptr_t* parameters
- Cross-platform scheduler: *EA JobManager*
- Uses *work stealing*

2 types of Jobs in Frostbite:

- *CPU job* (good)
 - General code moved into job instead of threads
- *SPU job* (great!)
 - Stateless pure functions, no side effects
 - Data-oriented, explicit memory DMA to local store
 - Designed to run on the PS3 SPUs = also very fast on in-order CPU
 - Can hot-swap → quick iterations ☺

Frostbite CPU job graph

Build big job graphs:

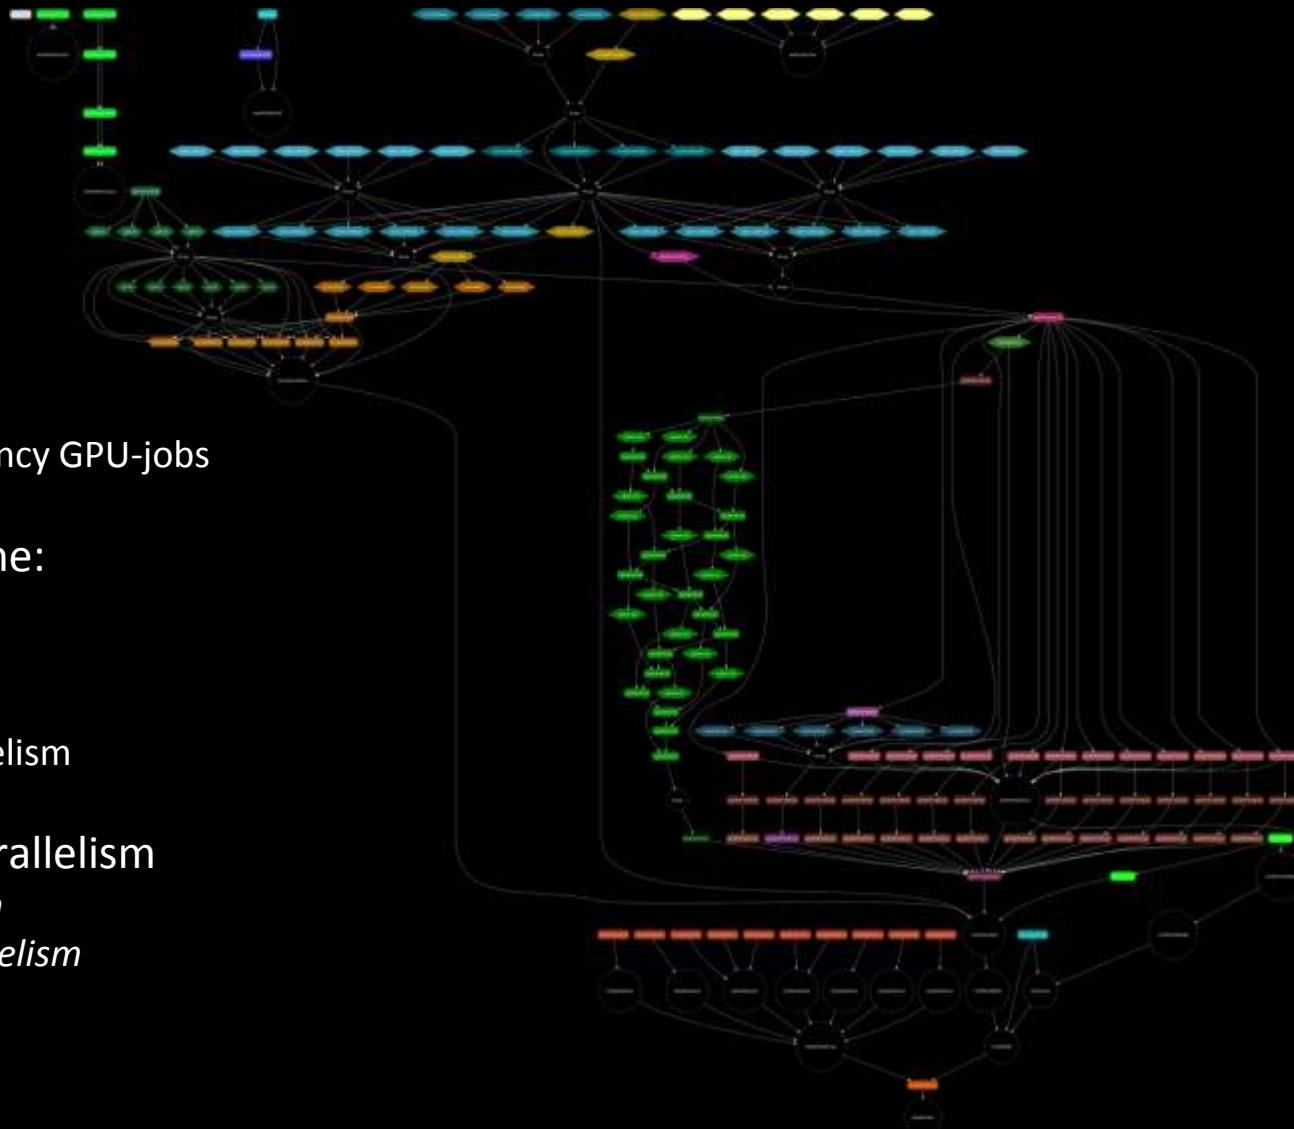
- Batch, batch, batch
- Mix CPU- & SPU-jobs
- Future: Mix in low-latency GPU-jobs

Job dependencies determine:

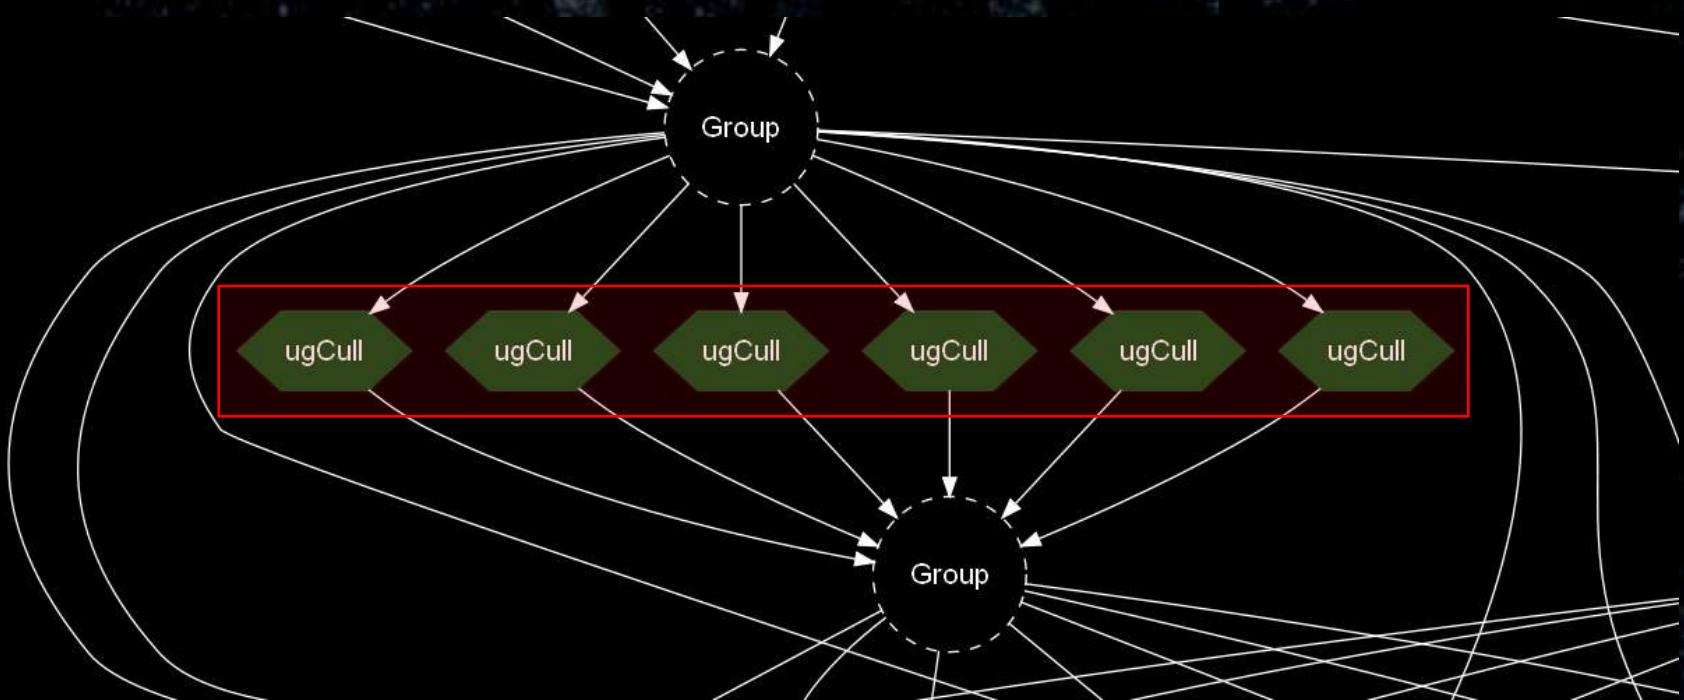
- Execution order
- Sync points
- Load balancing
- i.e. the *effective* parallelism

Intermixed task- & data-parallelism

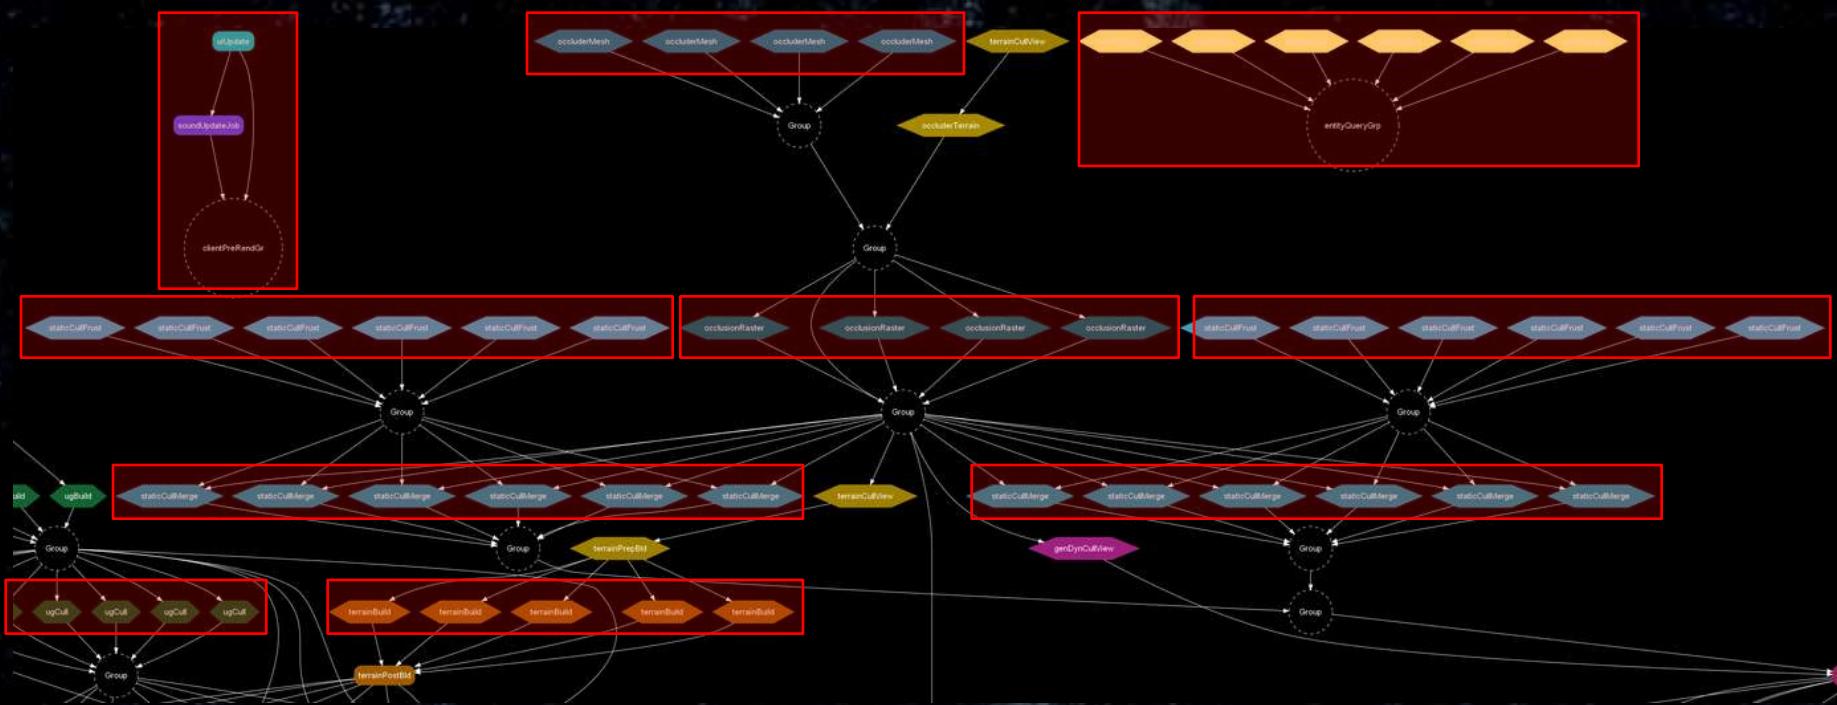
- aka *Braided Parallelism*
- aka *Nested Data-Parallelism*
- aka *Tasks and Kernels*



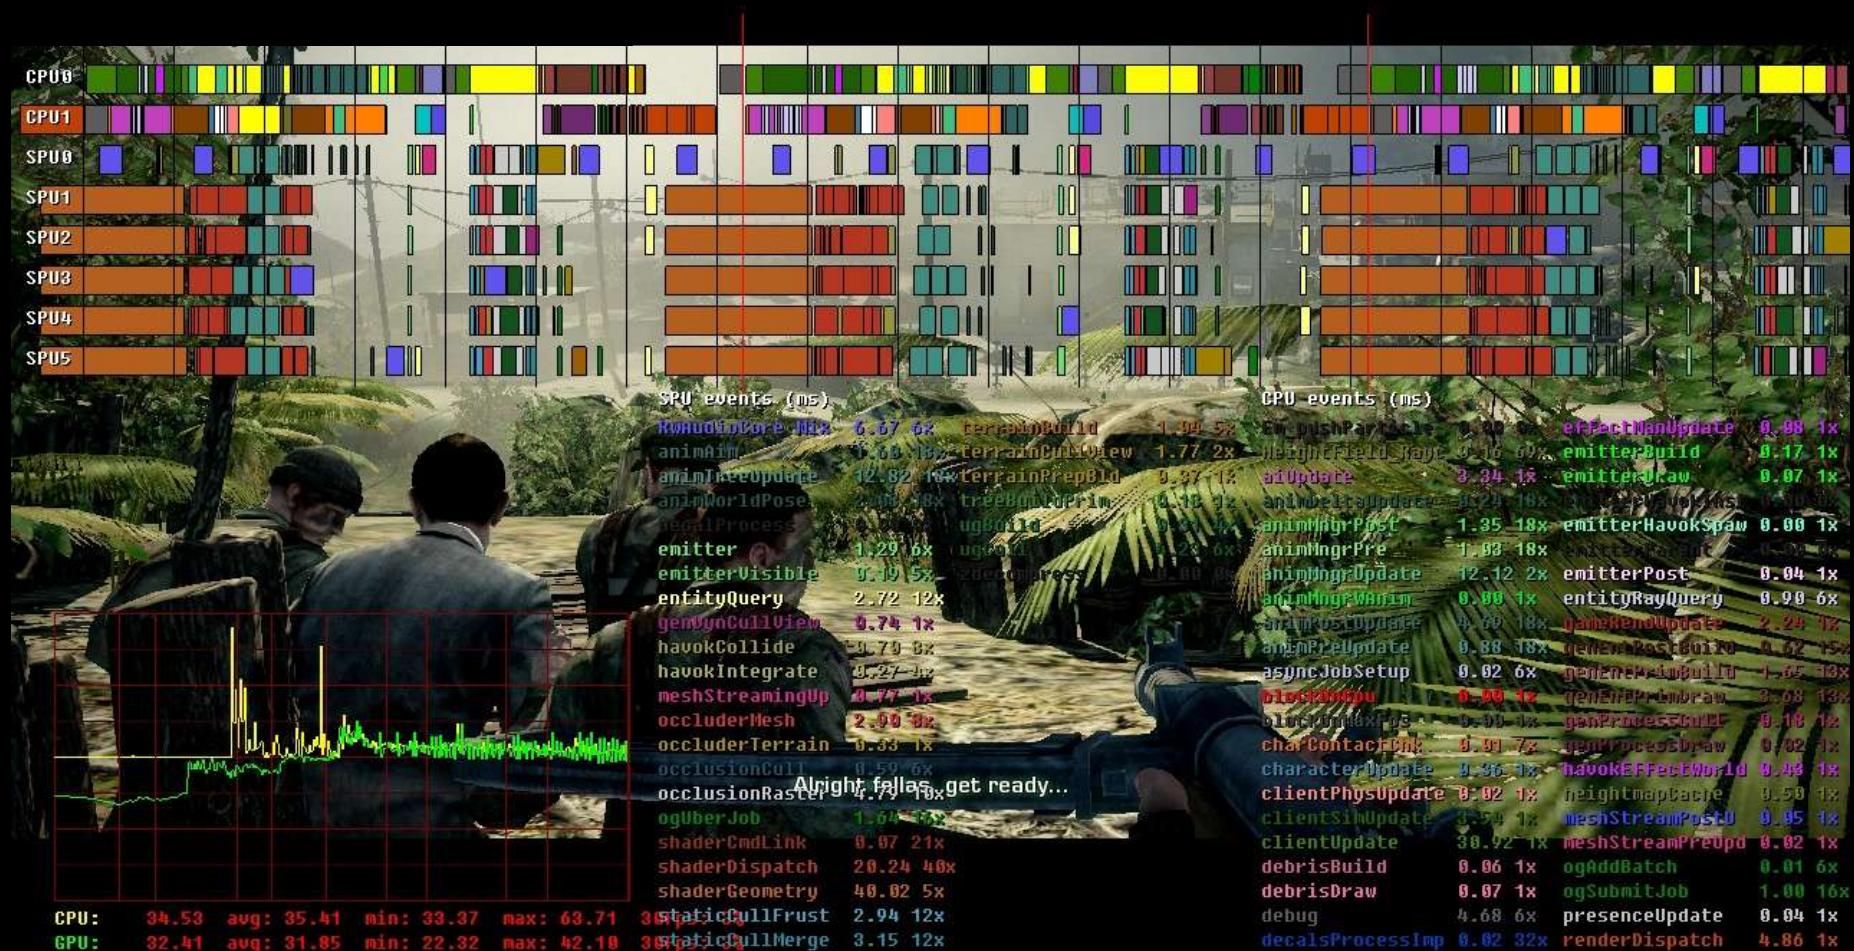
Data-parallel jobs



Task-parallel algorithms & coordination



Timing View: Battlefield Bad Company 2 (PS3)



Rendering jobs

Rendering systems are heavily divided up into **CPU- & SPU-jobs**

Most will move to GPU

- Eventually.. A few have already!
- PC CPU<->GPU *latency wall* 😞
- Mostly one-way data flow

Jobs:

- Terrain geometry [3]
- Undergrowth generation [2]
- Decal projection [4]
- Particle simulation
- Frustum culling
- Occlusion culling
- Occlusion rasterization
- Command buffer generation [6]
- PS3: Triangle culling [6]

Occlusion culling job example

Problem: Buildings & env occlude large amounts of objects

Obscured objects still have to:

- Update logic & animations
- Generate command buffer
- Processed on CPU & GPU
- = expensive & wasteful ☹

Difficult to implement full culling:

- Destructible buildings
- Dynamic occludees
- Difficult to precompute



From Battlefield: Bad Company PS3

Solution: Software occlusion culling

Rasterize coarse zbuffer on SPU/CPU

- 256x114 float
- Low-poly occluder meshes
 - 100 m view distance
 - Max 10000 vertices/frame
- Parallel vertex & raster SPU-jobs
- Cost: a few milliseconds



Cull all objects against zbuffer

- Screen-space bounding-box test
- Before passed to all other systems
- Big performance savings!



GPU occlusion culling

Ideally want to use the GPU, but current APIs are limited:

- Occlusion queries introduces overhead & latency
- Conditional rendering only helps GPU
- Compute Shader impl. possible, but same *latency wall*

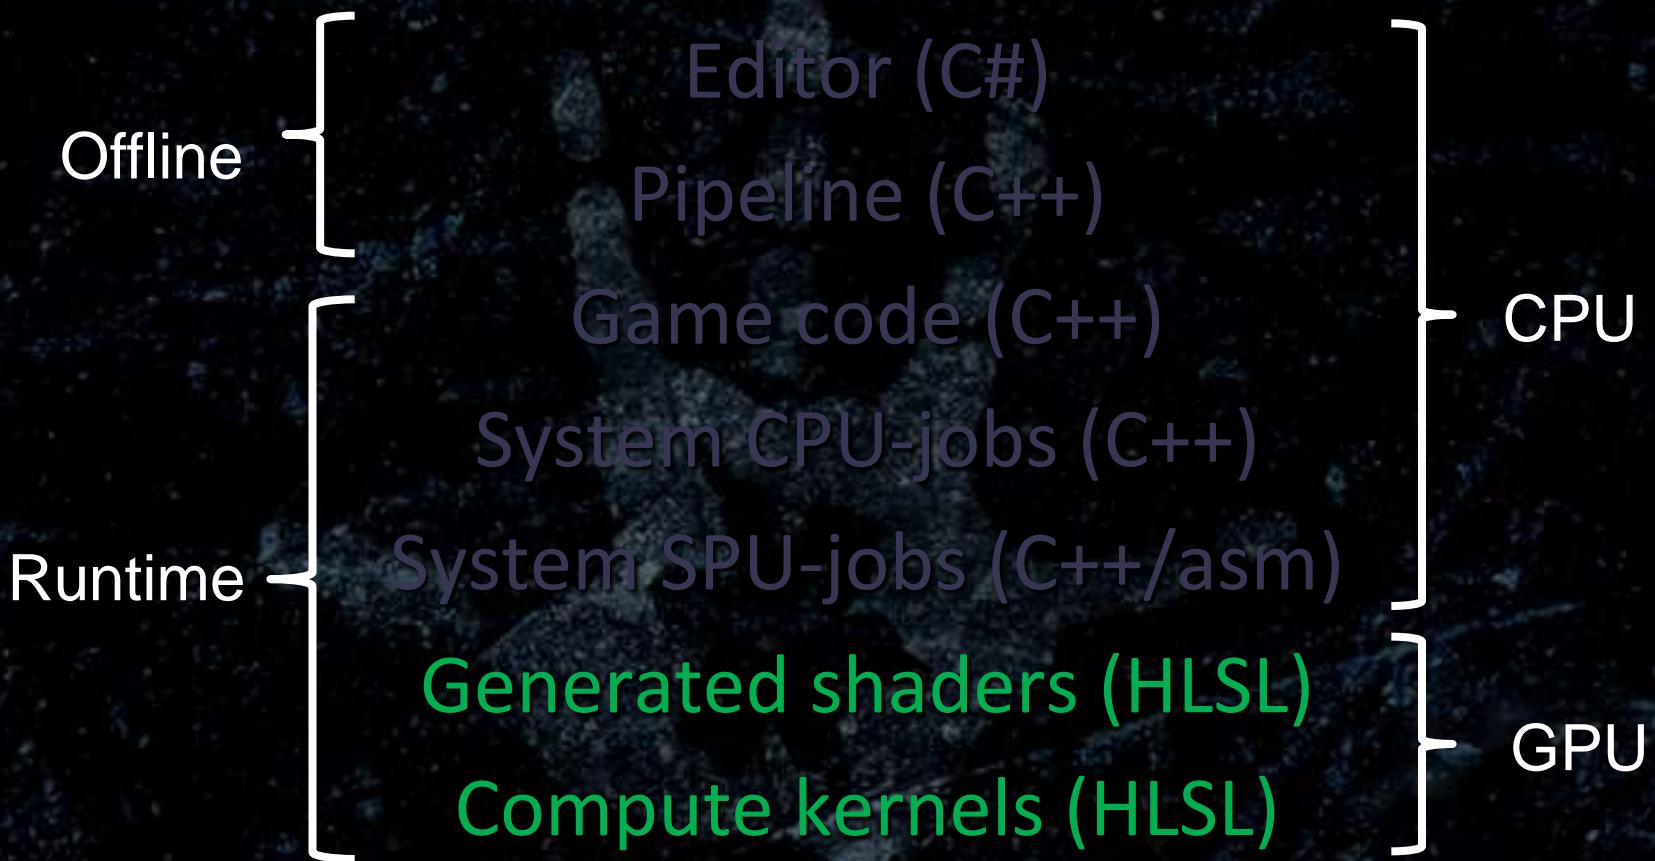
Future 1: Low-latency GPU execution context

- Rasterization and testing done on GPU where it belongs
- Lockstep with CPU, need to read back within a few ms
- Should be possible on *Larrabee* & *Fusion*, want on all PC

Future 2: Move entire cull & rendering to "GPU"

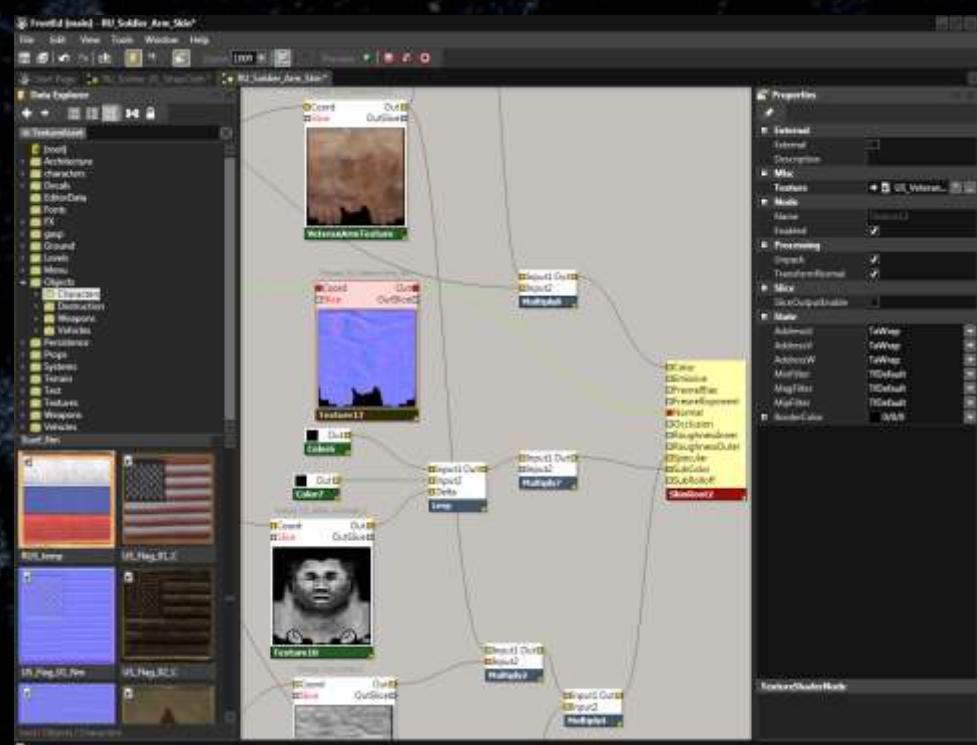
- World, cull, systems, dispatch. End goal
- *Very thin* GPU graphics driver – GPU feeds itself

Levels of code in Frostbite



Shader types

- Generated shaders [1]
 - Graph-based *surface shaders*
 - Treated as content, not code
 - Artist created
 - Generates HLSL code
 - Used by all meshes and 3d surfaces
- Graphics / Compute kernels
 - Hand-coded & optimized HLSL
 - Statically linked in with C++
 - Pixel- & compute-shaders
 - Lighting, post-processing & special effects



Graph-based surface shader in FrostEd 2

Futures



Challenges

3 major challenges/goals going forward:

1. How do we make it easier to develop, maintain & parallelize **general game code**?
2. What do we need to continue to innovate & scale up real-time **computational graphics**?
Most likely the same solution(s)!
3. How can we move & scale up advanced simulation and **non-graphics tasks** to data-parallel manycore processors?

Challenge 1

“How do we make it easier to develop, maintain & parallelize general game code?”

- *Shared State Concurrency* is a killer
 - Not a big believer in *Software Transactional Memory* either
 - Because of performance and too “optimistic” flow
- A more strict & adapted C++ model
 - Support for true **immutable** & **r/w-only** memory access
 - Per-thread/task memory access **opt-in**
 - Reduce the possibility for side effects in parallel code
 - As much compile-time validation as possible
 - Micro-threads / coroutines as first class citizens
 - More?
- Other languages?

Challenge 1 - Task parallelism

- Multiple task libraries
 - EA JobManager
 - Dynamic job graphs, on-demand dependencies & continuations
 - Targeted to cross-platform SPU-jobs, key requirement for this generation
 - Not geared towards super-simple to use CPU parallelism
 - MS ConcRT, Apple GCD, Intel TBB
 - All has some good parts!
 - Neither works on all of our current platforms
 - OpenMP
 - Just say no. Parallelism can't be tacked on, need to be an explicit core part
- Need C++ enhancements to simplify usage
 - C++ 0x lambdas / GCD blocks 😊
 - Glacial C++ development & deployment 😞
 - Want on all platforms, so lost on this console generation

Challenge 2 - Definition

”Real-time interactive graphics & simulation at Pixar level of quality”

- Needed visual features:
 - Complete anti-aliasing + natural motion blur & DOF
 - Global indirect lighting & reflections
 - Sub-pixel geometry
 - OIT
 - Huge improvements in character animation

These require massively more compute, BW and improved model!

(animation can't be solved with just more/better compute, so pretend it doesn't exist for now)



Challenge 2 - Problems

Problems & limitations with current GPU model:

- Fixed rasterization pipeline
 - Compute pipeline not fast enough to replace it
- GPU is handicapped by being **spoon-fed** by CPU
- *Irregular workloads* are difficult & inefficient
- Can't express *nested data-parallelism* well
- Current **HLSL** is a very limited language

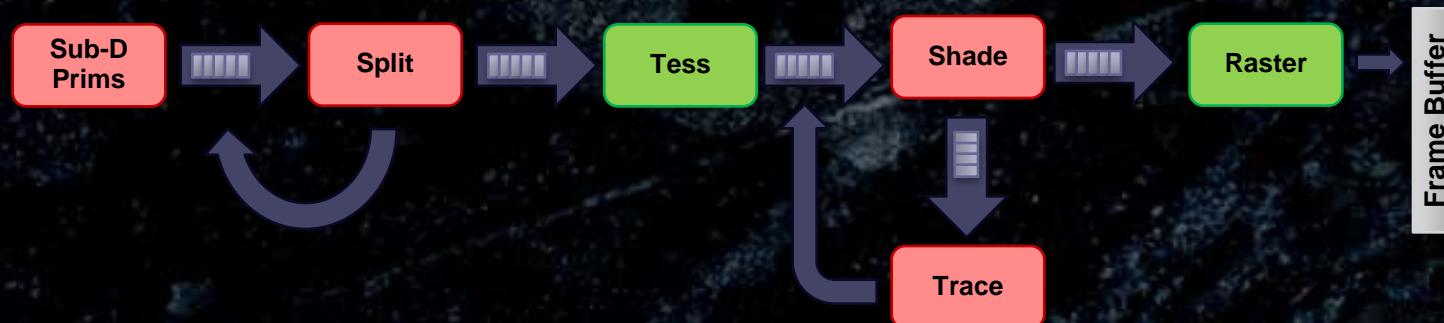
Challenge 2 - Solutions

- Absolutely a job for a *high-throughput oriented data-parallel processor*
 - With a highly flexible programming model
 - The CPU, as we know it, and APIs/drivers are **only in the way**
- Pure software solution not practical as next step after DX11 PC
 - Multi-vendor & multi-architecture marketplace
 - But for **future consoles**, the more flexible the better (perf. permitting)
- Long term goal:
 - Multi-vendor & multi-platform **standard/virtual DP ISA**
- Want a rich programmable compute model as next step
 - Nested data-parallelism & GPU self-feeding
 - Low-latency CPU<->GPU interaction
 - Efficiently target varying HW architectures

"Pipelined Compute Shaders"

- **Queues** as streaming I/O between compute kernels
 - Simple & expressive model supporting *irregular workloads*
 - Keeps data on chip, supports variable sized caches & cores
 - Can target multiple types of HW & architectures
- Hybrid **graphics/compute** user-defined pipelines
 - Language/API defining fixed stages inputs & outputs
 - Pipelines can feed other pipelines (similar to *DrawIndirect*)

Reyes-style Rendering with Ray Tracing



"Pipelined Compute Shaders"

- Wanted for next major *DirectX* and *OpenCL/OpenGL*
 - As a standard, as soon as possible
 - Run on all: discrete GPU, integrated GPU and CPU
- Model is also a good fit for many of our CPU/SPU jobs
 - Parts of job graph can be seen as queues between stages
 - Easier to write kernels/jobs with **streaming I/O**
 - Instead of explicit fixed-buffers and "memory passes"
 - Or dynamic memory allocation

Language?

- Future DP language is a big question
 - But the **concepts & infrastructure** are what is important!
- Could be an **extended HLSL** or "**data-parallel C++**"
 - Data-oriented imperative language (i.e. not standard C++)
 - Think HLSL would probably be easier & the most explicit
 - Amount of code is small and written from scratch
- Shader-like implicit SoA vectorization (**SIMT**)
 - Instead of SSE/LRBni-like explicit vectorization
 - Need to be first class language citizen, not tacked on C++
 - Easier to target multiple evolving architectures implicitly

Future hardware (1/2)

- Single main memory & address space
 - Critical to share resources between graphics, simulation and game in **immersive dynamic worlds**
- Configurable kernel local stores / cache
 - Similar to *Nvidia Fermi* & *Intel Larrabee*
 - Local stores = reliability & good for **regular loads**
 - Together with user-driven **async DMA** engines
 - Caches = essential for **irregular data structures**
- Cache coherency?
 - Not always important for kernels
 - But essential for general code, can partition?

Future hardware (2/2)

- 2015 = **40 TFLOPS**, we would spend it on:
 - 80% graphics
 - 15% simulation
 - 4% misc
 - 1% game (wouldn't use all **400 GFLOPS** for game logic & glue!)
- OOE CPUs more efficient for the majority of our game code
 - But for the *vast* majority of our FLOPS these are fully irrelevant
 - Can evolve to a **small dot** on a sea of DP cores
 - Or run **scalar** on general DP cores wasting some compute
- In other words: **no need for separate CPU and GPU!**
-> single heterogeneous processor

Conclusions

- Future is an interleaved mix of task- & data-parallelism
 - On both the HW and SW level
 - But programmable DP is where the massive compute is done
- Data-parallelism requires *data-oriented design*
- Developer productivity can't be limited by model(s)
 - It should enhance productivity & perf on all levels
 - Tools & language constructs play a critical role
- We should welcome our parallel future!

Thanks to

- DICE, EA and the Frostbite team
- The graphics/gamedev community on Twitter
- Steve McCalla, Mike Burrows
- Chas Boyd
- Nicolas Thibieroz, Mark Leather
- Dan Wexler, Yury Uralsky
- Kayvon Fatahalian

References

Previous Frostbite-related talks:

- [1] Johan Andersson. "Frostbite Rendering Architecture and Real-time Procedural Shading & Texturing Techniques ". GDC 2007.
<http://repi.blogspot.com/2009/01/conference-slides.html>
- [2] Natasha Tatarchuk & Johan Andersson. "Rendering Architecture and Real-time Procedural Shading & Texturing Techniques". GDC 2007.
[http://developer.amd.com/Assets/Andersson-Tatarchuk-FrostbiteRenderingArchitecture\(GDC07_AMD_Session\).pdf](http://developer.amd.com/Assets/Andersson-Tatarchuk-FrostbiteRenderingArchitecture(GDC07_AMD_Session).pdf)
- [3] Johan Andersson. "Terrain Rendering in Frostbite using Procedural Shader Splatting". Siggraph 2007. [http://developer.amd.com/media/gpu_assets/Andersson-TerrainRendering\(Siggraph07\).pdf](http://developer.amd.com/media/gpu_assets/Andersson-TerrainRendering(Siggraph07).pdf)
- [4] Daniel Johansson & Johan Andersson. "Shadows & Decals – D3D10 techniques from Frostbite". GDC 2009. <http://repi.blogspot.com/2009/03/gdc09-shadows-decals-d3d10-techniques.html>
- [5] Bill Bilodeau & Johan Andersson. "Your Game Needs Direct3D 11, So Get Started Now!". GDC 2009. <http://repi.blogspot.com/2009/04/gdc09-your-game-needs-direct3d-11-so.html>
- [6] Johan Andersson. "Parallel Graphics in Frostbite". Siggraph 2009, Beyond Programmable Shading course. <http://repi.blogspot.com/2009/08/siggraph09-parallel-graphics-in.html>

Questions?



Email: johan.andersson@dice.se **Blog:** <http://repi.se> **Twitter:** @repi

For more DICE talks: <http://publications.dice.se>

Bonus slides

Game development

- 2 year development cycle
 - New IP often takes *much* longer, 3-5 years
 - Engine is continuously in development & used
- AAA teams of 70-90 people
 - 40% artists
 - 30% designers
 - 20% programmers
 - 10% audio
- Budgets \$20-40 million
- Cross-platform development is market reality
 - Xbox 360 and PlayStation 3
 - PC (DX9, DX10 & DX11 for BC2)
 - Current consoles will stay with us for many more years

Game engine requirements (1/2)

- Stable real-time performance
 - Frame-driven updates, 30 fps
 - Few threads, instead **per-frame jobs/tasks** for everything
- Predictable memory usage
 - Fixed budgets for systems & content, fail if over
 - Avoid runtime allocations
 - Love *unified memory!*
- Cross-platform
 - The consoles determines our **base tech level** & focus
 - PS3 is design target, most difficult and good potential
 - Scale up for PC, dual core is min spec (**slow!**)

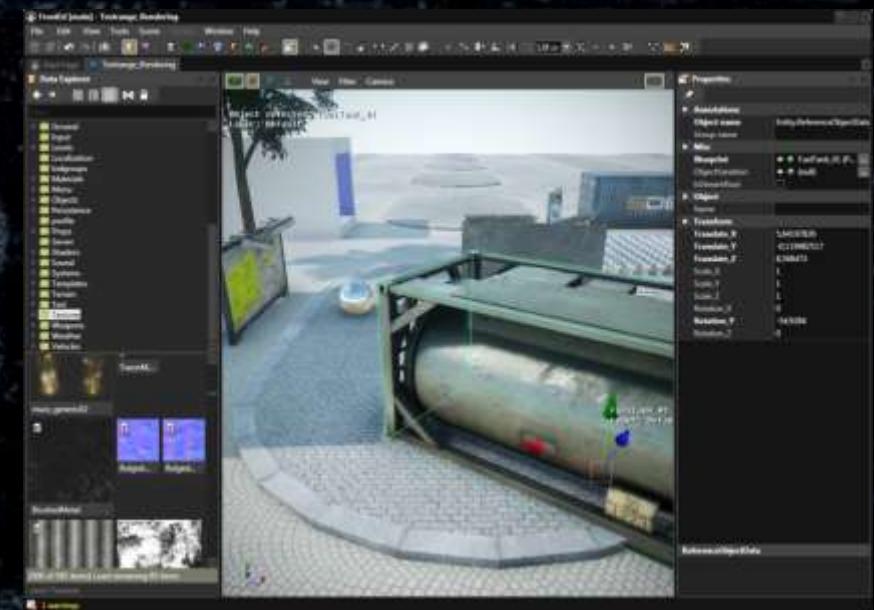
Game engine requirements (2/2)

- Full system profiling/debugging
 - Engine is a vertical solution, touches everywhere
 - PIX, xbtracedump, SN Tuner, ETW, GPUView
- Quick iterations
 - Essential in order to be creative
 - Fast building & fast loading, **hot-swapping** resources
 - Affects both the tools and the game
- Middleware
 - Use when it make senses, cross-platform & optimized
 - Parallelism have to go through *our* systems

Editor & Pipeline

Editor ("FrostEd 2")

- WYSIWYG editor for content
- C#, Windows only
- Basic threading / tasks



Pipeline

- Offline/background data-processing & conversion
- C++, some MC++, Windows only
- Typically IO-bound
- A few compute-heavy steps use CUDA
- Texture compression uses CUDA, prefer OpenCL or CS

CPU parallelism models are generally **not** a problem here

EntityRenderCull job example

```
struct FB_ALIGN(16) EntityRenderCullJobData
{
    enum
    {
        MaxSphereTreeCount = 2,
        MaxStaticCullTreeCount = 2
    };

    uint sphereTreeCount;
    const SphereNode* sphereTrees[MaxSphereTreeCount];

    u8 viewCount;
    u8 frustumCount;

    u8 viewIntersectFlags[32];
    Frustum frustums[32];

    .... (cut out 2/3 of struct for display size)

    u32 maxOutEntityCount;

    // Output data, pre-allocated by callee
    u32 outEntityCount;
    EntityRenderCullInfo* outEntities;
};

void entityRenderCullJob(EntityRenderCullJobData* data);

void validate(const EntityRenderCullJobData& data);
```

- Frustum culling of dynamic entities in sphere tree
- struct contain all input data needed
- Max output data pre-allocated by callee
- Single job function
 - Compile both as CPU & SPU job
- Optional struct validation func

EntityRenderCull SPU setup

```
// local store variables
EntityRenderCullJobData g_jobData;
float g_zBuffer[256*114];
u16 g_terrainHeightData[64*64];

int main(uintptr_t dataEa, uintptr_t, uintptr_t, uintptr_t)
{
    dmaBlockGet("jobData", &g_jobData, dataEa, sizeof(g_jobData));
    validate(g_jobData);

    if (g_jobData.zBufferTestEnable)
    {
        dmaAsyncGet("zBuffer", g_zBuffer, g_jobData.zBuffer, g_jobData.zBufferResX*g_jobData.zBufferResY*4);
        g_jobData.zBuffer = g_zBuffer;

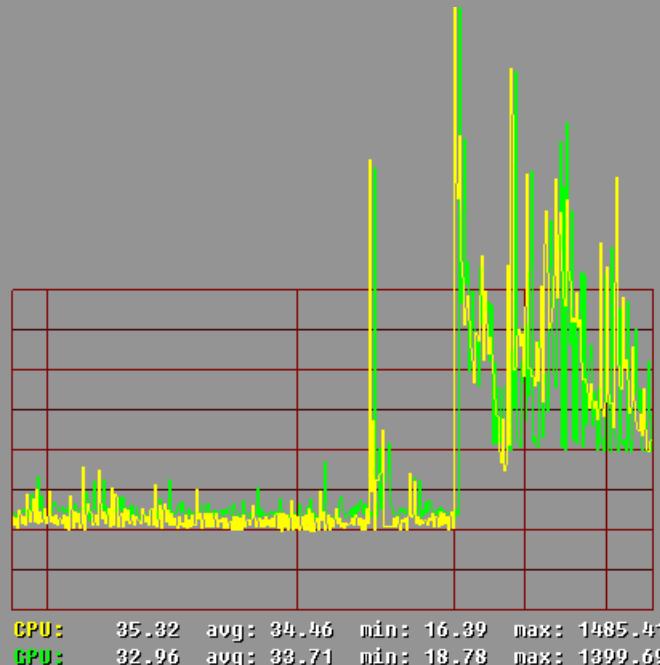
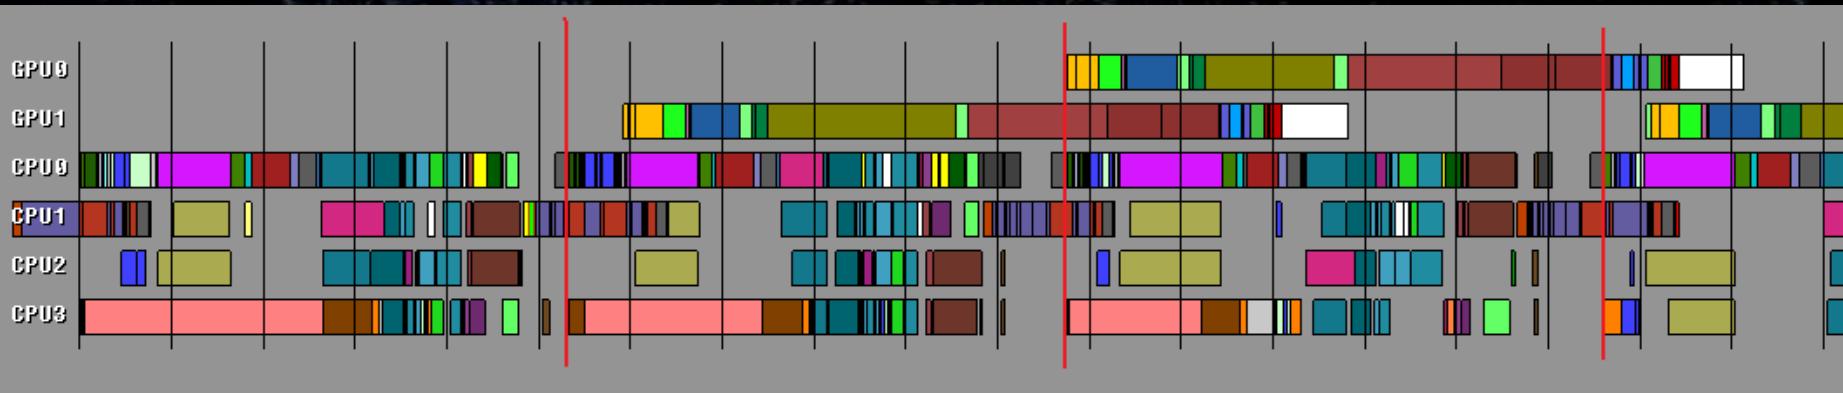
        if (g_jobData.zBufferShadowTestEnable && g_jobData.terrainHeightData)
        {
            dmaAsyncGet("terrainHeight", g_terrainHeightData, g_jobData.terrainHeightData, g_jobData.terrainHeightDataSize);
            g_jobData.terrainHeightData = g_terrainHeightData;
        }
        dmaWaitAll(); // block on both DMAs
    }

    // run the actual job, will internally do streaming DMAs to the output entity list
    entityRenderCullJob(&g_jobData);

    // put back the data because we changed outEntityCount
    dmaBlockPut(dataEa, &g_jobData, sizeof(g_jobData));
    return 0;
}
```

Timing view

Example: PC, 4 CPU cores, 2 GPUs in AFR (AMD Radeon 4870x2)



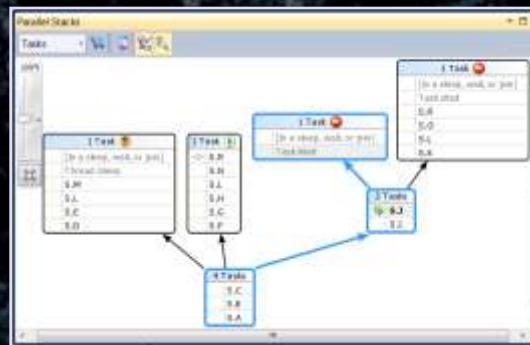
- Real-time in-game overlay
 - See timing events & effective parallelism
 - On **CPU, SPU & GPU** – for all platforms
 - Use to reduce sync-points & optimize load balancing
- GPU timing through DX event queries
- Our main performance tool!

Language (cont.)

- Requirements:
 - Full rich debugging, ideally in Visual Studio
 - Asserts
 - Internal kernel profiling
 - Hot-swapping / edit-and-continue of kernels
- Opportunity for IHVs and platform providers to innovate here!
 - Goal: Cross-vendor standard
 - Think of the co-development of *Nvidia Cg* and HLSL

Unified development environment

- Want to debug/profile task- & data-parallel code seamlessly
 - On all processors! CPU, GPU & manycore
 - From any vendor = **requires standard APIs or ISAs**
- *Visual Studio 2010* looks promising for task-parallel PC code
 - Usable by our offline tools & hopefully PC runtime
 - Want to integrate our own JobManager
- *Nvidia Nexus* looks great for data-parallel GPU code
 - Eventual must have for all HW, how?
 - Huge step forward!



VS2010 Parallel Tasks