



GDC Tutorial:
Advanced OpenGL Game Development

**A Practical and Robust
Bump-mapping Technique for
Today's GPUs**

March 8, 2000



Mark J. Kilgard
Graphics Software Engineer
NVIDIA Corporation

“Hardware Bump Mapping” mostly hype so far




Previous techniques

- Prone to aliasing artifacts
- Do not correctly handle surface self-shadowing
- Cut too many corners, i.e. fail to renormalize, etc.
- Too expensive: many passes (OIM), needs accumulation buffer or `glCopyPixels`
- Only handle infinite, non-attenuated, non-spotlight lights
- Assume local viewer
- Correct only for essentially flat surfaces
- Difficult to implement and author

GAMEDevelopers
CONFERENCE 2000

Demo!



Overview of the Technique



Uses new NVIDIA GPU features

- Cube maps - for per-pixel normalization
 - EXT_texture_cube_map extension
- Register combiners - for per-pixel dot products and additional math
 - NV_register_combiners
- And dual-textured - for normalization cube map and 2D normal map
 - ARB_multitexture



Features of the technique

Correctly handles

- Surface self-shadowing for both diffuse and specular contributions
- Diffuse minification so that bumped objects in the distance look dimmer than equivalent smooth distant objects
- Mipmap filtering of normal maps minimizes aliasing
- Local, attenuated, and/or spot-light light sources with local viewer
- Ambient and surface decals

Anatomy of the technique

Three passes (or just two if no specular)





High-level view

Tangent space per-pixel lighting [Peercy 97]

- CPU supplies tangent space light vector or half-angle vector as (s,t,r) texture coordinates
 - linearly interpolated
 - normalized by an RGB8 cube map texture
- Normal map is RGBA8 2D texture
 - RGB contains “compressed” normalized perturbed normal
 - Alpha contains averaged normal shortening
- But no hard-wired per-pixel “lighting engine”!

Per-pixel lighting building blocks



Cube maps enable per-pixel vector normalization

- Cube map normalizes interpolated (s,t,r) into RGB
 - (s,t,r) is interpolated per-vertex light or half-angle vector
- Combiners expand RGB from $[0,1]$ range to $[-1,1]$ range
 - expansion done by register combiners
- *Optimization:* If infinite viewer or infinite light used, half-angle can be computed within the cube map!
- 32x32x6 RGB8 cube map texture is sufficient

Per-pixel lighting building blocks




Register combiners perform per-pixel math

- Signed math for dot products
 - $L \cdot N'$ for diffuse, $H \cdot N'$ for Blinn specular
 - L and H properly normalized by cube map
 - Pre-normalized N' supplied by normal map
- Uses tangent space L_Z for geometric surface self-shadowing
- Constant color for adding in ambient contribution
- Successive squaring for specular exponent


Surface self-shadowing

Two kinds of self-shadowing

- $\max(0, L \cdot N')$ based on the perturbed normal
- Also should clamp when $L \cdot N$ goes negative!



Surface should self-shadow
due to perturbed normal,
i.e., $L \cdot N' < 0$



Surface should self-shadow
due to unperturbed normal,
i.e., $L \cdot N < 0$

Surface self-shadowing

Self-shadowing computation

- Technique supports bumped self-shadowing by computing
 $\min(8 * \max(0, L_z), 1) * \max(0, L \cdot N')$
- Also stashes this self-shadowing term in destination alpha
 $\min(8 * \max(0, L_z), 1)$
- Specular pass blends with DST_ALPHA, ONE computing
 $\min(8 * \max(0, L_z), 1) * \max(0, H \cdot N')^8$
- Illumination does not appear on geometric “back side”
 - Steep ramp avoids winking and popping at self-shadow boundary



Complete lighting model


More per-pixel math than you might think!

- Ambient = constant
- Diffuse = $\min(8 * \max(0, L_z), 1) * \max(0, L \cdot N')$
- Specular = $\min(8 * \max(0, L_z), 1) * \max(0, H \cdot N')^8$
- Final = (Diffuse + Ambient) * Decal + Specular


Attenuation and/or spotlight

- Final = (Attenuate * Diffuse + Ambient) * Decal + Attenuate * Specular

Self-shadowing example



*Light leaks onto torus “back side”
due to bad geometric self-shadowing*



*Proper geometric
self-shadowing!*



Some CPU work required

Per-vertex work

- Model requires per-vertex tangent space basis
 - any bump mapping scheme needs surface normals (N), tangents (T), and binormals (B)
- If light or viewer (in the specular case) changes relative to the object, CPU re-compute tangent space light and half-angle
 - Transform object-space vector by [T B N] 3x3 matrix
- *Nice split:* CPU does per-vertex work, GPU does transform and per-pixel work

Issues


Truth in advertising

- Specular exponent only 8 so not too shiny
 - banding evident in specular highlights, though easy to hide artifacts in surface decal and diffuse contribution
- Perturbed normal slightly de-normalized by filtering
 - causes very slight dimming, alternative is aliasing
 - most objectionable when magnifying, “grated” look
- Not “bumped environment mapping”
 - but DX6 bump-env support is a joke
- Requires GeForce 256, Quadro, or any future NVIDIA GPU

Claim: practical

Justification for this claim

- Real-time frame rates on current generation GPUs
 - tweakable: eliminate specular or decal pass, etc.
- Mixable with other multi-pass rendering techniques
 - stenciled shadows, multiple lights, fog, etc.
- Bump map easy to author as an 8-bit gray-scale image
 - examples:





Claim: robust

Justification for this claim

- Animates with minimal temporal aliasing
- Models both self-shadowing terms, including for specular
- Diffuse illumination filtered reasonably
- Light and half-angle vectors renormalized per-pixel
- Local, attenuated, and spot-light light sources
- Local viewer
- Overall, very reasonable fidelity to Blinn's original formulation

Source available on the web already



Bumpdemo source code

- Go to www.nvidia.com Developer Section
 - http://www.nvidia.com/developer.nsf/htmlmedia/devMainFR_top.html
- Find and download bumpdemo.zip
 - requires GeForce 256 or Quadro to run
- Also download “NVIDIA OpenGL Extension Specifications”
 - nvOpenGLspecs.pdf
 - Adobe Acrobat, 197 pages
 - documents OpenGL extensions used by bumpdemo

Other novel uses for cube mapping




Stable specular highlights

- Encode specular & diffuse lighting solution
 - Diffuse cube map using *normal* map texgen
 - Specular cube map using *reflection* map texgen
 - Encode unlimited number of directional lights
- Result is stable specular highlights
- Less significant for diffuse lighting
 - Average of dot products \approx dot product of averages
 - Excepting clamping


Example for stable specular highlights

GAMEDevelopers
CONFERENCE 2000

Cube map
lighting



Standard
per-vertex lighting




Bright, stable
specular highlight,
even at low tessellation.

Poor per-vertex
sampling of the highlight.
Wobbles during animation.


Another example of stable specular highlights



High tessellation does not completely fix per-vertex artifacts



Still bright and stable.



Better sampled, but still has streaky artifacts.


Still more novel uses for cube maps & combiners



Per-pixel specular normal mapping

- Assumes fancy per-pixel dot product operations!
- Cube map encodes normalized half-angle
 - Texgen supplies view vector
 - Cube map generates normalize($V+L$) vector
- Another 2D texture supplies per-pixel surface normals
- Per-pixel specular lighting
- Per-pixel “normalized($V+L$) dot N” !

One pass per-pixel diffuse & specular example!




OpenGL per-vertex
lighting with mesh of
thousands of vertices
(note displacement).




Per-pixel specular normal
mapping with cube maps
(flat, but lighting all there).

Example in wireframe




Dense, dense mesh.



One polygon!


Object-space per-pixel lighting example

GAMEDevelopers
CONFERENCE 2000



Object-space per-pixel example in wire-frame

GAMEDevelopers
CONFERENCE 2000



The detail is per-pixel!


Register combiners overview



- overrides texture stages/environment, color sum, and fog in current APIs
- signed math (negative one to positive range)
 - extended range through scaling
- dot products for lighting and image processing applications
 - designed for specular, diffuse, and ambient per-pixel lighting
 - object space bump map lighting
 - tangent space bump map lighting
 - post-filtering 3x3 color matrix for color space conversions
- register model supports non-linear data flows
 - superior to linear chain in current APIs
- effectively, a VLIW instruction set for fragment coloring
- very efficient hardware implementation

Register combiners operational overview


GAMEDevelopers
CONFERENCE 2000



General combiners

RGB operation


GAMEDevelopers
CONFERENCE 2000



General combiners

Alpha operation

GAMEDevelopers
CONFERENCE 2000



Final combiner operation

GAMEDevelopers
CONFERENCE 2000

