



Shiny Pixels and Beyond: Real-Time Raytracing at SEED

Johan Andersson & Colin Barré-Brisebois
Electronic Arts




SEED

Watch the trailer here:
<https://www.youtube.com/watch?v=LXo0WdIEJk>

“PICA PICA”

Exploratory mini-game & world

- For our self-learning AI agents to play, not for humans ☺
- Uses SEED's **Halcyon** R&D engine
- Goals
 - Explore hybrid raytracing with DXR
 - Clean and consistent visuals
 - Procedurally-generated worlds
 - No precomputation



Why raytracing?

- Flexible new tool in the toolbox
- Solve sparse & incoherent problems
- Unified API + performance (DXR + RTX)
- Simple high quality - easy ground truth




Self-Learning AI


Using deep reinforcement learning

- “Imitation Learning with Concurrent Actions in 3D Games” [Harmer 2018]
- 36 semantic views - 1550 fps
- Training with TensorFlow
- Future: Inference with WinML

See “Deep Learning - Beyond the Hype” tomorrow



Hybrid Rendering Pipeline



Deferred shading (**raster**)



Direct shadows
(**raytrace** or **raster**)



Direct lighting (**compute**)



Reflections (**raytrace**)



Global Illumination (**raytrace**)



Ambient occlusion
(**raytrace** or **compute**)



Transparency & Translucency
(**raytrace**)




Post processing (**compute**)



- **Spawn a Mesh?**
 - DXR: build its bottom acceleration structure
 - Multiple geometries for multiple materials
 - Triangles, AABBs, custom
 - Mesh instances specified in top acceleration


- **Move a Mesh?**
 - Update the instance's position/orientation in the top acceleration

- **Spawn [some] Rays?**
 - Multiple Hit and Miss shaders possible



Raytraced Reflections

- Rasterize primary visibility
- Launch rays from the G-Buffer
- Raytrace at half resolution
- Reconstruct at full resolution
 - Spatiotemporal filtering
- Works on both flat and curved surfaces



Raytraced Reflections

Reflection Rays

Let's launch some reflection rays:

1. Select one of the (2x2) pixels to trace
2. Reconstruct position and vectors
3. Initialize Halton & random number seq.
4. Initialize the payload
5. Prepare a new ray
6. Trace
7. Gather results from ray payload
 - Reflection Color, Direction, HitT, 1/pdf

```

RaytracingAccelerationStructure g_rtScene : register(t0, space0);
ConstantBuffer<RaytracingConstants> g_rt : register(b0, space0);
RWTexture2D<float4> g_tx0 : register(u0, space0);
RWTexture2D<float4> g_tx1 : register(u1, space0);

[shader("raygeneration")]
void reflectionRaygen()
{
    uint2 px = DispatchRaysIndex();
    uint2 launchDim = DispatchRaysDimensions();

    // Select one of the four full-res pixels to trace from
    const uint2 noiseCoord = (px / 2 + (g_rt.frameIndex / 2) * uint2(3, 7)) & 31;
    const uint subpixelIdx = g_blueNoise[noiseCoord.y * 32] + g_rt.frameIndex;
    const uint2 gbufferPx = px * 2 + uint2(subpixelIdx & 1, (subpixelIdx >> 1) & 1);

    // Reconstruct position and the various vectors
    const Gbuffer gbuffer = loadbuffer(gbufferPx, g_gbuffer);
    const float depth = g_depth[gbufferPx];
    float2 uvPos = (gbufferPx + 0.5f) / g_rt.viewDimensions.xy;
    float4 csPos = float4(uvToCs(uvPos), depth, 1.0f);
    float4 wsPos = mul(g_rt.clipToWorld, csPos);
    const float3 P = wsPos.xyz / wsPos.w;
    const float3 E = g_rt.eyeworldPosition;
    const float3 V = normalize(E - position);
    const float3 N = gbuffer.normal;

    // Initialize the Halton sequence for each ray and random number generator to rotate the sequence
    Halton halton = haltonInit(px, g_rt.frameIndex, ...);
    uint seed = randomInit(px);

    // Initialize a new payload
    Payload payload = payloadInit(seed, halton);

    // Initialize a new ray
    RayDesc ray;
    ray.Origin = position + V * max(1, length(frame.position)) * 1e-4f; // epsilon
    ray.Direction = sampleDirectionFromBrdf(gbuffer, V, N, halton, seed); // stochastic BRDF
    ray.TMin = 0;
    ray.TMax = 1000; // perf: ray stops at 1000 meters

    // Launch the ray
    TraceRay(g_rtScene, RAY_FLAG_CULL_FRONT_FACING_TRIANGLES, RaytracingInstanceMaskAll,
              HitType_Primary, SbtRecordStride, MissType_Primary,
              ray, hitData);


    g_tx0[px] = float4(hitData.lighting.rgb, depth);
    g_tx1[px] = float4(ray.Direction * hitData.hitT, 1.0f / brdfSample.pdf);
}

```

Reflection Filtering

Inspired by Stochastic Screen-Space Reflections
[Stachowiak 2015]

- For every full-res pixel, sample 16 pixels in half-res ray results
 - Blue Noise offsets, decorrelated every 2x2 pixels
- Build color bounding box of ray-hit results
 - Clamp temporal history to bounding box
- Followed by a variance-driven bilateral filter
 - Helps with rough reflections



Unfiltered (Top) and Filtered (Bottom) Results





Screen-Space Reflections



G-Buffer Raytraced



Path Tracing Reference



Screen-Space Reflections



G-Buffer Raytraced



Path Tracing Reference

Materials

Combine multiple microfacet surface layers into a single, unified & expressive BRDF

- Inspired by Arbitrarily Layered Micro-Facet Surfaces [Weidlich 2007]
- Unified for all lighting & rendering modes
 - Raster, path-traced reference, and hybrid
- Energy conserving & Fresnel
- Rapidly experiment with different looks
 - Bake down number of layers for production



Objects with Multi-Layered Materials

Materials

Standard

- Aluminum
- Brushed Aluminum
- Coated Carbon
- Copper
- Silver Satin
- Shiny / Mat Plastic
- Dark Rubber

Exotics*

- Glass
- Rough Glass
- Jade / Marble


* Integrated in BRDF, but built separately



Multi-Layered Materials

BRDF Sampling

- General idea: Launch one ray for the whole stack
 - Stochastically select a layer & sample
 - When evaluating material, estimate visibility vs other layers
 - Some energy is reflected, some is refracted:
 - Microfacet: Fresnel ($\text{refracted} = 1 - \text{Fresnel}$)
 - Diffuse: reflected fraction equal to albedo, remaining light absorbed; no refraction
 - Sample BSDF of selected layer
 - Attenuated by layer(s) on top
- Result is temporally filtered
 - A single value for many layers requires clever filtering



Multi-Layered Materials [Weidlich 2007]

Transparency & Translucency

Raytracing allows to unify reflections and refractions



- **Glass**
 - Order-independent (OIT)
 - Handles multiple IOR transitions
- **Translucency**
 - Inspired from *Translucency in Frostbite* [Barré-Brisebois 2011]
 - Inner structure scattering based on lighting traveling inside the medium
- Performance: we do it in texture-space



Glass and Translucency


Translucency Breakdown

- For every valid position & normal




Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside




Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
 - Alternatively, front + back cosine lobes
 - Perf: only do n-rays per frame




Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
 - Alternatively, front + back cosine lobes
 - Perf: only do n-rays per frame
- Compute lighting at intersection
 - Sample previous translucency result




Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
 - Alternatively, front + back cosine lobes
 - Perf: only do n-rays per frame
- Compute lighting at intersection
 - Sample previous translucency result
- Gather
 - Modulate with Beer-Lambert or Henyey-Greenstein phase function



Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
 - Alternatively, front + back cosine lobes
 - Perf: only do n-rays per frame
- Compute lighting at intersection
 - Sample previous translucency result
- Gather
 - Modulate with Beer-Lambert or Henyey-Greenstein phase function
- Store new BTDF value



Translucency Filtering

Result converges over a couple frames

- Denoised or temporally-accumulated
- Temporal: build an update heuristic
 - Exponential moving average can be OK
 - Game-specific threshold to update
 - Reactive enough for moving lights & objects
- Variance-adaptive mean estimation



Translucency Shadowing

Raytracing allows for globally shadowed translucency

- Global phenomena
 - Integration via feedback
 - Objects occlude each other
- Overall more grounded and visually-convincing translucency



Direct Lighting




Self-Shadowed Translucency

Transparency

Similar approach is used for glass

- Launch ray using view's origin and direction
- Refract based on medium's index-of-refraction (IOR)
 - Snell's law: `refract(ray, N, iorInput / iorOutput)`
 - DXR: `HitKind()` to handle IOR transitions (Air \leftrightarrow Glass)
- Trace a ray in the scene & sample lighting
- Tint the result by glass color + chromatic aberration



Transparency

Similar approach is used for glass



- Launch ray using view's origin and direction
- Refract based on medium's index-of-refraction (IOR)
 - Snell's law: `refract(ray, N, iorInput / iorOutput)`
 - DXR: `HitKind()` to handle IOR transitions (Air \leftrightarrow Glass)
- Trace a ray in the scene & sample lighting
- Tint the result by glass color + chromatic aberration
- Pen: we don't handle transparent shadows yet



Transparency

Works for clear and rough glass

- Clear
 - No filtering required
- Rough / Blurry
 - Open cone angle with Uniform Cone Sampling [PBRT]
 - Wider cone → more samples
 - Or temporal filtering
- Tint with phase function and more complex BTDF



Global Illumination

- We want a technique that:
 - Doesn't require any precomputation
 - Doesn't require parametrization (UVs, proxies)
 - Works for both static and dynamic scenes
 - Adaptive & refines itself on-the-fly
- Point-based / surfels for a dynamic world
 - Runtime Monte-Carlo integration



S E E D // Shiny Pixels and Beyond: Rendering Research at SEED


Surfel Placement



Surfel Spawning From Camera @ 1% speed

S E E D // Shiny Pixels and Beyond: Rendering Research at SEED


Surfel Placement



Skinned Surfels for Dynamic Objects

Sampling & Integration

- Compute surfel irradiance by PT
- When shooting rays/frame
 - Limit depth and number of paths
- Limiting depth != fewer bounces
 - Reuse results from previous frames
 - 1 = radiosity, ∞ = path tracing
- Variance-adaptive mean estimator
- Application (Binning, culling, half-res apply + upsample)



Progressive Integration (slowed down for clarity)

Shadows


Accumulated and filtered in screen-space




- Raygen: Launch a ray towards light
 - Payload's **miss** flag set to **true** (from Miss Shader) if it doesn't hit geometry
 - Penumbra driven by *uniform cone sampling* [PBRT]
- Temporal Reprojection
 - Accumulates shadow and variance + TAA-style bounding box clamping
- Filter (SVGF-like) [Schied and NVIDIA 2017]
 - Multipass weighted spatial blur, driven by variance from temporal accumulation




Hard Raytraced Shadows



Soft Raytraced Shadows (Unfiltered)



Soft Raytraced Shadows (Filtered)



Hard Raytraced Shadows


Soft Raytraced Shadows (Unfiltered)

Soft Raytraced Shadows (Filtered)

Texture Level-of-Detail

What about texture level of detail?

- Mipmapping [Williams 1983] is the standard method to avoid texture aliasing:




Left: level-of-detail (λ), partial derivatives and the parallelogram-approximated texture-space footprint of a pixel. Right: mipmap chain

- Screen-space pixel maps to approximately one texel in the mipmap hierarchy
- Supported by all GPUs for rasterization via shading quad and derivatives

Texture Level-of-Detail

No shading quads for ray tracing!

- Traditionally: *Ray Differentials*
 - Estimates the footprint of a pixel by computing world-space derivatives of the ray with respect to the image plane
 - Have to differentiate (virtual offset) rays
 - Heavier payload (12 floats) for subsequent rays (can) affect performance. Optimize!




Ray Differentials [Igehy99]

- Alternative: always sample mip 0 with bilinear filtering (with extra samples)
 - Leads to aliasing and additional performance cost

Texture Level-of-Detail

Together with  **NVIDIA**. Research, we developed a texture LOD technique for raytracing:

- Heuristic based on **triangle properties**, a **curvature estimate**, **distance**, and **incident angle**
 - Similar quality to ray differentials with single trilinear lookup
 - Single value stored in the payload for subsequent rays
- Upcoming publication by:
 - Tomas Akenine-Möller (NV), Magnus Andersson (NV), Colin Barré-Brisebois (EA), Jim Nilsson (NV), Robert Toth (NV)



Ground Truth, Ray Differentials, Ours, Mip0

Summary

- Just the beginning – important new tool going forward
- Unified API – easy to experiment and integrate
- Flexible but complex tradeoffs - noise vs ghosting vs perf
- Can enable very high quality cinematic visuals
- Lots more to explore – perf, raster vs trace, sparse render, denoising, new techniques

SEED @ GDC 2018

- **DirectX: Evolving Microsoft's Graphics Platform (presented by Microsoft)**
 - Johan Andersson and Colin Barré-Brisebois
 - Content will be available online soon at www.ea.com/seed
- **Deep Learning - Beyond the Hype**
 - Magnus Nordin
 - Room 2016, West Hall, Thursday, March 22nd, 11:30am - 12:30pm
- **Creativity of Rules and Patterns: Designing Procedural Systems**
 - Anastasia Opara
 - GDC Show Floor, Thursday, March 22nd, 12:30PM-1:00PM and Friday, March 23rd @ 11:00AM-11:30AM



Thanks

- **SEED**

- Joakim Bergdahl
- Ken Brown
- Dean Calver
- Dirk de la Hunt
- Jenna Frisk
- Paul Greveson
- Henrik Halen
- Effeli Holst
- Andrew Lauritzen
- Magnus Nordin
- Niklas Nummelin

- Anastasia Opara
- Kristoffer Sjöö
- Tomasz Stachowiak
- Ida Winterhaven
- Graham Wihlidal

- **Microsoft**

- Chas Boyd
- Ivan Nevraev
- Amar Patel
- Matt Sandy

- **NVIDIA**

- Tomas Akenine-Möller
- Nir Benty
- Jiho Choi
- Peter Harrison
- Alex Hyder
- Jon Jansen
- Aaron Lefohn
- Ignacio Llamas
- Henry Moreton
- Martin Stich

References

- **[Harmer 2018]** Jack Harmer, Linus Gisslén, Henrik Holst, Joakim Bergdahl, Tom Olsson, Kristoffer Sjöö and Magnus Nordin“ Imitation Learning with Concurrent Actions in 3D Games”. [available online](#)
- **[Barré-Brisebois 2011]** Barré-Brisebois, Colin and Bouchard, Marc. “Approximating Translucency for a Fast, Cheap and Convincing Subsurface Scattering Look”, [available online](#).
- **[Barré-Brisebois 2017]** Barré-Brisebois, Colin. “A Certain Slant of Light: Past, Present and Future Challenges of Global Illumination in Games”, [available online](#).
- **[Igehy 1999]** Igehy, Homan. “Tracing Ray Differentials”, [available online](#).
- **[PBRT]** Pharr, Matt. Jakob, Wenzel and Humphreys, Greg. “Physically Based Rendering”, Book, <http://www.pbrt.org/>.
- **[Schied 2017]** Schied, Christoph et. Al. “Spatiotemporal Variance-Guided Filtering: Real-Time Reconstruction for Path-Traced Global Illumination”, NVIDIA Research, [available online](#).
- **[Stachowiak 2015]** Stachowiak, Tomasz. “Stochastic Screen-Space Reflections”, [available online](#).
- **[Weidlich 2007]** Weidlich, Andrea and Wilkie, Alexander. “Arbitrarily Layered Micro-Facet Surfaces”, [available online](#).
- **[Williams 1983]** Williams, Lance. “Pyramidal Parametrics”, [available online](#).



SEED // SEARCH FOR EXTRAORDINARY EXPERIENCES DIVISION

STOCKHOLM – LOS ANGELES – MONTRÉAL – REMOTE

WWW.EA.COM/SEED

WE'RE HIRING!


Questions?

Bonus


Ambient Occlusion

Ambient Occlusion (AO) [Langer 1994] [Miller 1994]
maps and scales directly with real-time ray tracing:


- Integral of the **visibility** function over the hemisphere Ω for the point \bar{p} on a **surface** with normal \hat{n} with respect to the projected **solid angle**
- Games often approximate this in screen-space
- With RT, more grounded & improves visual fidelity!
 - Random directions \hat{w}
 - Can be temporally accumulated or denoised



$$A_{\bar{p}} = \frac{1}{\pi} \int_{\Omega} V_{\bar{p}, \hat{w}}(\hat{n} \cdot \hat{w}) d\omega$$




Screen-Space AO




Raytraced AO


Screen-Space AO



Raytraced AO
(Same Radius as SSAO)




Raytraced AO
(Far-field)



mGPU

Explicit Heterogenous Multi-GPU

- Parallel Fork-Join Style
- Resources copied through system memory using copy queue
- Minimize PCI-E transfers
- Approach
 - Run ray generation on primary GPU
 - Copy results in sub-regions to other GPUs
 - Run tracing phases on separate GPUs
 - Copy tracing results back to primary GPU
 - Run filtering on primary GPU



Ray Tracing Gems

Call for Papers

- A new book series with focus on real-time and interactive ray tracing for game development using the DXR API.
- We invite articles on the following topics:
Basic ray tracing algorithms, effects (shadows, reflections, tec), non-graphics applications, reconstruction, denoising, & filtering, efficiency and best practices, baking & preprocessing, ray tracing API & design, rasterization and ray tracing, global Illumination, BRDFs, VR, deep learning, etc.
- Important dates
 - 15th of October 2018: submission deadline for full papers
 - GDC 2019: publication of Ray Tracing Gems (paper version + e-book)
- Tomas Akenine-Möller will lead the editorial team
<http://developer.nvidia.com/raytracinggems/>

