



# Stochastic all the things: Raytracing in hybrid real-time rendering

Tomasz Stachowiak  
SEED – Electronic Arts



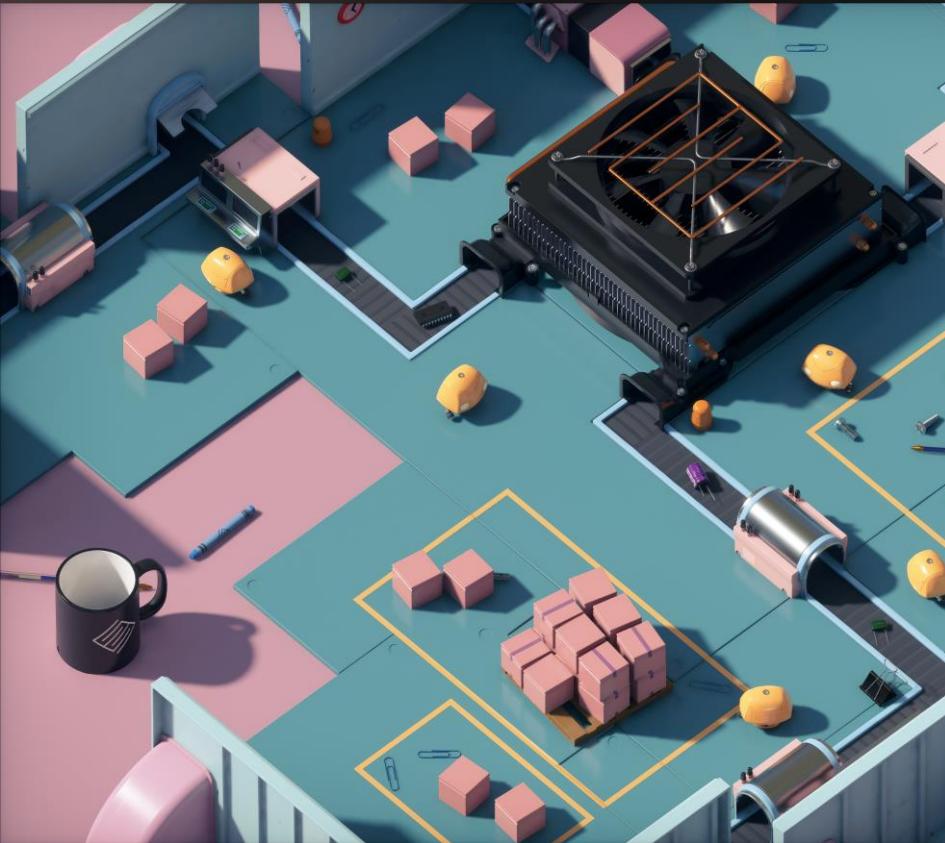
SEED

Watch the trailer here:

<https://www.youtube.com/watch?v=LXo0WdIELJk>

# “PICA PICA”

- Exploratory mini-game & world
- For our self-learning AI agents to play, not for humans ☺
  - “Imitation Learning with Concurrent Actions in 3D Games” [Harmer 2018]
- Uses SEED’s **Halcyon** R&D engine
- Goals
  - Explore hybrid rendering with DXR
  - Clean and consistent visuals
  - Procedural worlds [Opara 2018]
  - No precomputation

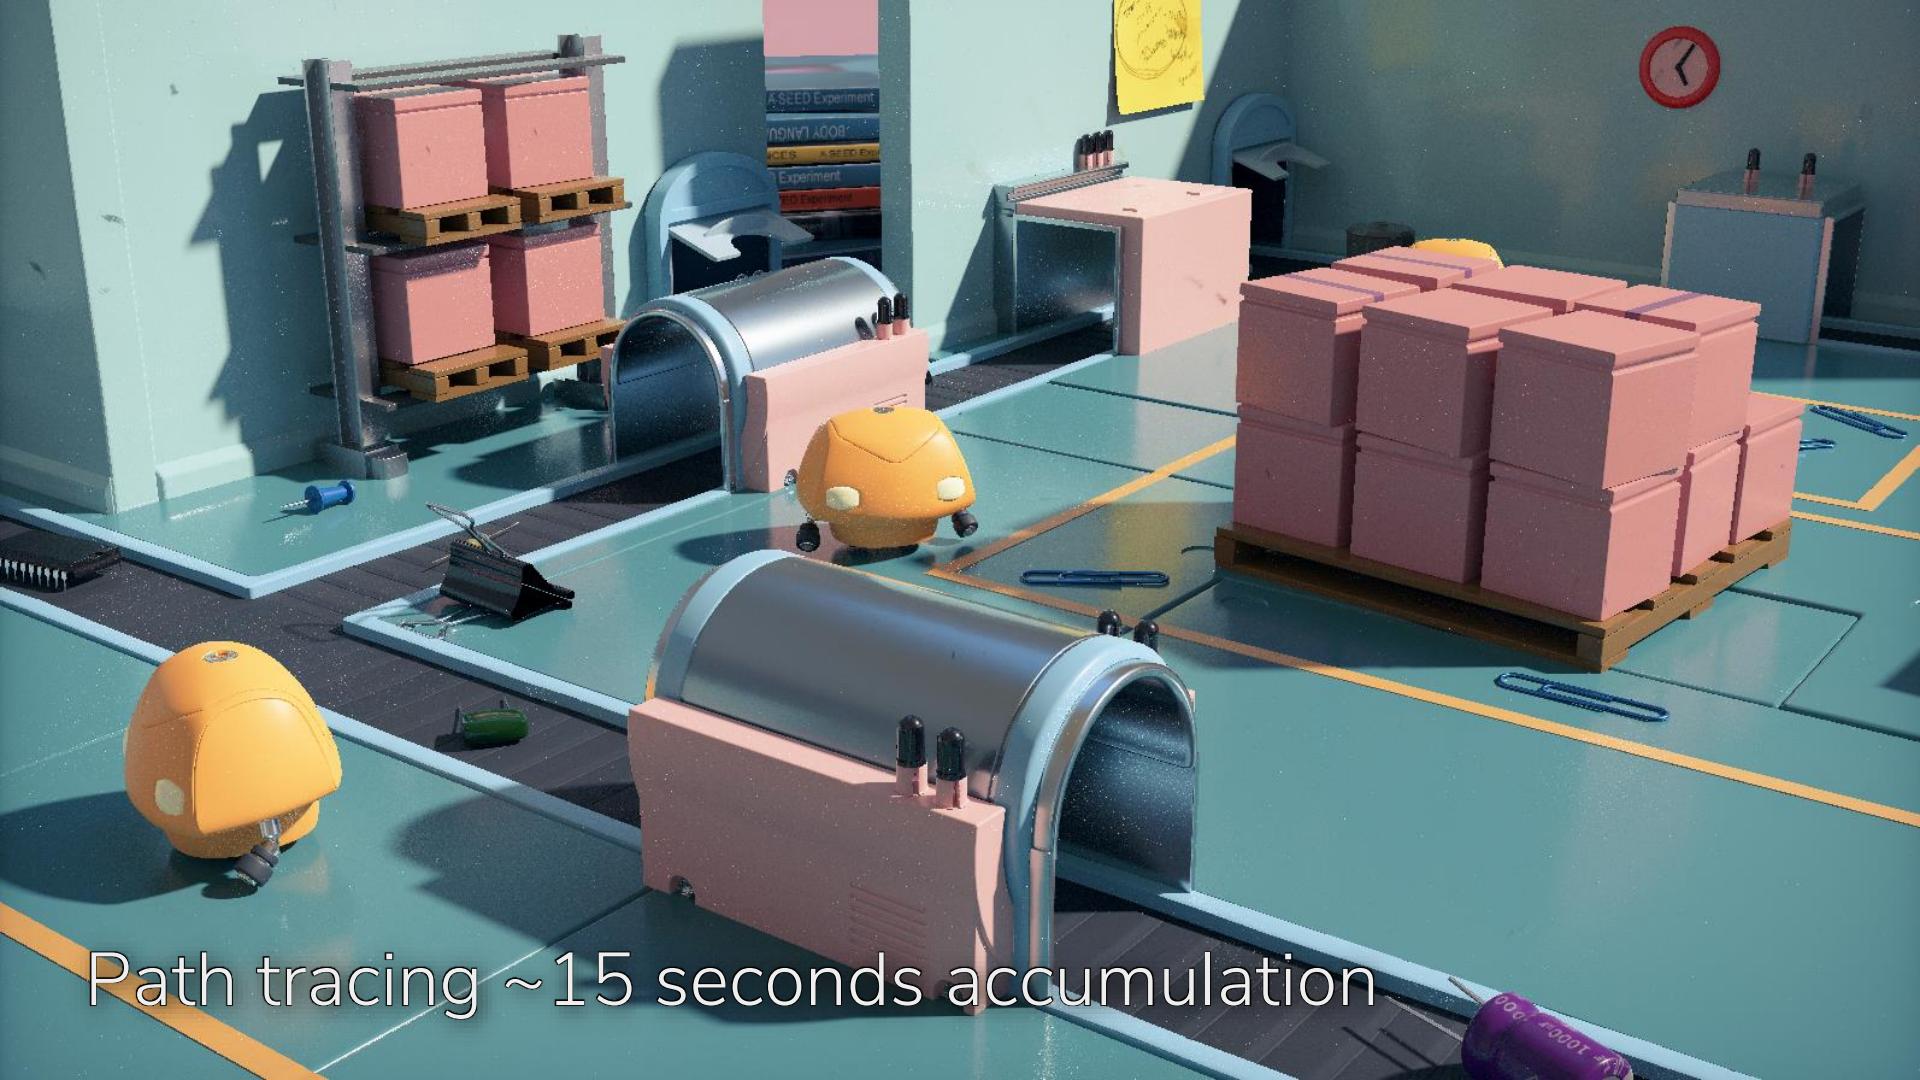


# Agenda

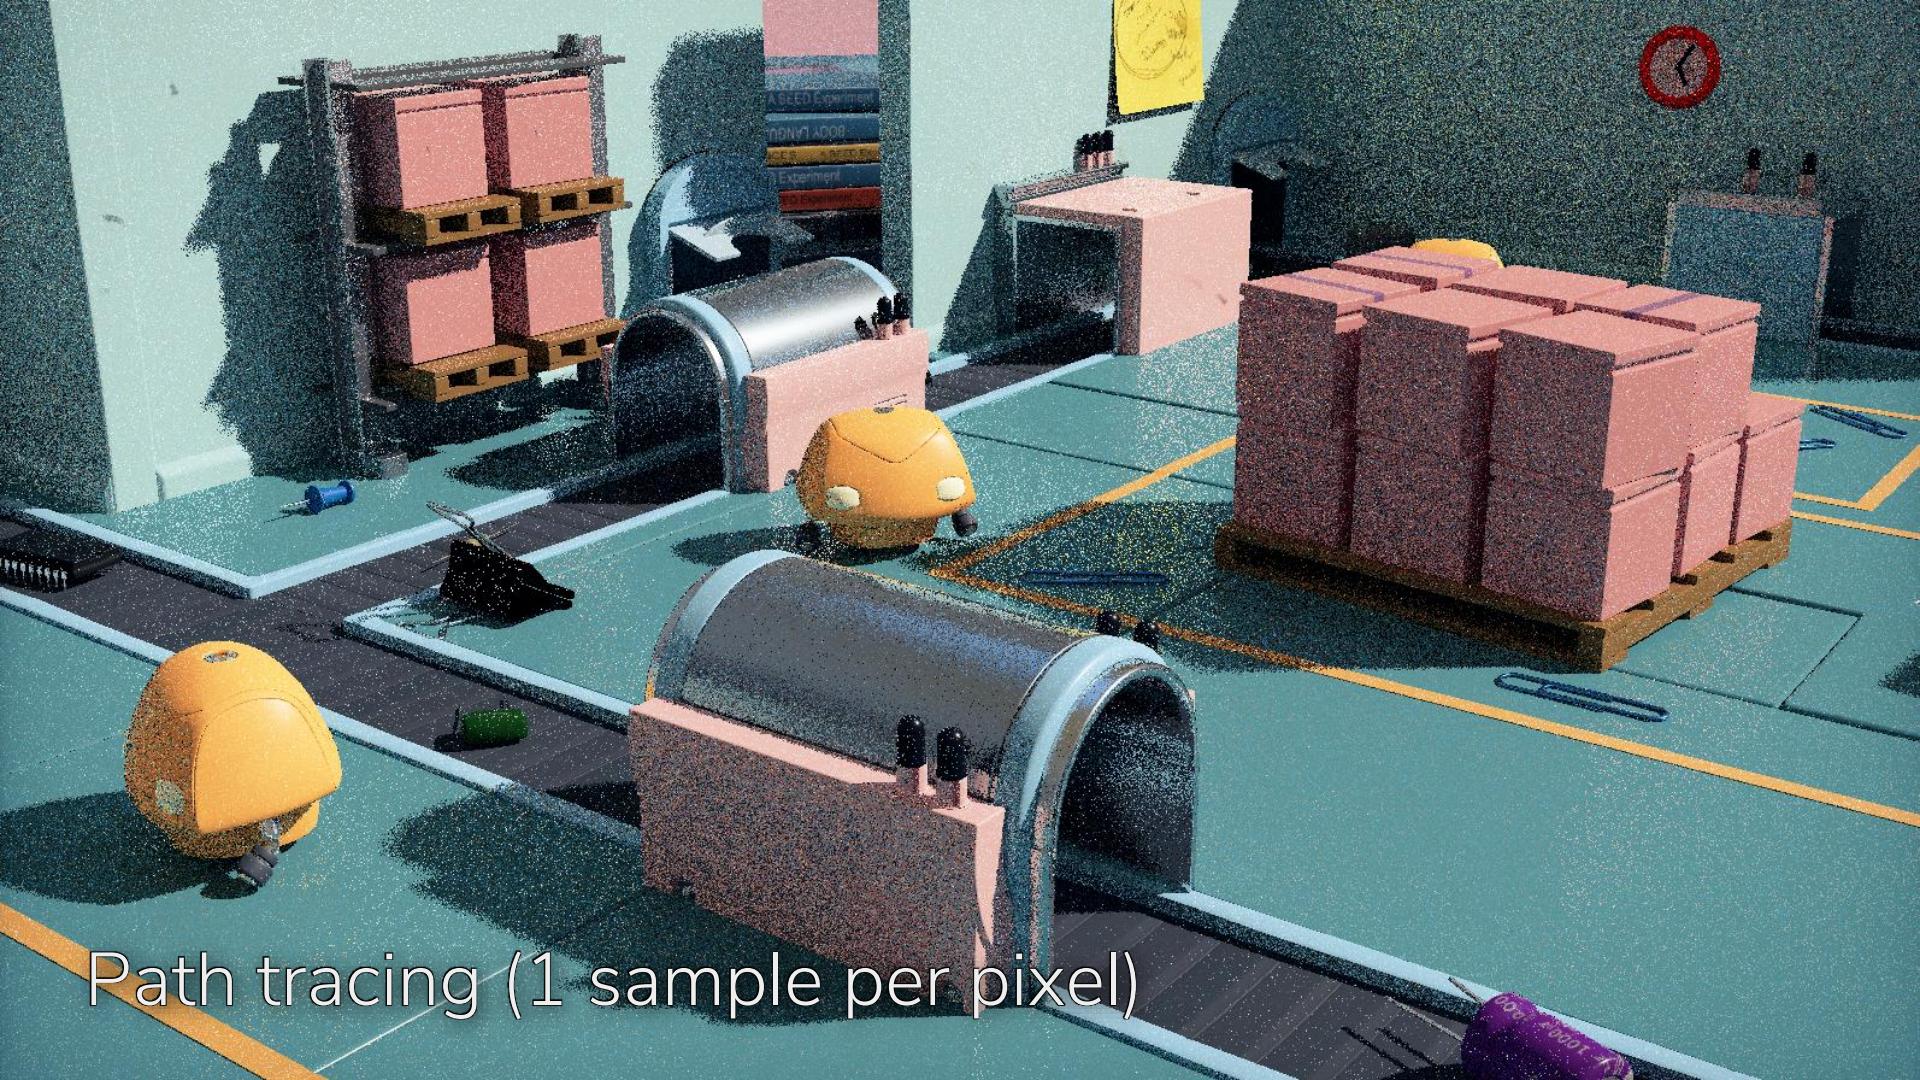
- Motivation
- Reflections
- Translucency and Transparency
- Global Illumination
- Shadows

A close-up of a small, yellow, spherical robotic device with a textured surface. It features a black circular base with two small black feet. On the top surface, there is a white 'CE' mark and the text 'MADE IN SWEDEN'. The device is positioned on a light blue surface with some yellow markings, possibly a floor or table. In the background, there are several stacked shipping containers in various colors like red, brown, and blue, suggesting a port or industrial setting.

Why use raytracing?



Path tracing ~15 seconds accumulation



Path tracing (1 sample per pixel)



Real-time hybrid with raytracing



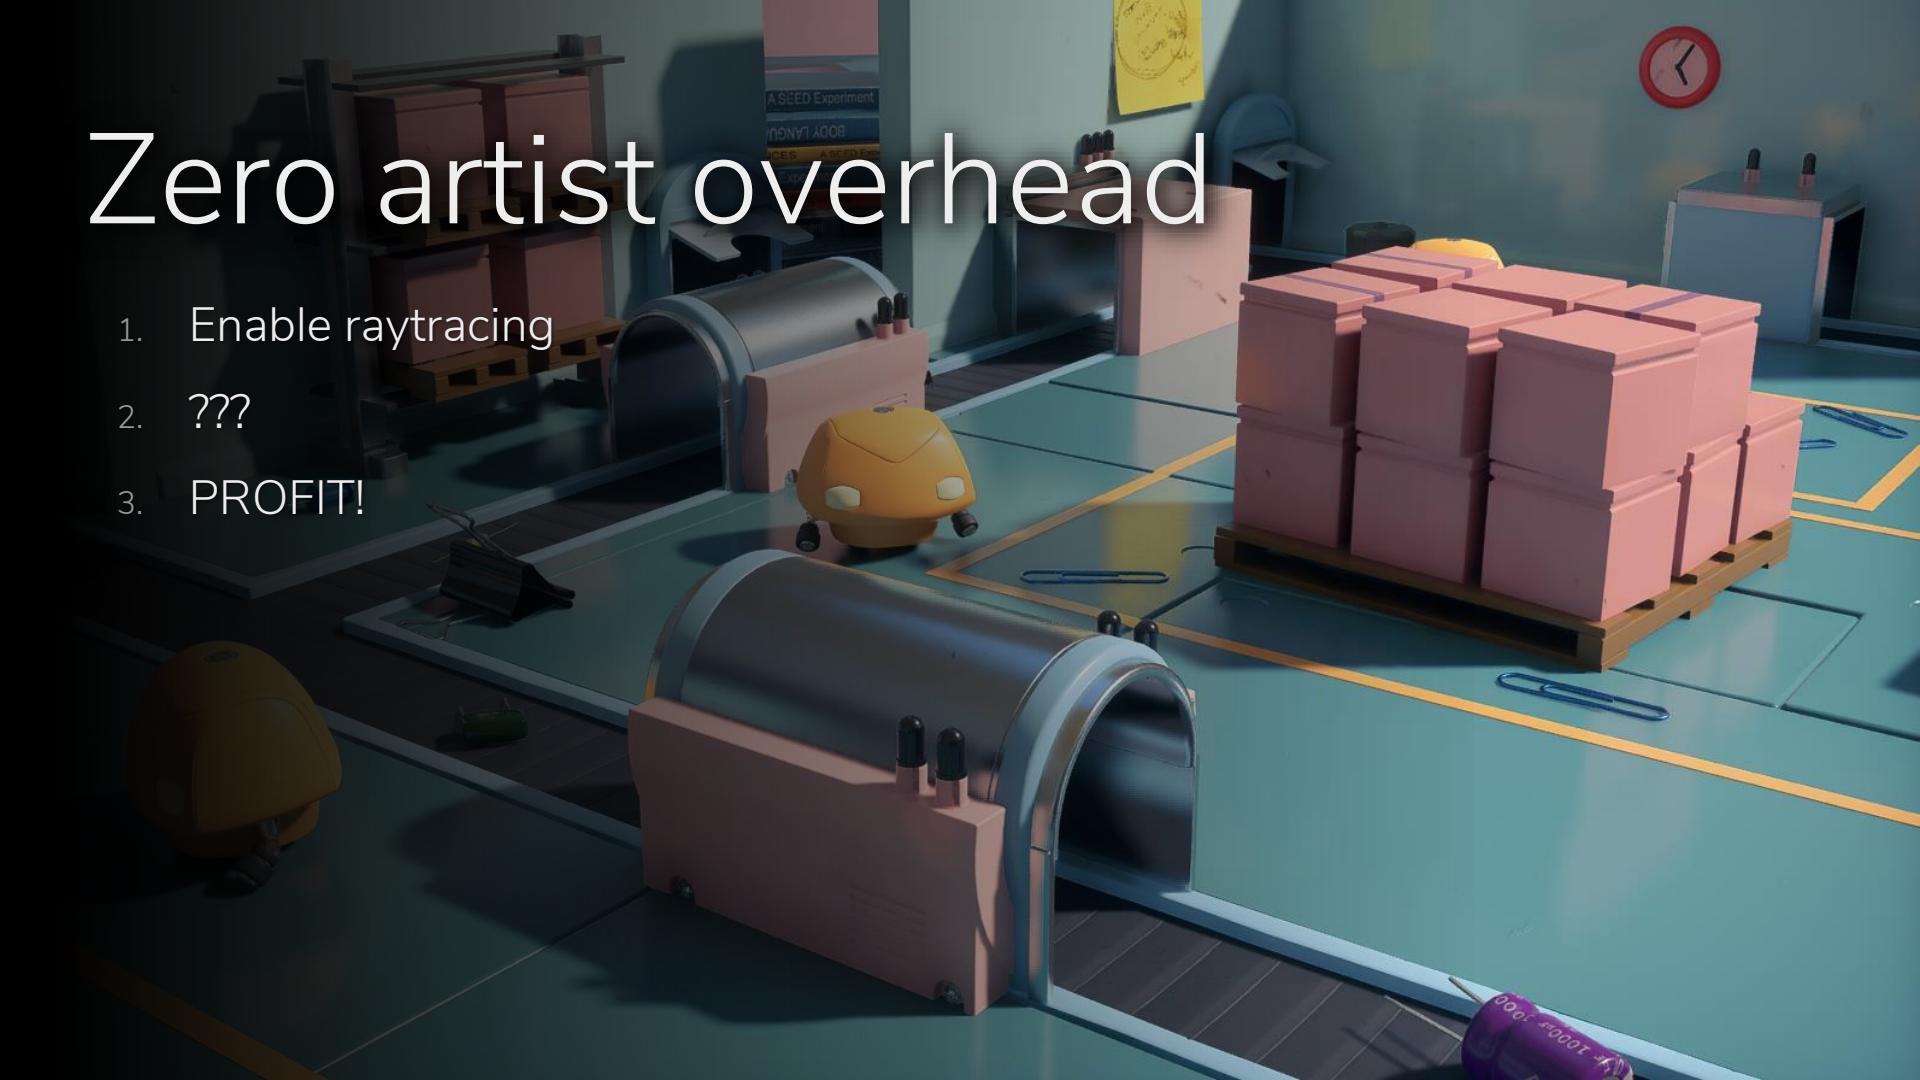
No raytracing

# Is this a fair comparison?

- “Classic” game solutions exist
  - Local reflection volumes
  - Planar reflections
  - Lightmaps
  - “Sky visibility” maps

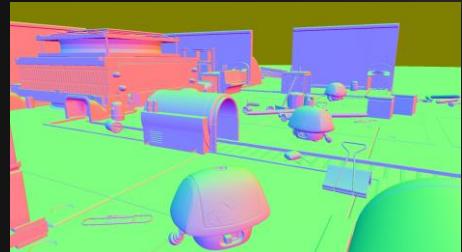
# Zero artist overhead

1. Enable raytracing
2. ???
3. PROFIT!



S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

# Hybrid Rendering Pipeline



Deferred shading (raster)



Direct shadows  
(raytrace or raster)



Direct lighting (compute)



Reflections  
(raytrace or compute)



Global Illumination (raytrace)



Ambient occlusion  
(raytrace or compute)



Transparency & Translucency  
(raytrace)



Post processing (compute)

# Materials

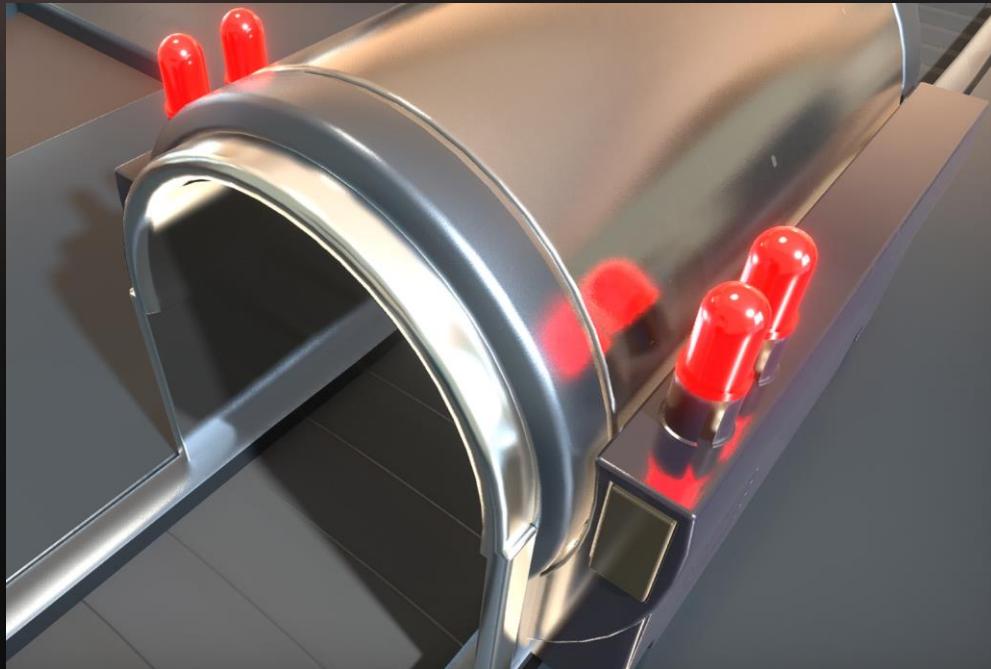
- Multiple microfacet layers
- Rapidly experiment with different looks
  - Bake down a number of layers for production
- Energy conserving
  - Automatic Fresnel between layers
- Inspired by Arbitrarily Layered Micro-Facet Surfaces [Weidlich 2007]
- Unified for all lighting & rendering modes
  - Raster, hybrid, path-traced reference



Objects with Multi-Layered Materials

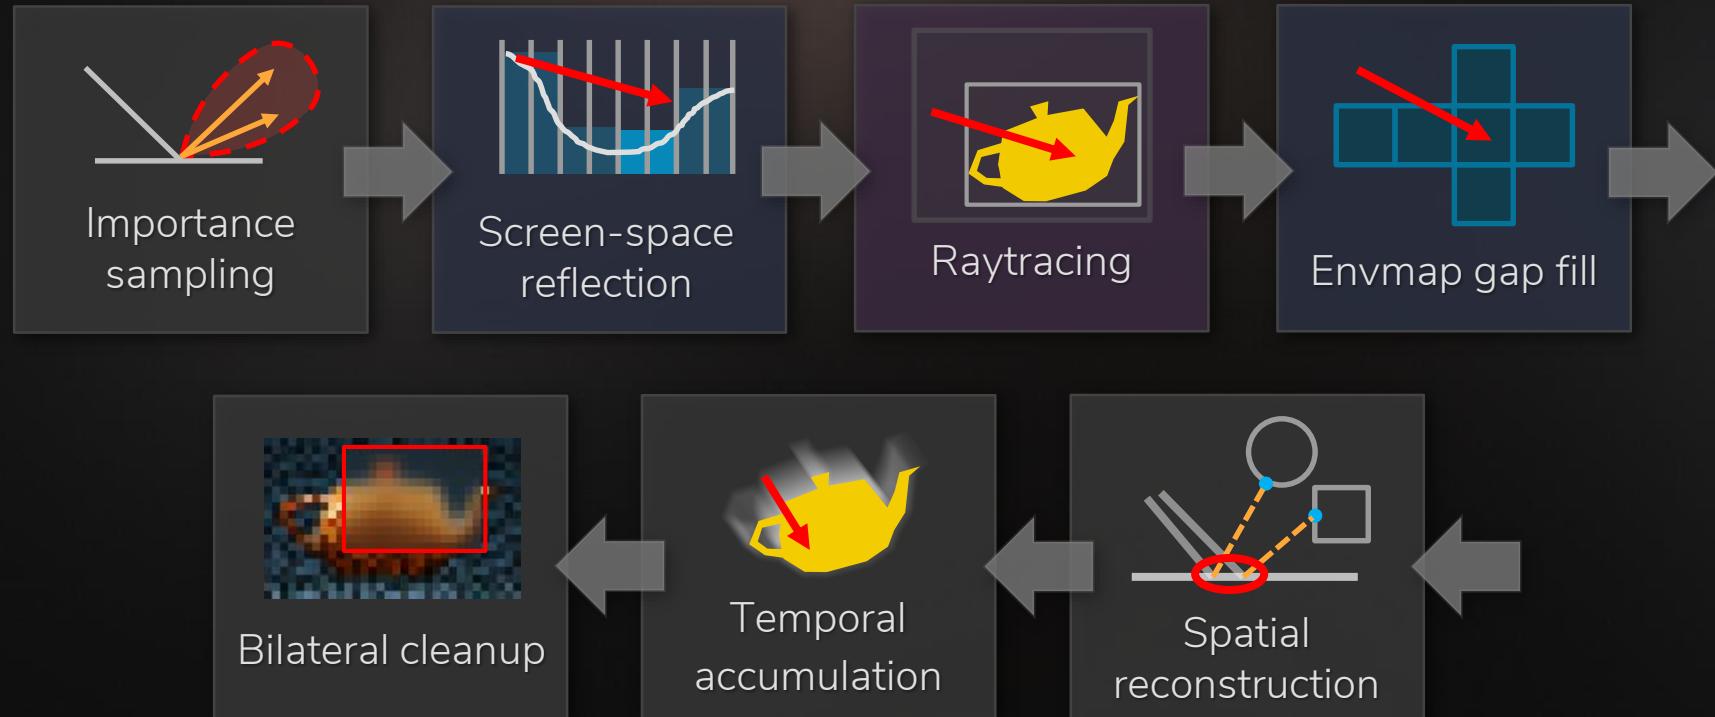
# Raytraced Reflections

- Launch rays from G-Buffer
- Raytrace at half resolution
  - $\frac{1}{4}$  ray/pixel for reflection
  - $\frac{1}{4}$  ray/pixel for reflected shadow
- Reconstruct at full resolution
- Arbitrary normals
- Spatially-varying roughness



Raytraced Reflections

# Reflection Pipeline



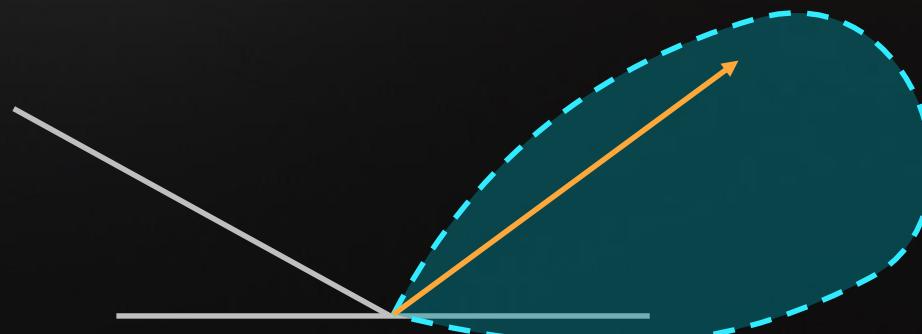
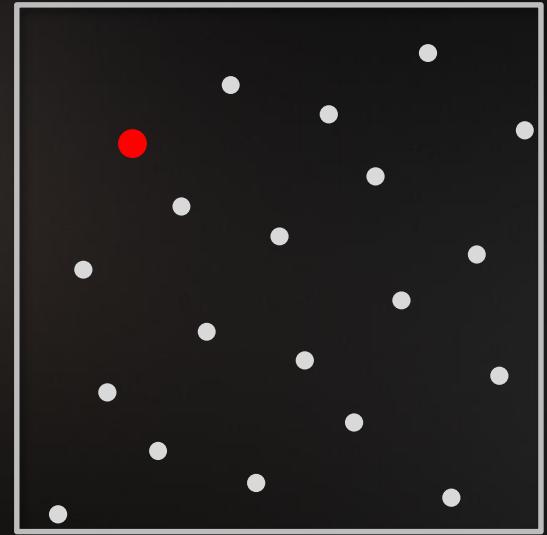
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]



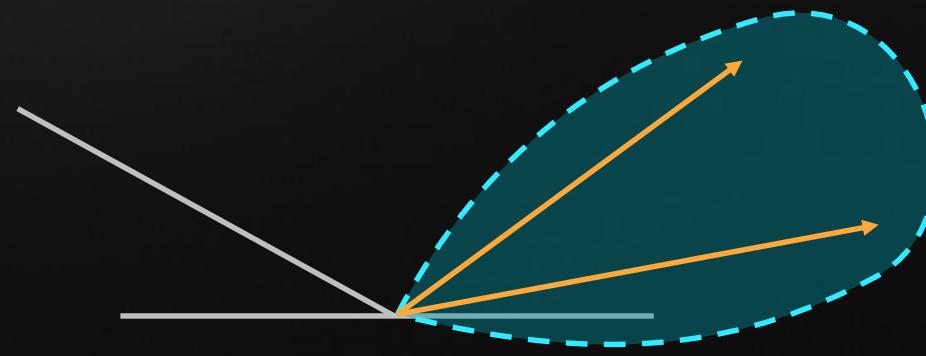
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]



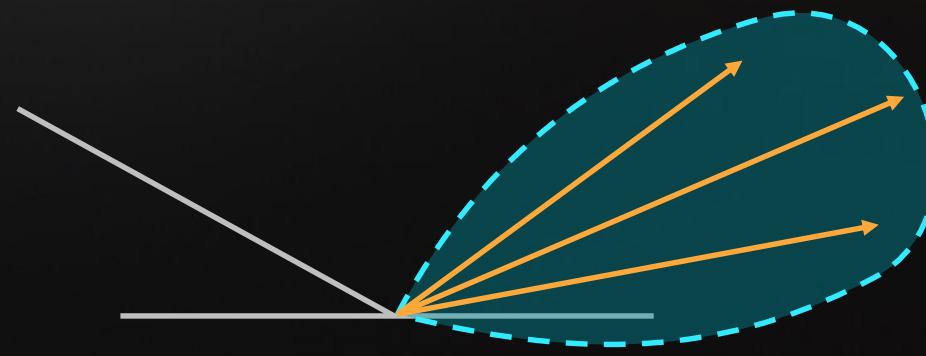
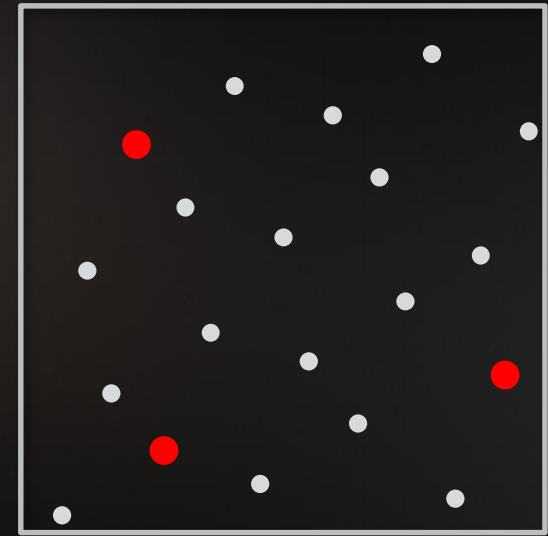
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]



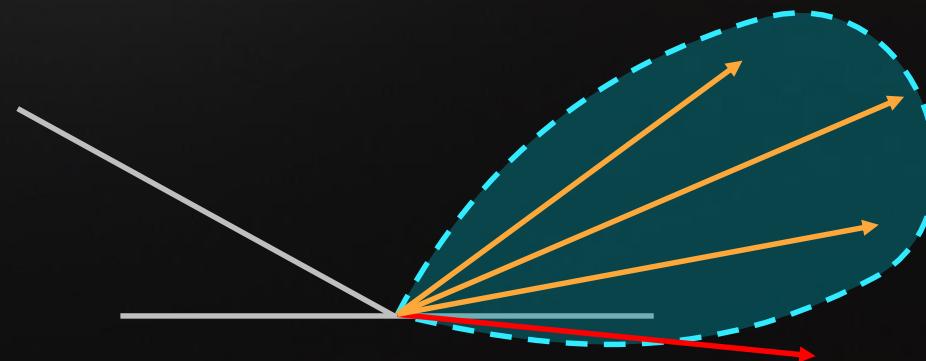
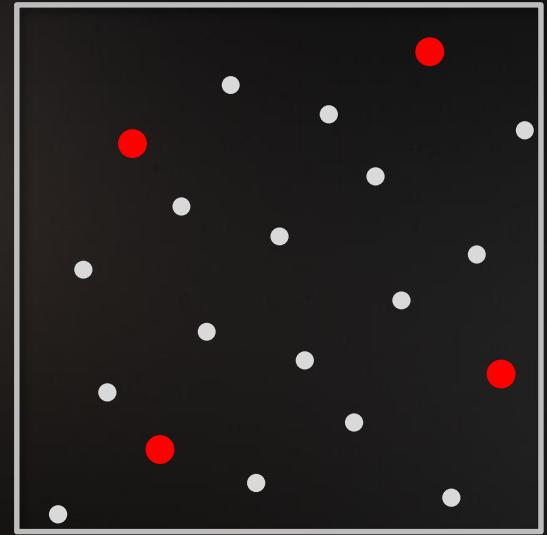
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]



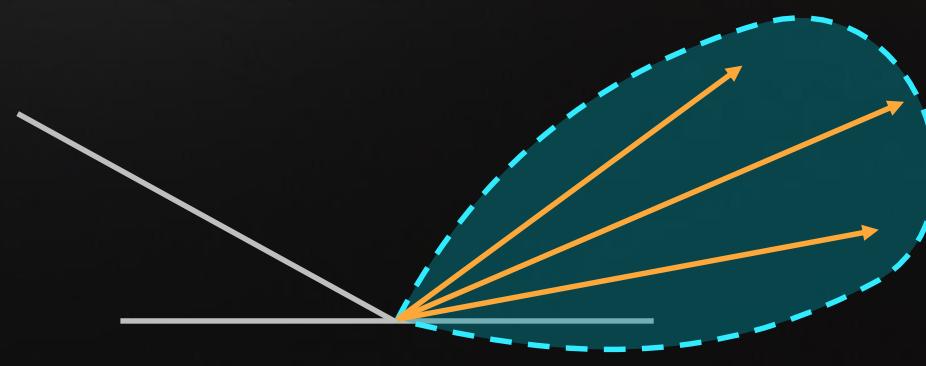
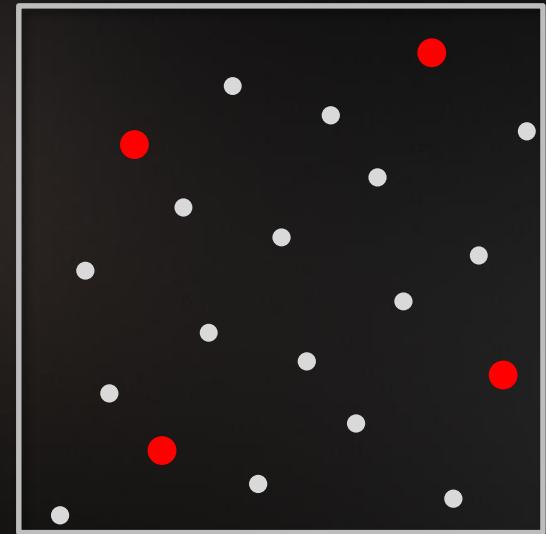
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
  - Cranley-Patterson rotation per pixel [PBRT]



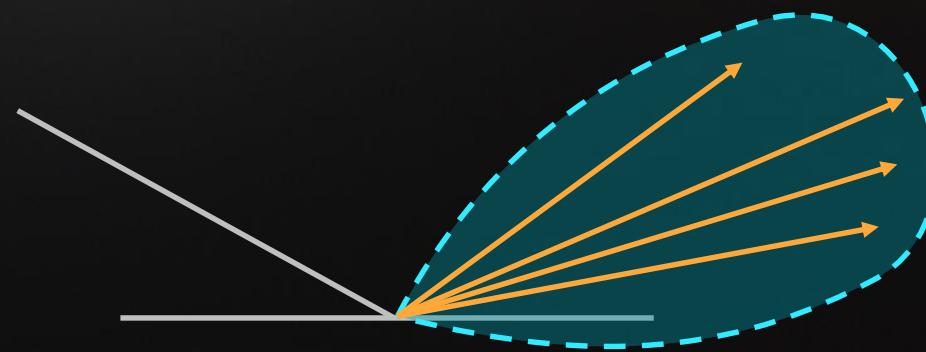
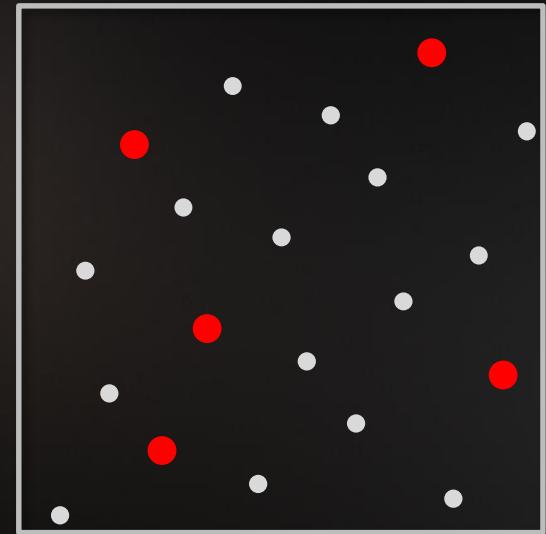
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]
  - Re-roll if ray below horizon



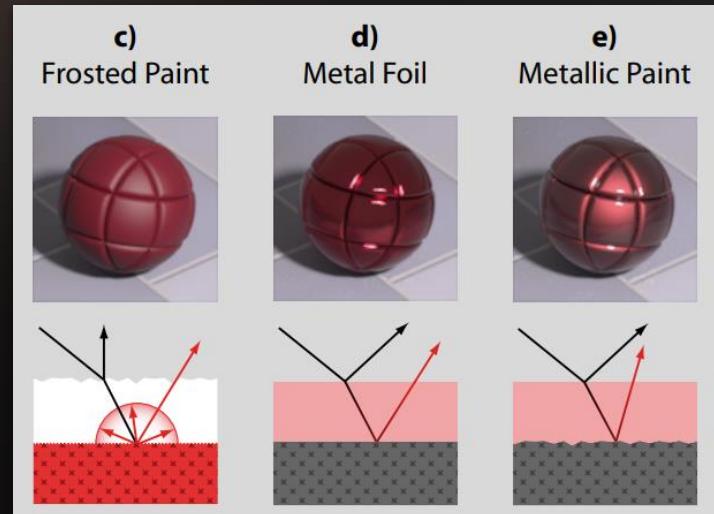
# Reflection Sampling

- Microfacet importance sampling
  - Normal distribution function (NDF) part of BRDF
- Very few rays, so quality matters
  - Low-discrepancy Halton sequence
    - Cranley-Patterson rotation per pixel [PBRT]
  - Re-roll if ray below horizon



# Layered Material Sampling

- One ray for the whole stack
  - Stochastically select a layer
  - Based on layer visibility estimate
    - Bottom layers occluded by top layers
    - View-dependent due to Fresnel
    - Approximate with Schlick
  - Sample BRDF of selected layer
  - Multiple layers can generate same direction
    - Query each layer about generated direction
    - Add up probabilities

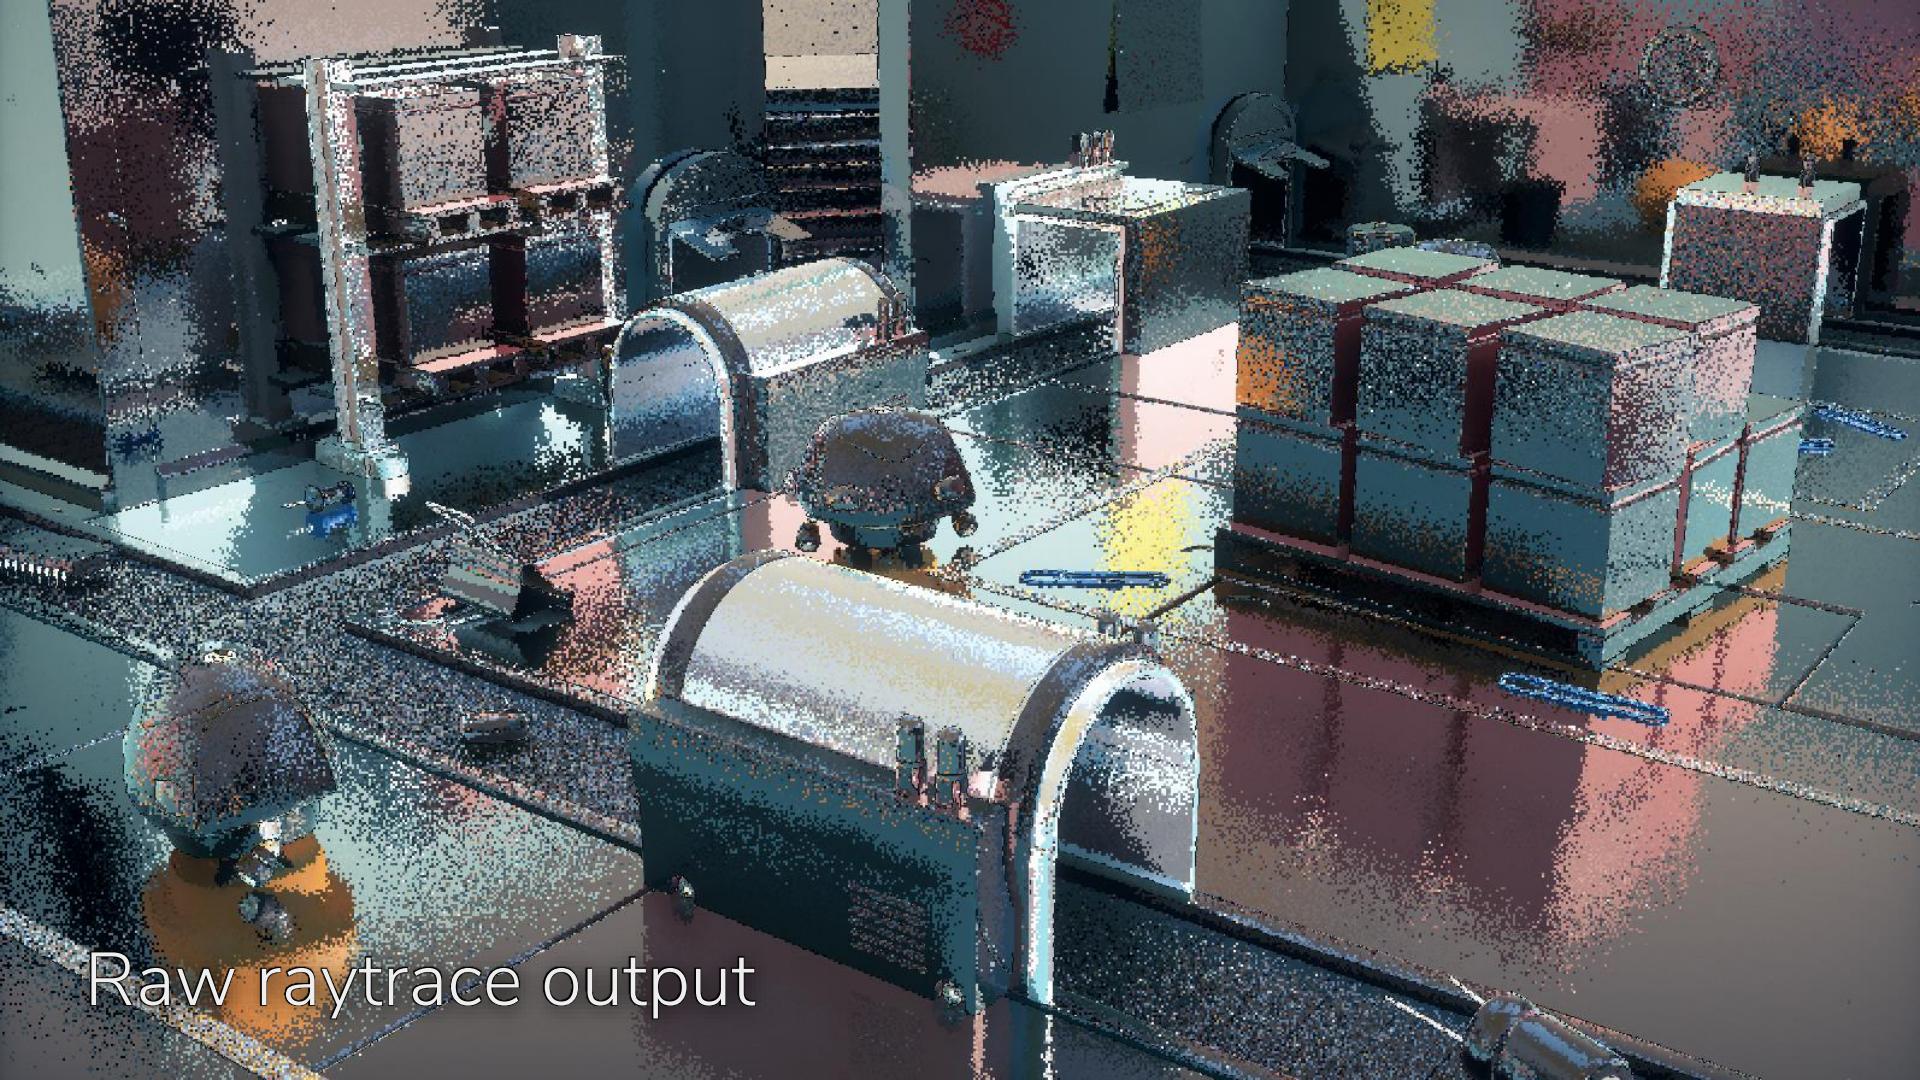


Multi-Layered Materials [Weidlich 2007]

# Reflection Raygen

- Read generated ray
- TraceRay()
  - Same hit shaders as path-tracing reference
- Output packed in 2x RGBA 16F

Color R	Color G	Color B	G-buffer Depth
Dir X * Hit T	Dir Y * Hit T	Dir Z * Hit T	Inverse PDF



Raw raytrace output

# Why Stochastic?

- Gives right answer, but
  - Produces noise
  - Thrashes caches
- Alternative: post-filter
  - Needs arbitrarily large kernel
  - Introduces bleeding
  - Sensitive to aliasing
- Meet in the middle?
  - Bias samplers
  - Combine with spatial filtering
  - High variance -> increase bias
  - More research needed!

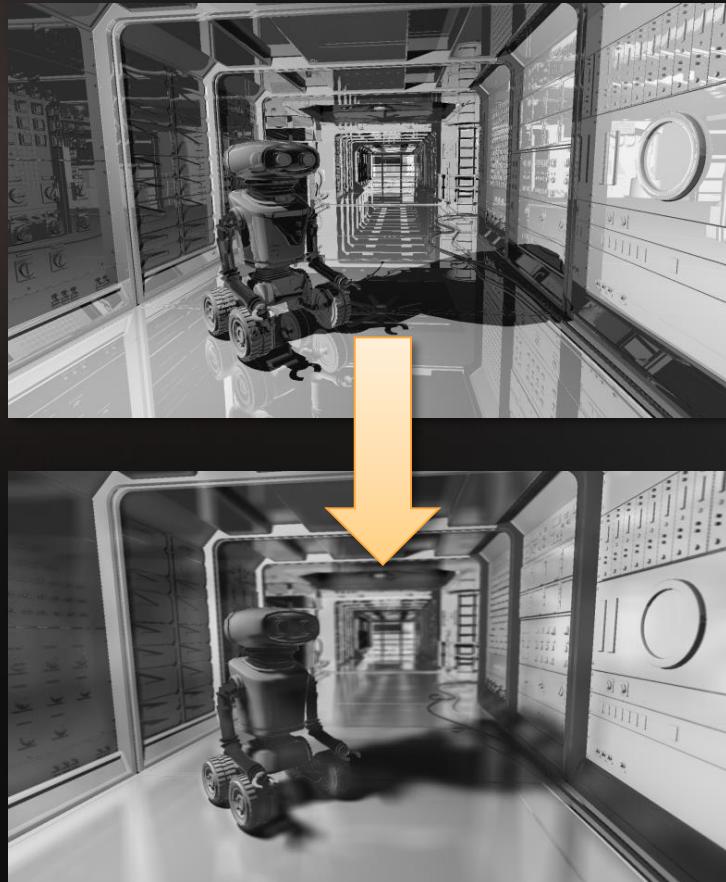
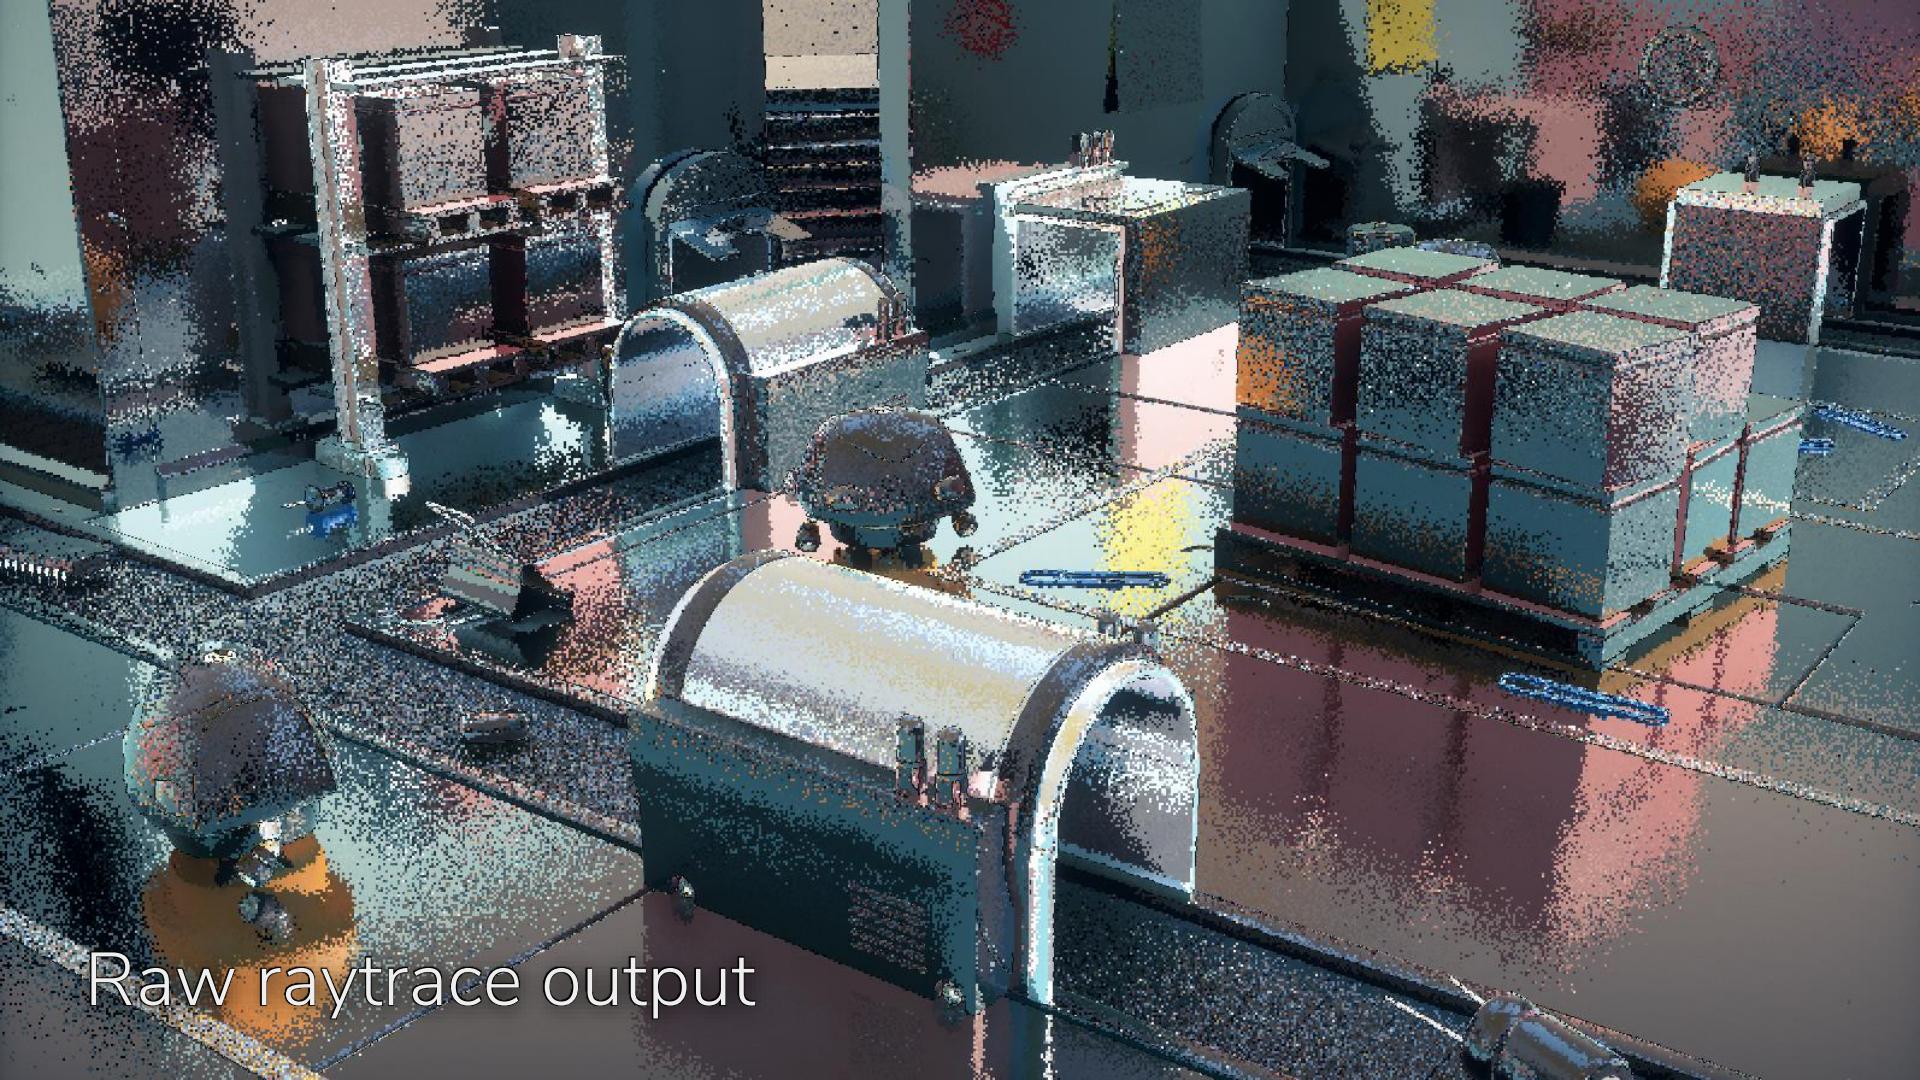


Image Space Gathering [Robison 2009]



Raw raytrace output



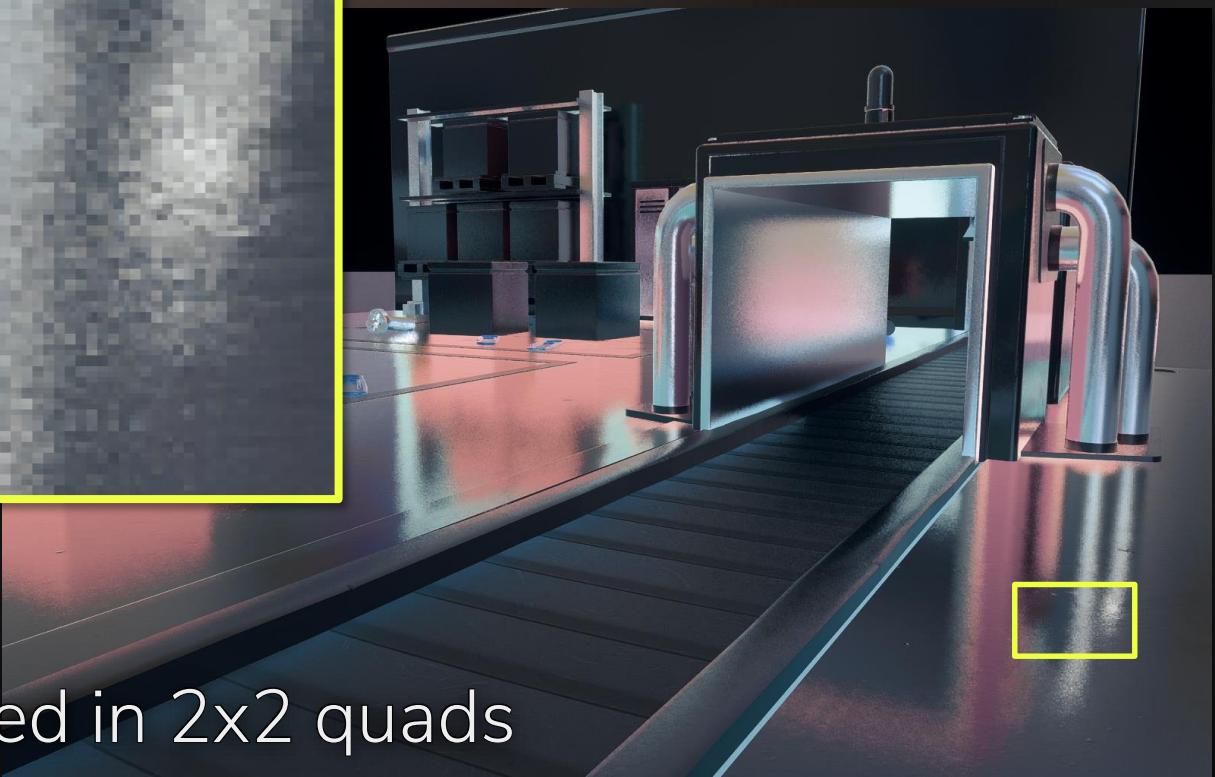
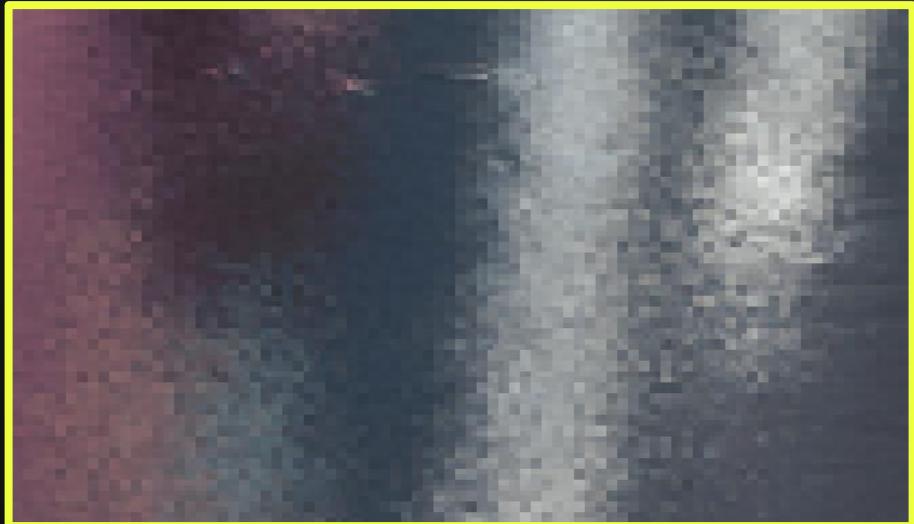
+Spatial reconstruction

- Based on Stochastic Screen-Space Reflections [Stachowiak 2015]
- For every full-res pixel, use 16 half-res ray hits
  - Poisson-disk distribution around pixel
  - Scale by local BRDF
  - Divide by ray PDF
- Ratio estimator [Heitz 2018]

```
result      = 0.0
weightSum  = 0.0
for pixel in neighborhood:
    weight = localBrdf(pixel.hit) / pixel.hitPdf
    result += color(pixel.hit) * weight
    weightSum += weight
result /= weightSum
```

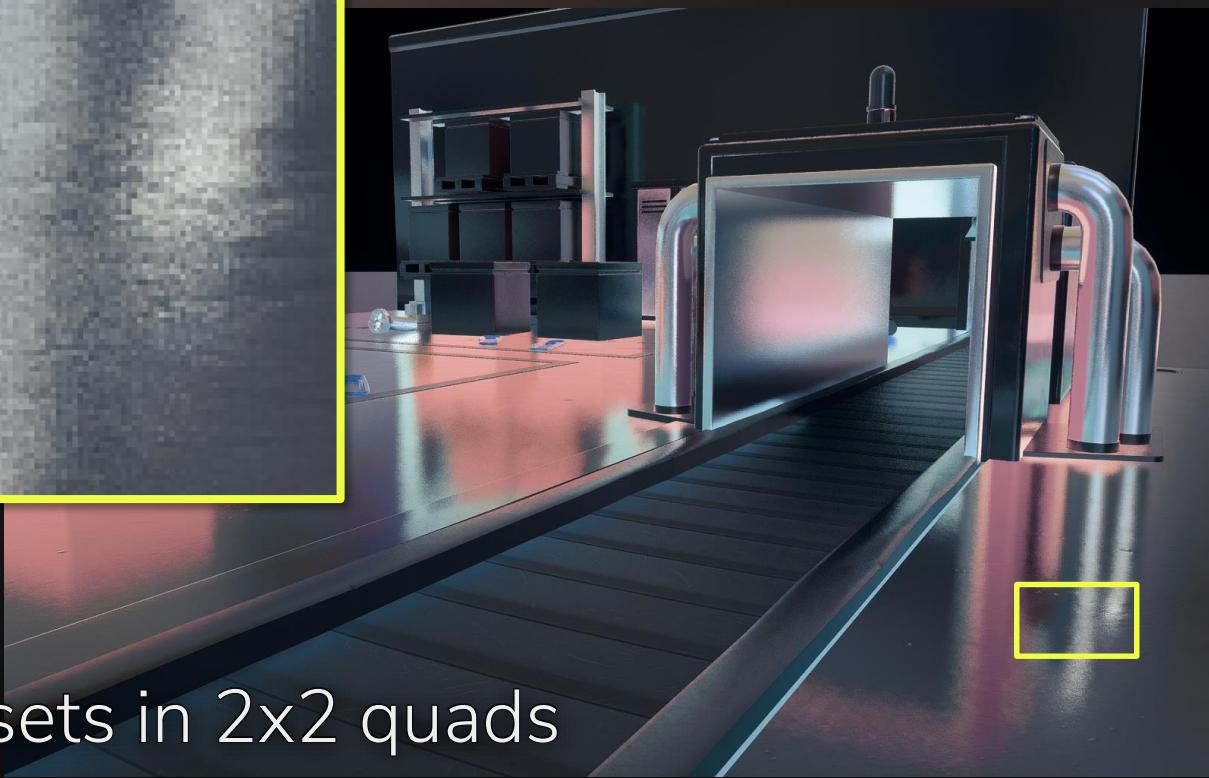
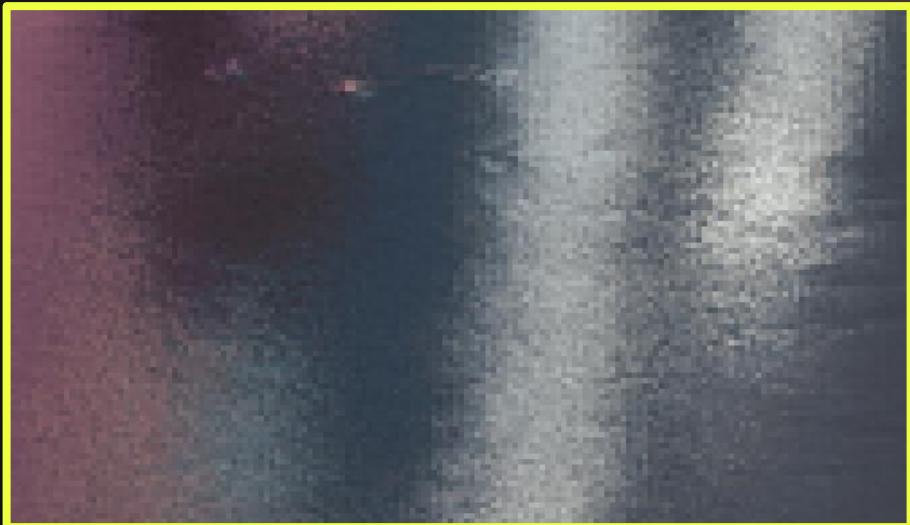
+Spatial reconstruction

# Spatial upsampling



**Same** ray hits used in 2x2 quads

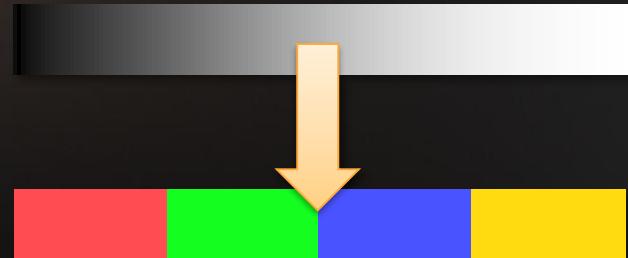
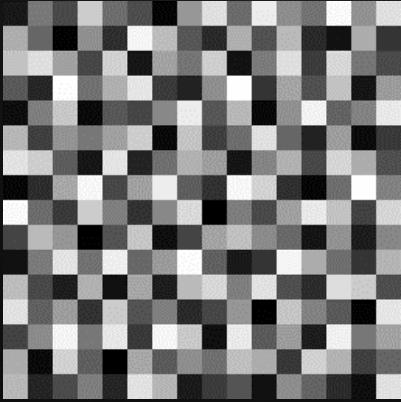
# Spatial upsampling



**Disjoint** sample sets in  $2 \times 2$  quads

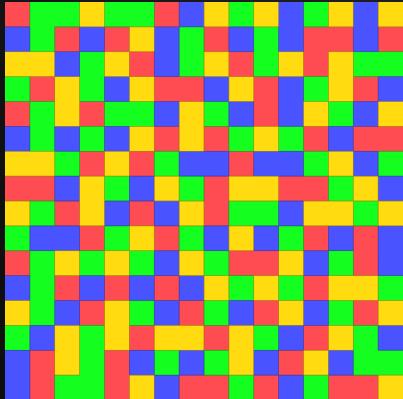
# Spatial Sample Pattern

- Four sample sets for pixels in 2x2 quad
- Threshold blue noise into four classes



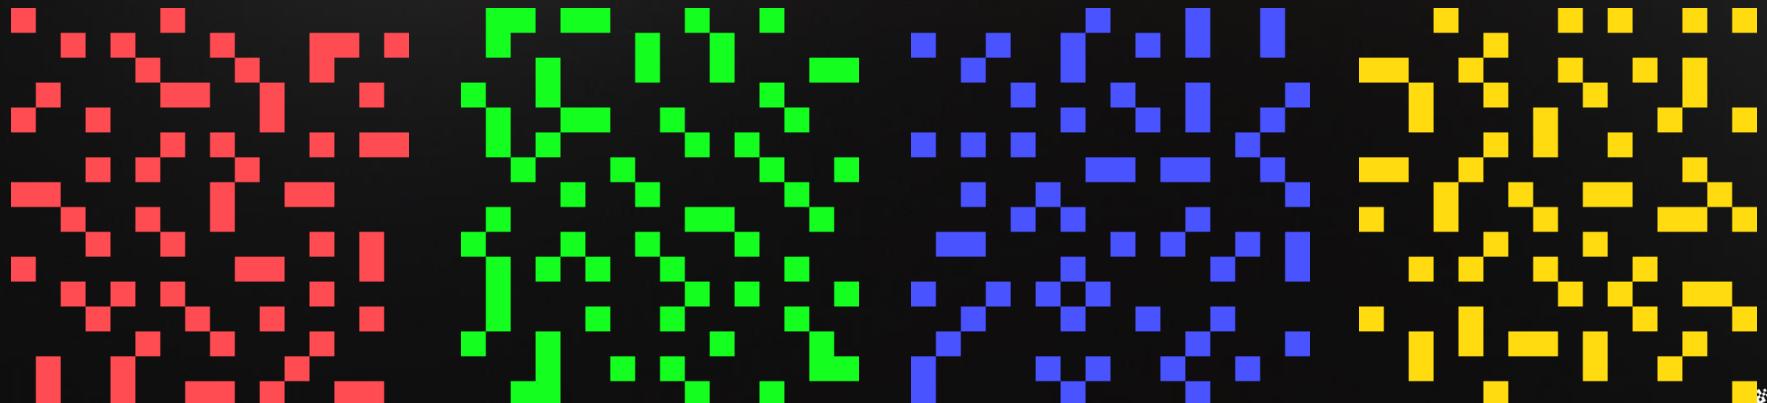
# Spatial Sample Pattern

- Four sample sets for pixels in 2x2 quad
- Threshold blue noise into four classes



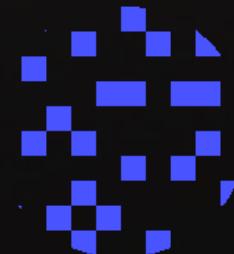
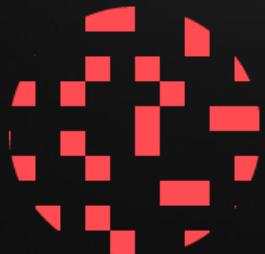
# Spatial Sample Pattern

- Four sample sets for pixels in 2x2 quad
- Threshold blue noise into four classes
- Class determines sample set



# Spatial Sample Pattern

- Four sample sets for pixels in 2x2 quad
- Threshold blue noise into four classes
- Class determines sample set
- Take 16 samples from each set in a disk
  - Disjoint high-quality samples





Spatial reconstruction



+Temporal accumulation



+Bilateral cleanup

- Secondary bilateral filter
  - Much dumber than reconstruction
  - Introduces blur
  - Only run where variance high
    - Variance from spatial reconstruction
    - Temporally smoothed with hysteresis = 0.5
  - Variable kernel width, sample count

+Bilateral cleanup

- Secondary bilateral filter
  - Much dumber than reconstruction
  - Introduces blur
  - Only run where variance high
    - Variance from spatial reconstruction
    - Temporally smoothed with hysteresis = 0.5
- Variable kernel width, sample count

Variance estimate



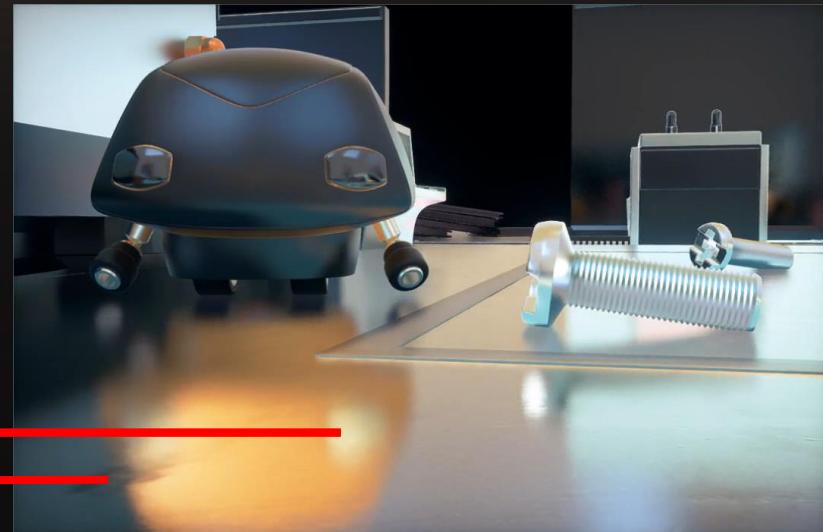
+Bilateral cleanup



+Temporal anti-aliasing

# Temporal Reprojection

- Two sources of velocity
  - Per-pixel motion vector
  - Reprojection of hit point

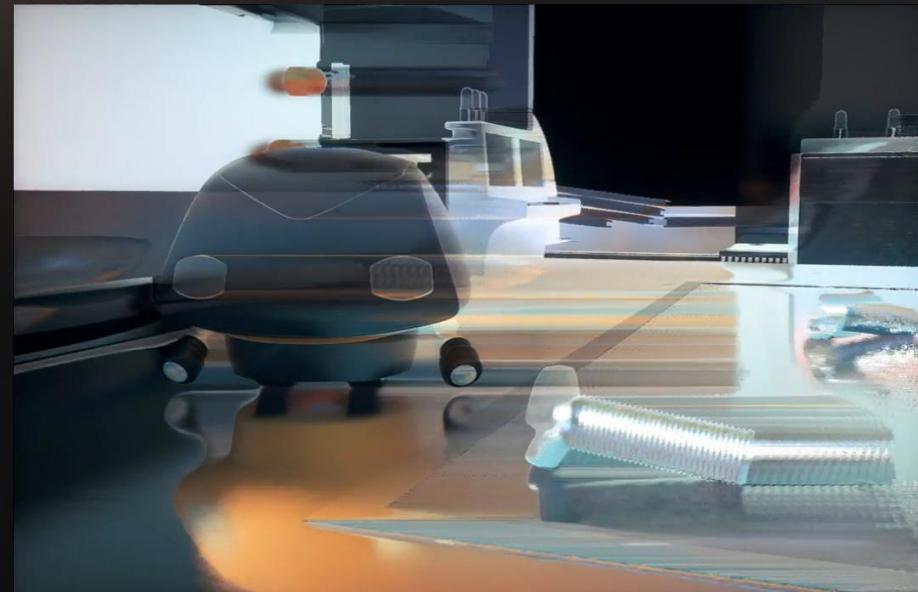


# Temporal Reprojection

- Neither reprojection method robust



Per-pixel motion vector



Reprojection of hit point

# Dual-Source Reprojection

- Sample history using both reprojection strategies
- Estimate local statistics during spatial reconstruction
  - RGB mean
  - RGB standard deviation
  - Of used hit samples
- Weigh contributions
  - $\text{rgb\_dist} = (\text{rgb} - \text{rgb\_mean}) / \text{rgb\_dev}$
  - $\text{dist} = (\text{rgb} - \text{rgb\_mean}) / \text{rgb\_dev}$
  - $w = \exp2(-10 * \text{luma}(\text{dist}))$



# Reprojection Clamp

- Build color box from local statistics
  - $\text{rgb\_min} = \text{rgb\_mean} - \text{rgb\_dev}$
  - $\text{rgb\_max} = \text{rgb\_mean} + \text{rgb\_dev}$
- Clamp reprojected values to box
  - $\text{rgb} = \text{clamp}(\text{rgb}, \text{rgb\_min}, \text{rgb\_max})$
  - Same idea as in temporal anti-aliasing [Karis 2014]
- Weigh clamped contributions
  - $\text{rgb\_sum} += \text{val} * \text{w}$
- Introduces bias
- Fixes most ghosting



# Transparency & Translucency

- Raytracing enables accurate light scattering
- Refractions
  - Order-independent (OIT)
  - Variable roughness
  - Handles multiple IOR transitions
  - Beer-Lambert absorption
- Translucency
  - Light scattering inside a medium
  - Inspired by Translucency in Frostbite [Barré-Brisebois 2011]



Glass and Translucency

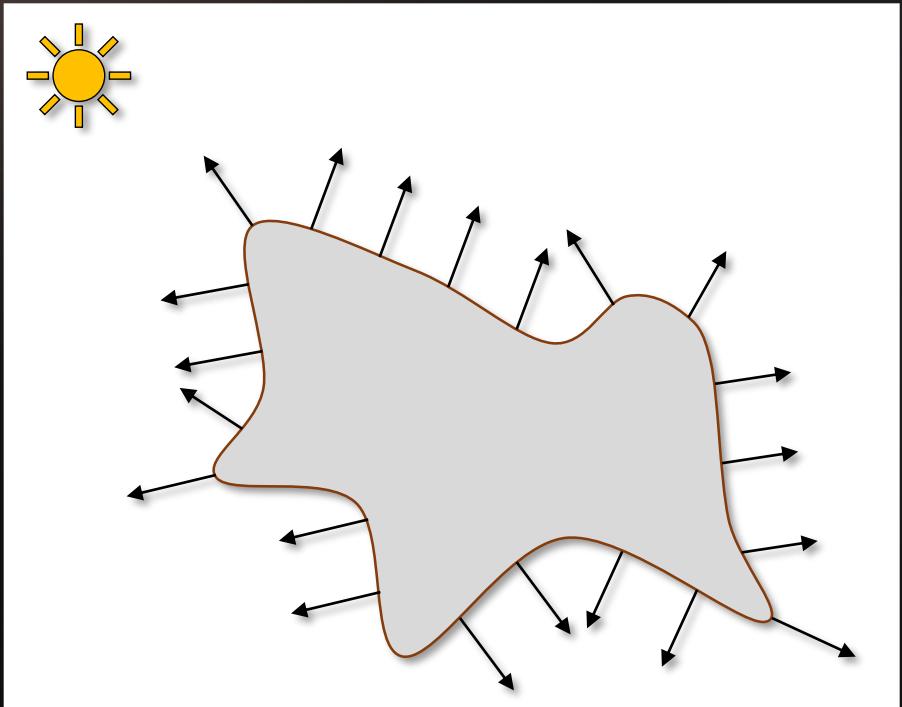
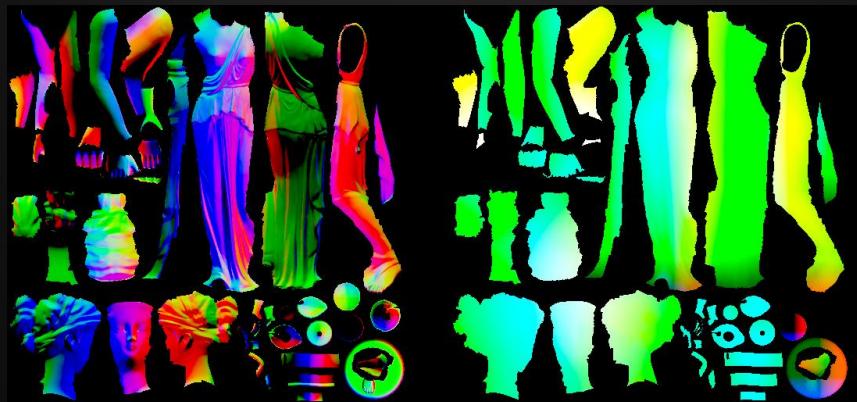
# Integration

- Multiple samples required for rough refractions and translucency
- Multi-layer screen-filtering complicated
  - Layered framebuffers?
  - Per-pixel linked lists?
  - Filter in 2D or between layers?
  - More research needed!
- Accumulate in texture-space instead
  - Stable integration domain
- Static texture allocation
  - 512\*512 per object
  - Time-sliced, ~1M rays/frame



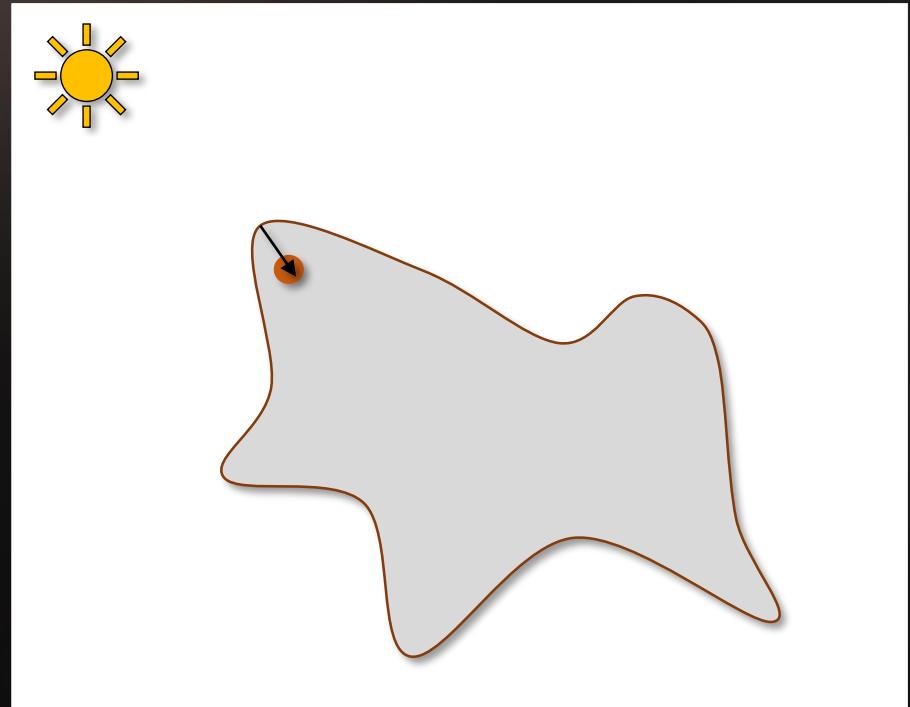
# Translucency Breakdown

- For every valid position & normal



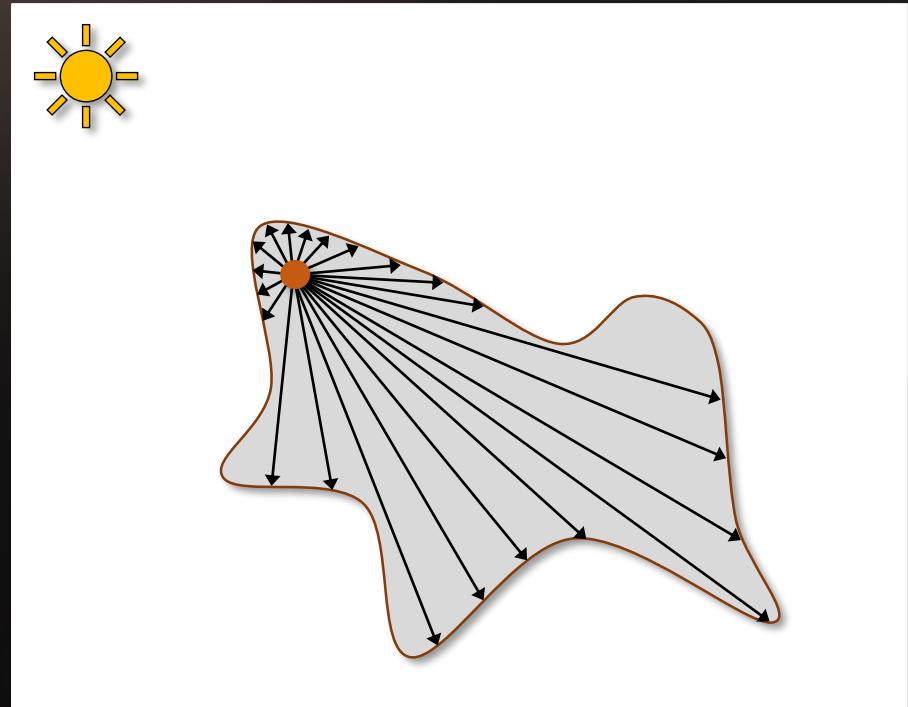
# Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside



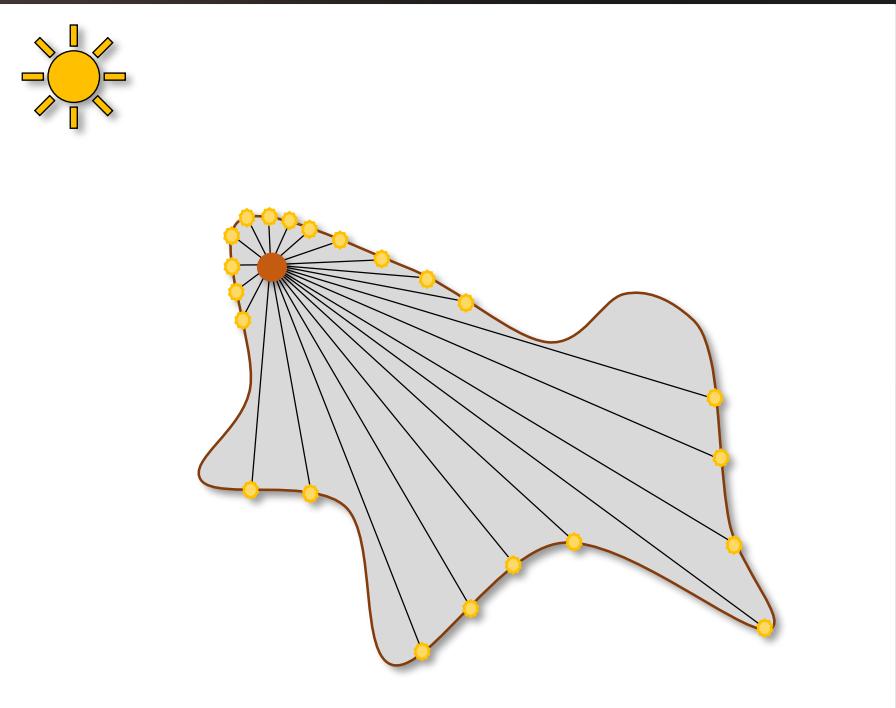
# Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
  - *(Importance-sample phase function)*



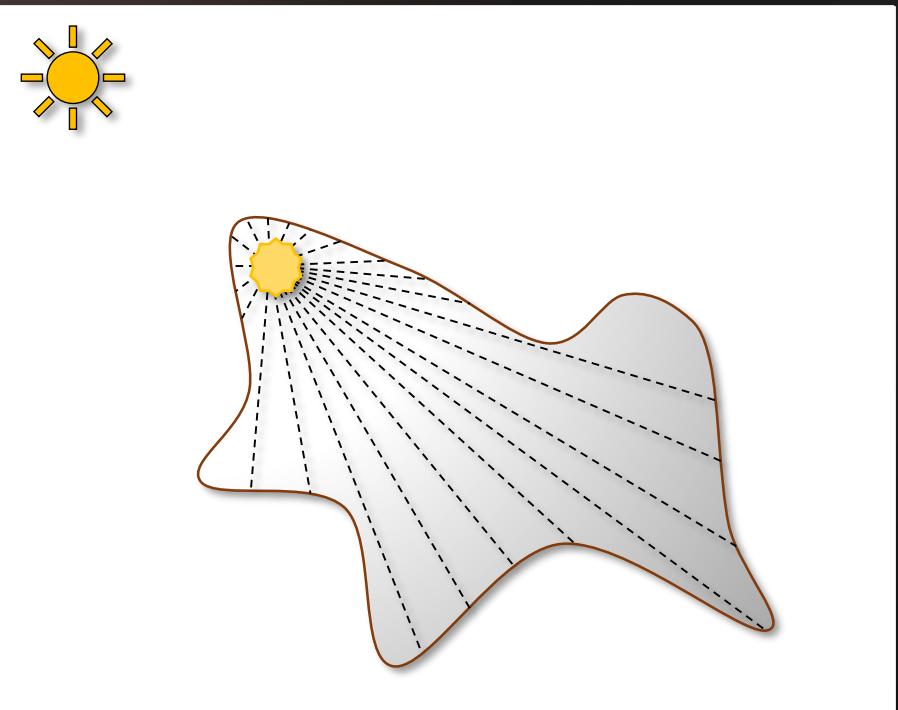
# Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
  - (*Importance-sample phase function*)
- Compute lighting at intersection



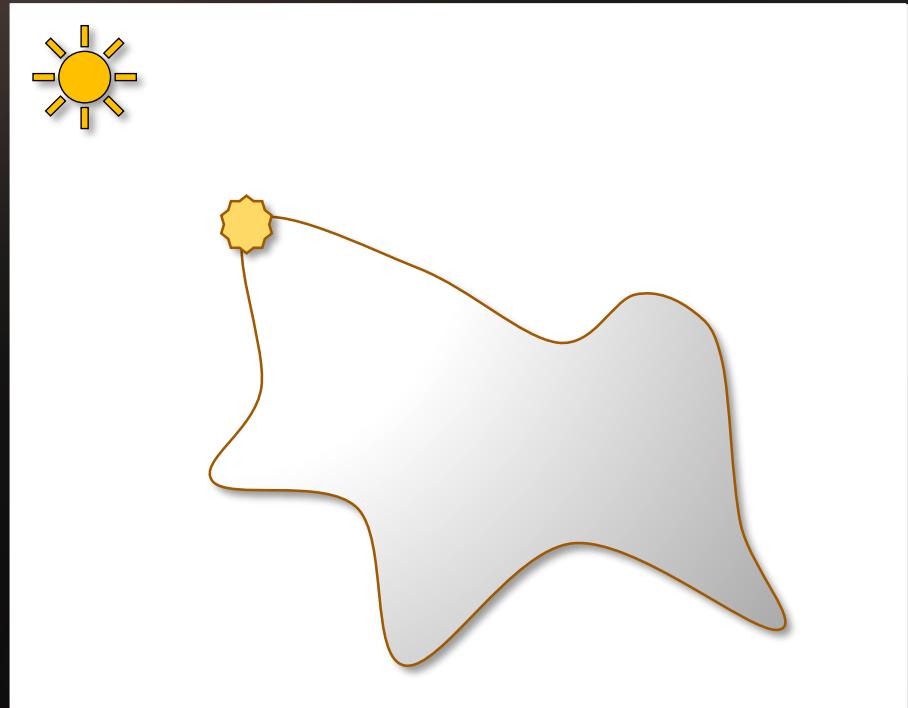
# Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
  - (*Importance-sample phase function*)
- Compute lighting at intersection
- Gather all samples



# Translucency Breakdown

- For every valid position & normal
- Flip normal and push (ray) inside
- Launch rays in uniform sphere dist.
  - (*Importance-sample phase function*)
- Compute lighting at intersection
- Gather all samples
- Update value in texture



# Translucency Filtering

- Can denoise spatially and/or temporally
- Temporal: build an update heuristic
  - Reactive enough for moving lights & objects
  - Exponential moving average can be OK
- Variance-adaptive mean estimation
  - Based on exponential moving average
  - Adaptive hysteresis



# Transparency

- Similar approach to translucency
- Launch ray using view's origin and direction
- Refract based on medium's index of refraction
  - Sample a BSDF for rough refraction
- Trace a ray in the scene & sample lighting
- Tint the result
  - Chromatic aberration from interface
  - Beer-Lambert absorption in medium



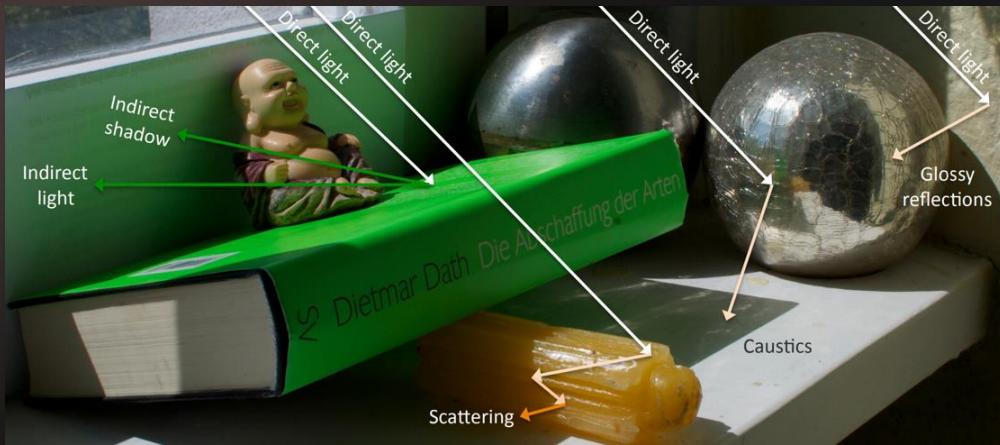
# Transparency

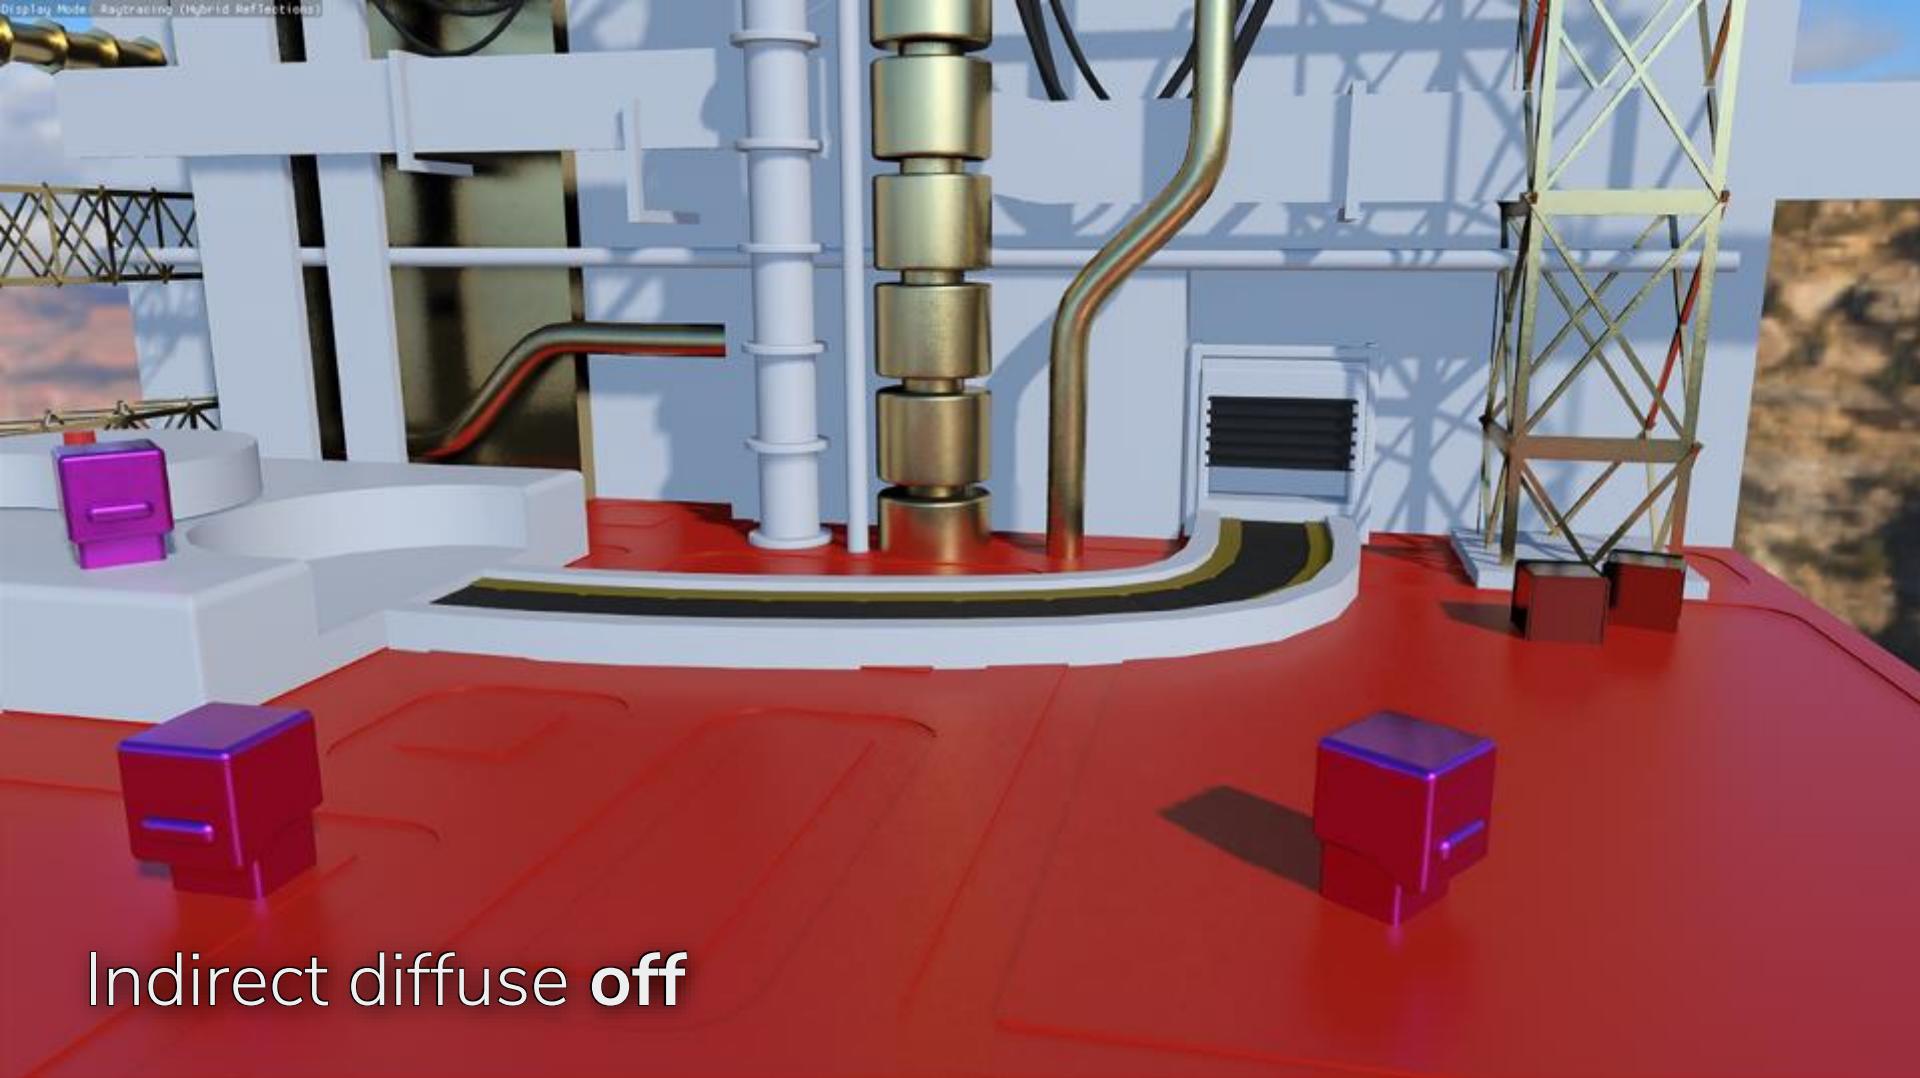
- Similar approach to translucency
- Launch ray using view's origin and direction
- Refract based on medium's index of refraction
  - Sample a BSDF for rough refraction
- Trace a ray in the scene & sample lighting
- Tint the result
  - Chromatic aberration from interface
  - Beer-Lambert absorption in medium
- Pen: we don't handle transparent shadows yet



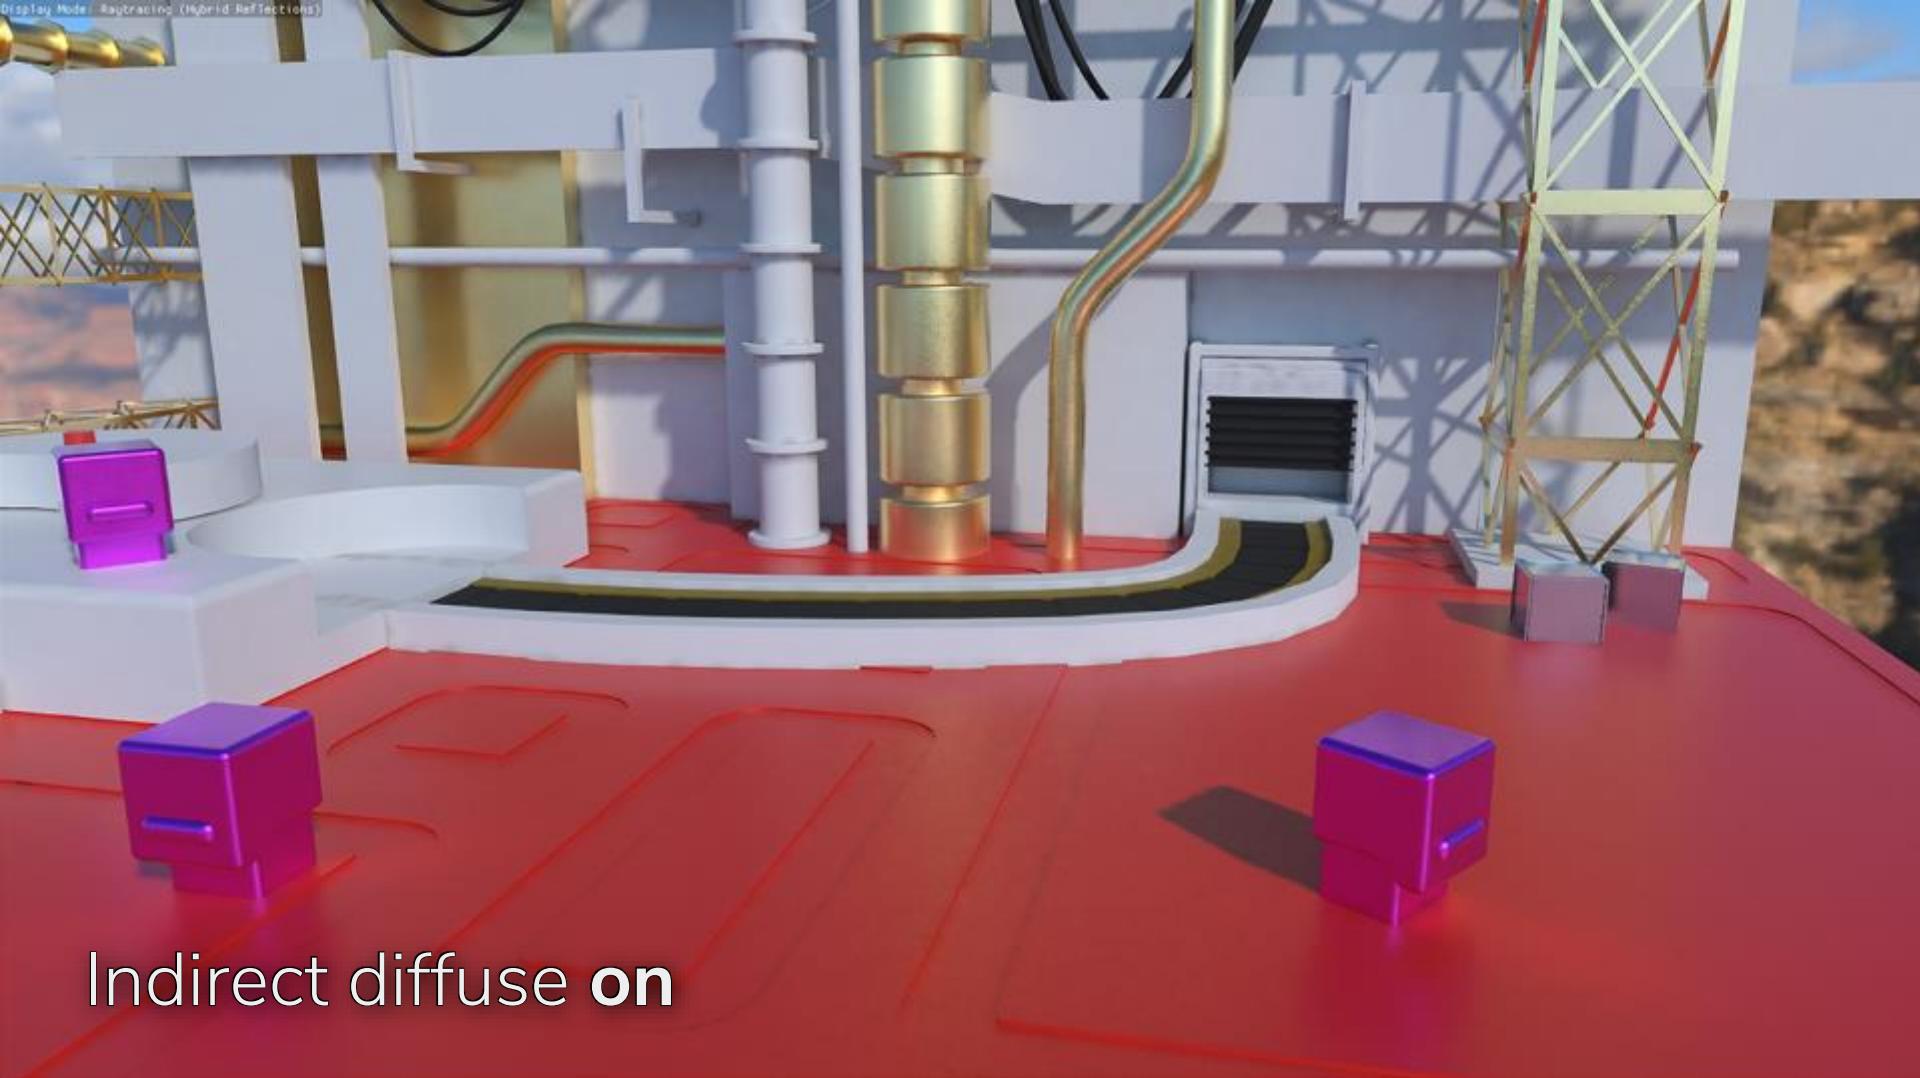
# Global Illumination

- Important for consistency
- Minimize artist overhead
- We want a technique with:
  - No precomputation
  - No parametrization (UVs, proxies)
  - Support for static and dynamic scenes
  - Adaptive refinement





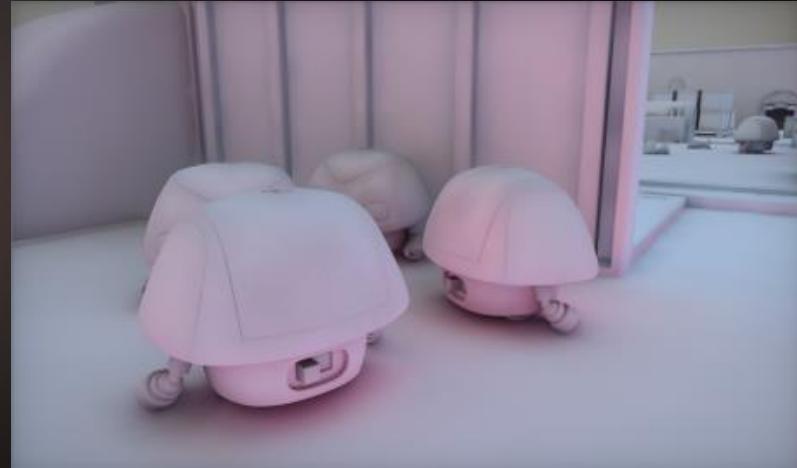
Indirect diffuse **off**



Indirect diffuse **on**

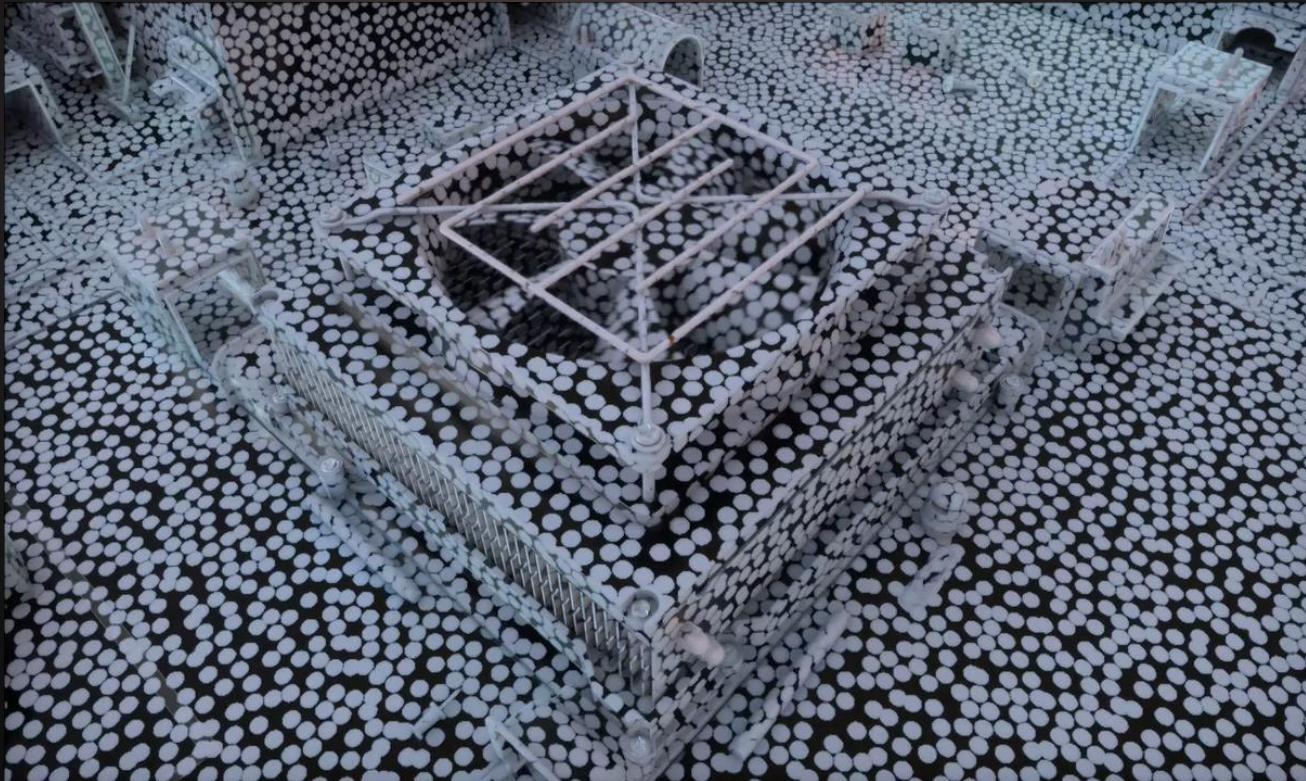
# Spatial Storage

- Can't afford to solve every frame
  - Our GI budget: 250k rays / frame
- World space vs screen space
  - World space is stable
  - Screen-space fully dynamic
- World space surfels
  - Easy to accumulate
  - Distribute on the fly
  - No parameterization
  - Smooth result by construction
  - Position, normal, radius, animation info



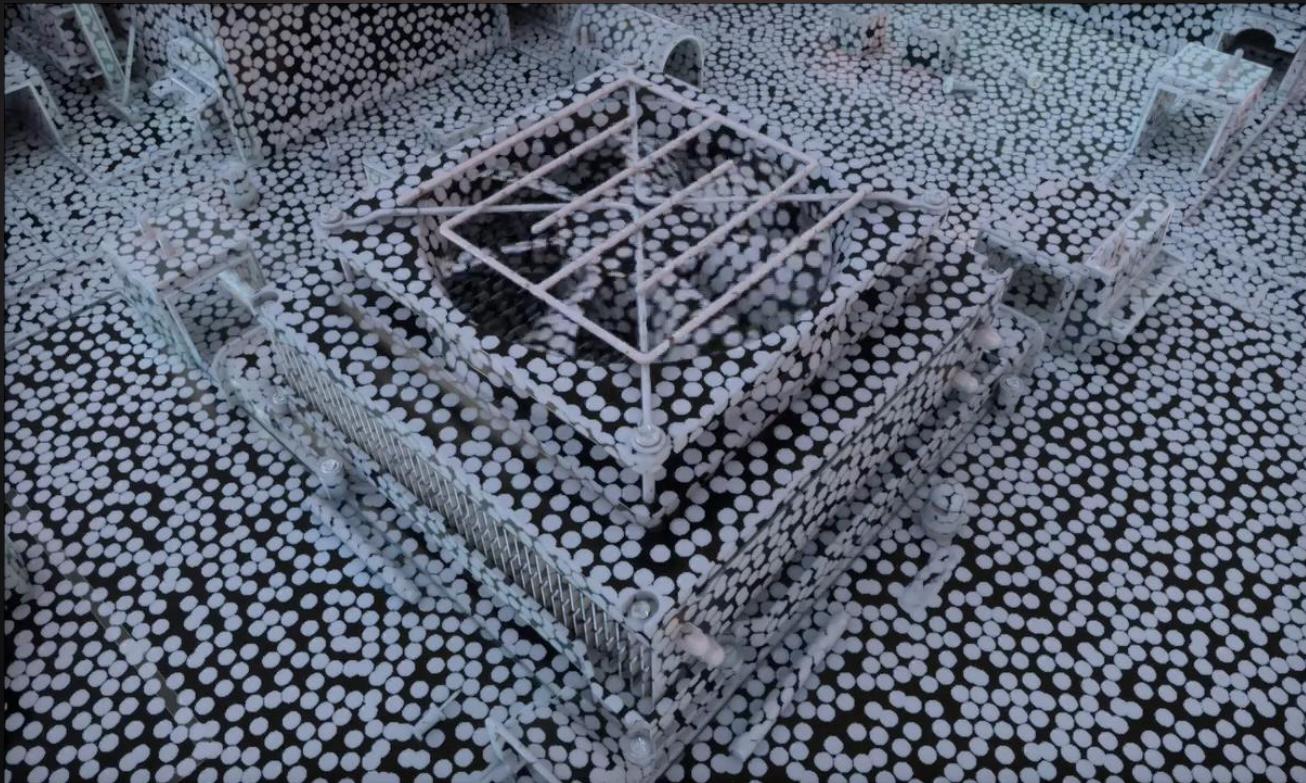
S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

# Surfel Skinning



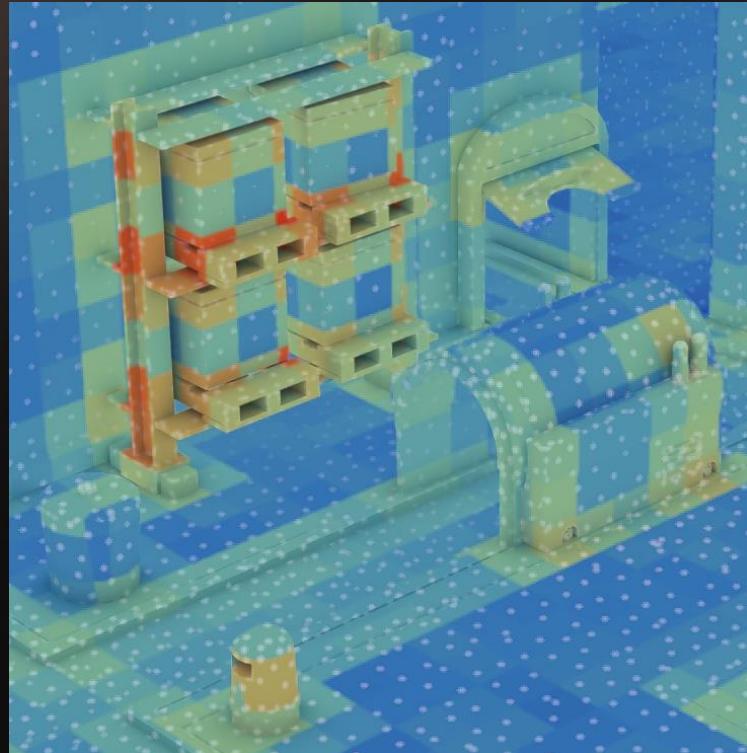
S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

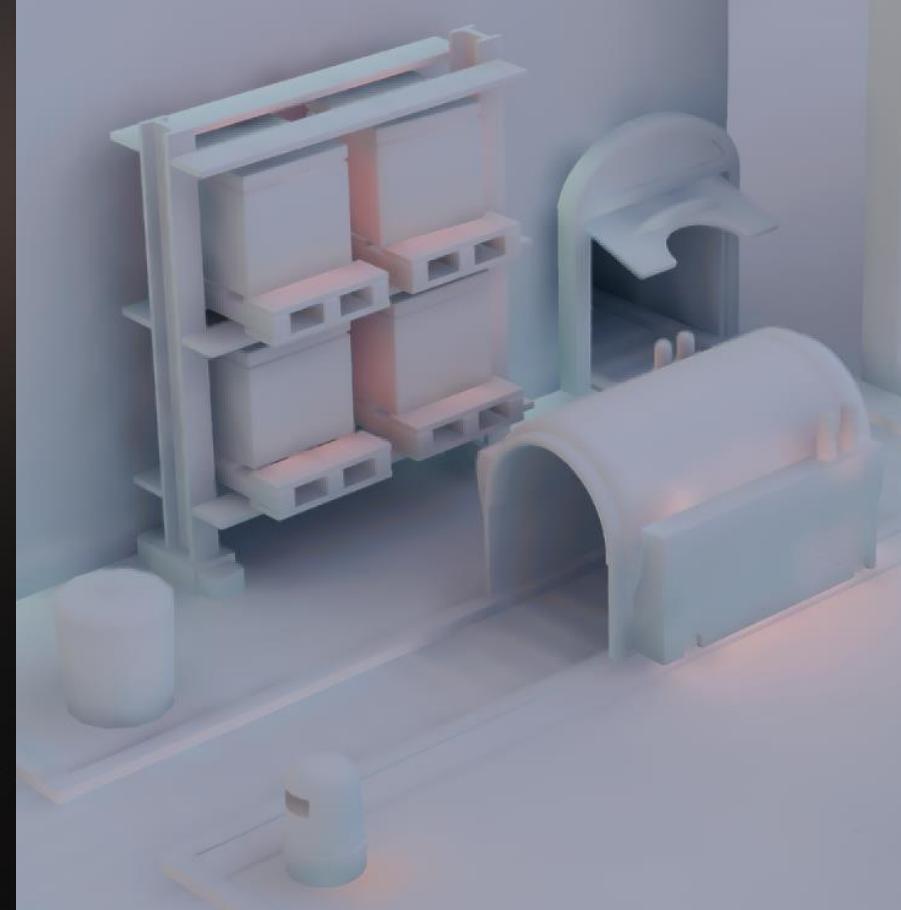
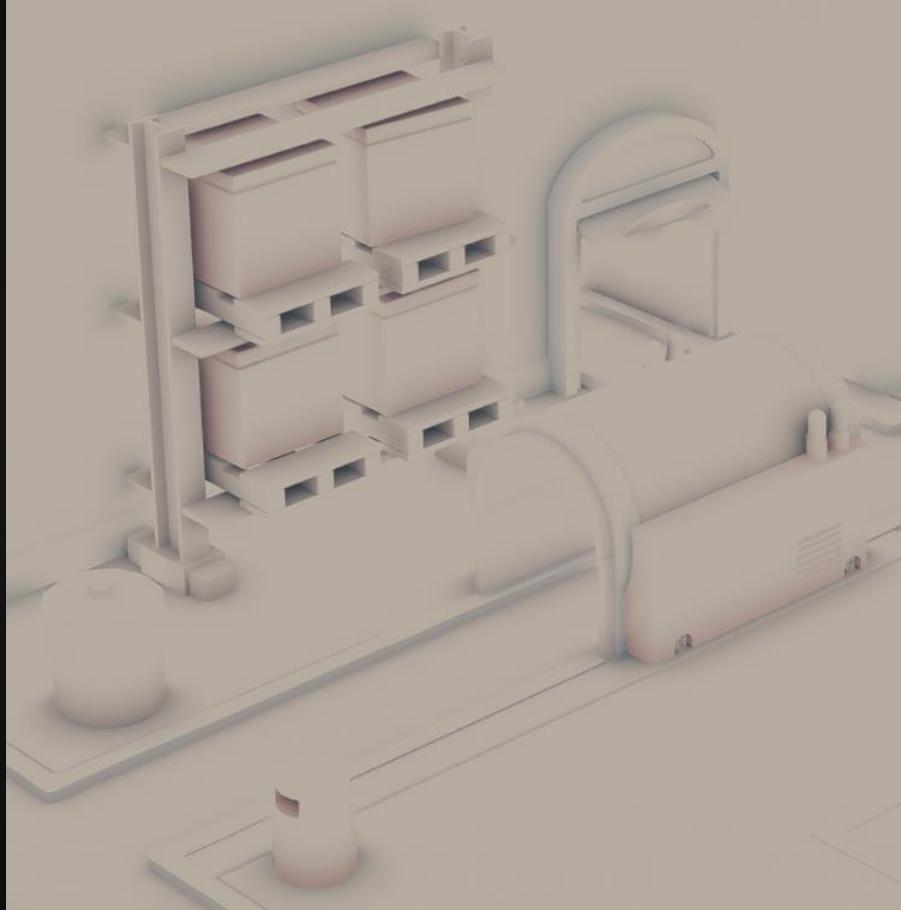
# Surfel Skinning



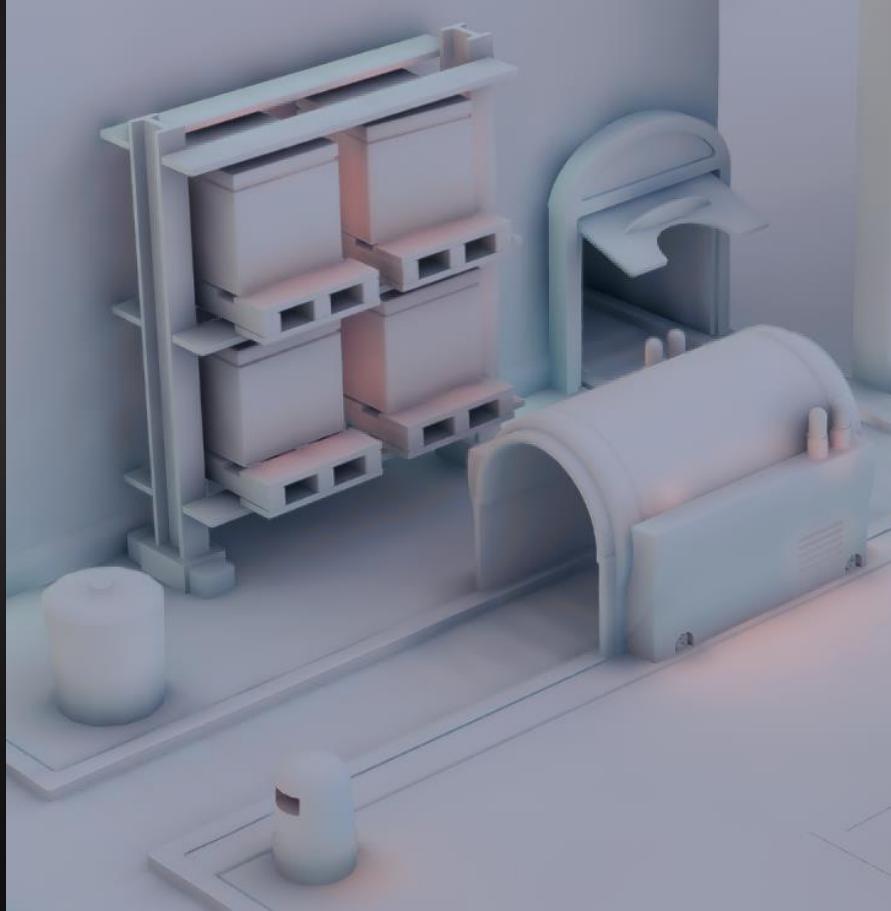
# Surfel Screen Application

- Render like deferred light sources
  - Diffuse only
  - Smoothstep distance attenuation
    - Mahalanobis metric
  - Squared dot of normals angular weight
  - Inspired by [Lehtinen 2008]
- World-space 3D culling
  - Uniform grid
  - Each cell holds list of surfels
  - Find one cell per pixel, use all from list





Combine with Screen Space AO [Jimenez 2016]



Combine with Screen Space AO [Jimenez 2016]

S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

# Surfel Placement



Surfel Spawning From Camera @ 1% speed

S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

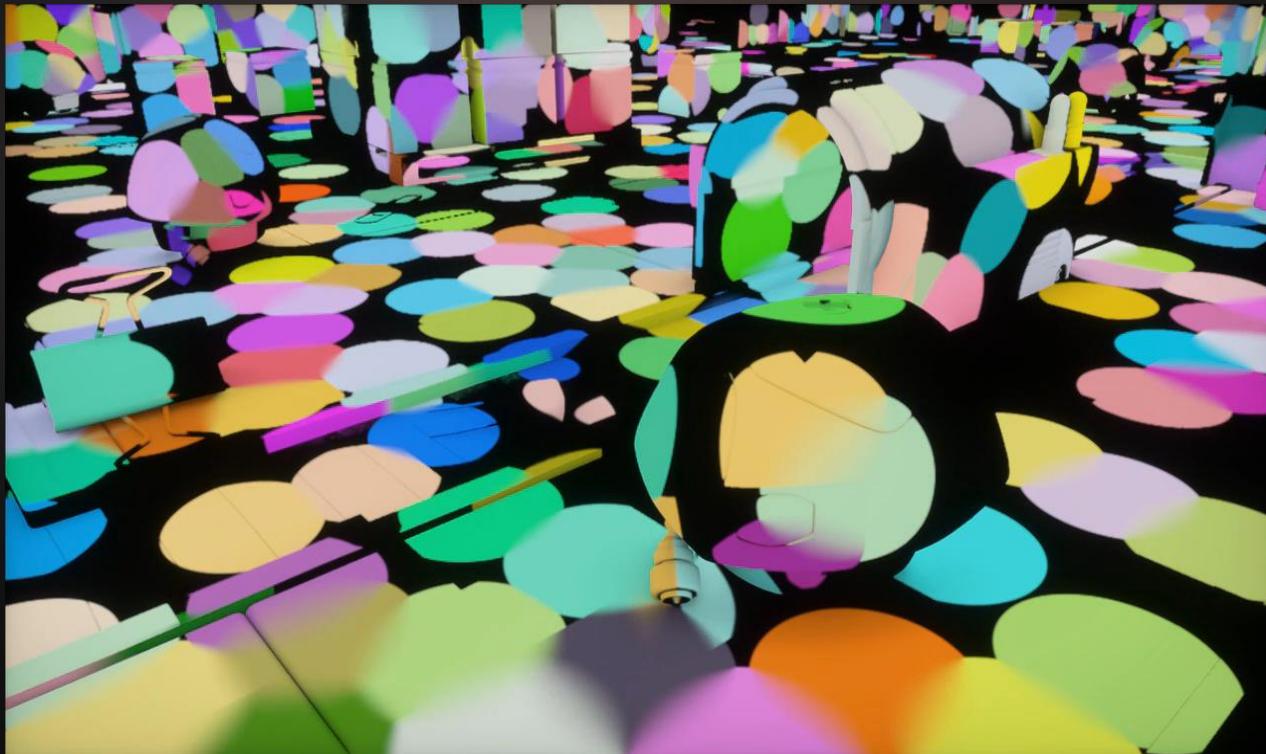
# Surfel Placement



Surfel Spawning From Camera @ 1% speed

S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

# Surfel Placement



Surfel Spawning From Camera @ 1% speed

S E E D // Stochastic all the things: Raytracing in hybrid real-time rendering

# Surfel Placement Algorithm

- Iterative screen-space hole filling

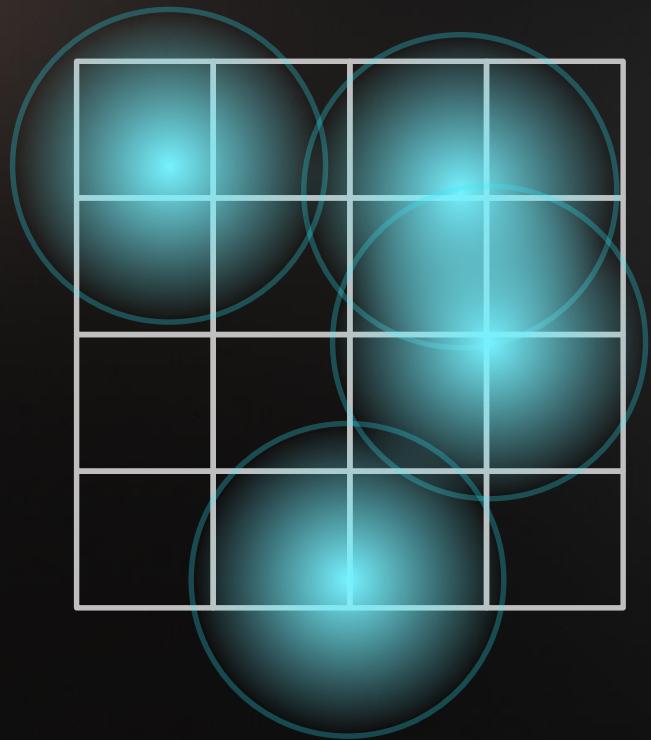
# Surfel Placement Algorithm

- Iterative screen-space hole filling
- Calculate surfel coverage for each pixel



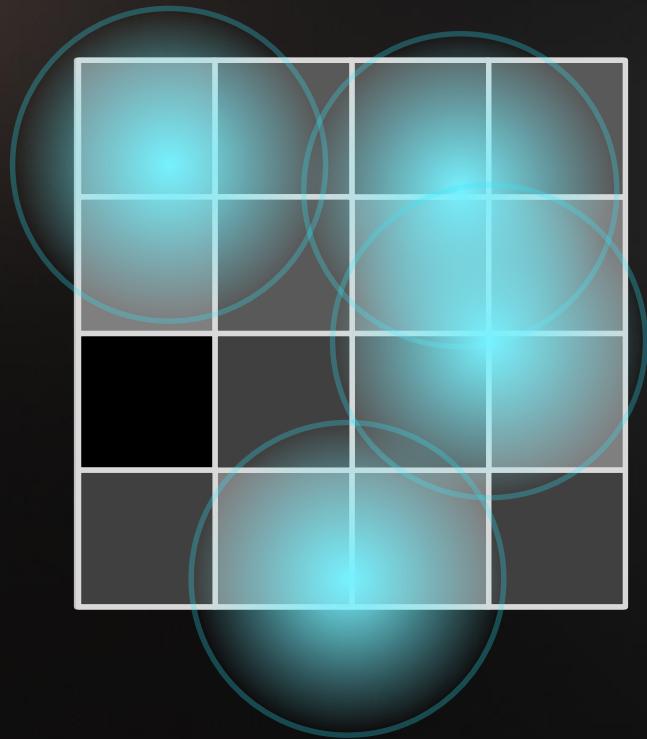
# Surfel Placement Algorithm

- Iterative screen-space hole filling
- Calculate surfel coverage for each pixel



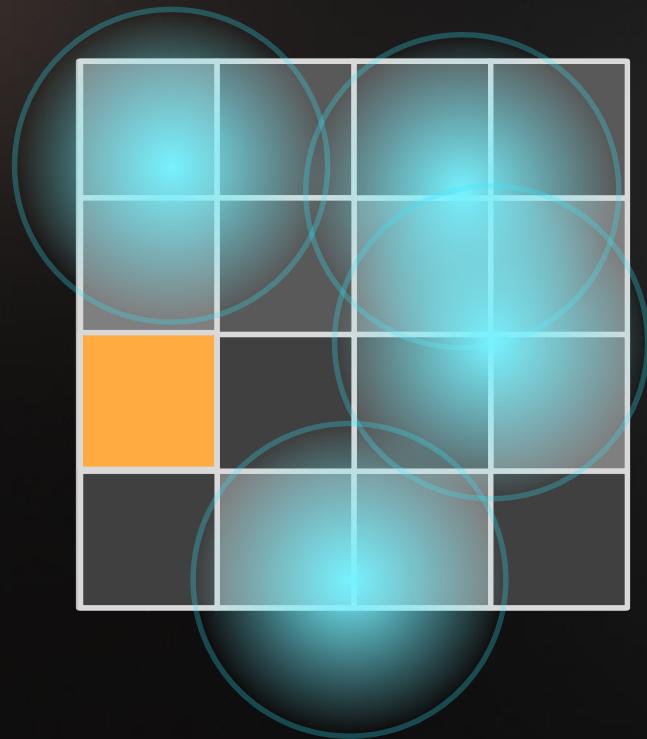
# Surfel Placement Algorithm

- Iterative screen-space hole filling
- Calculate surfel coverage for each pixel



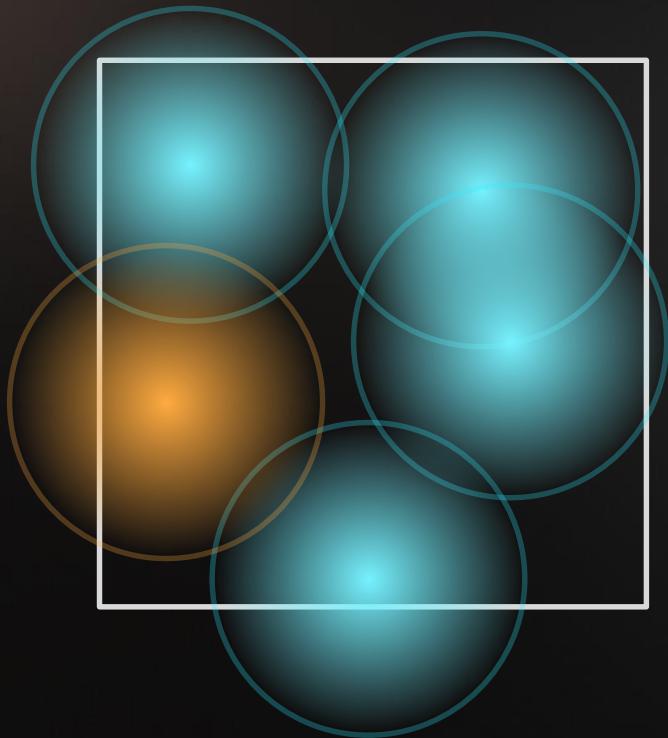
# Surfel Placement Algorithm

- Iterative screen-space hole filling
- Calculate surfel coverage for each pixel
- Find lowest coverage in tile (16x16 pixels)



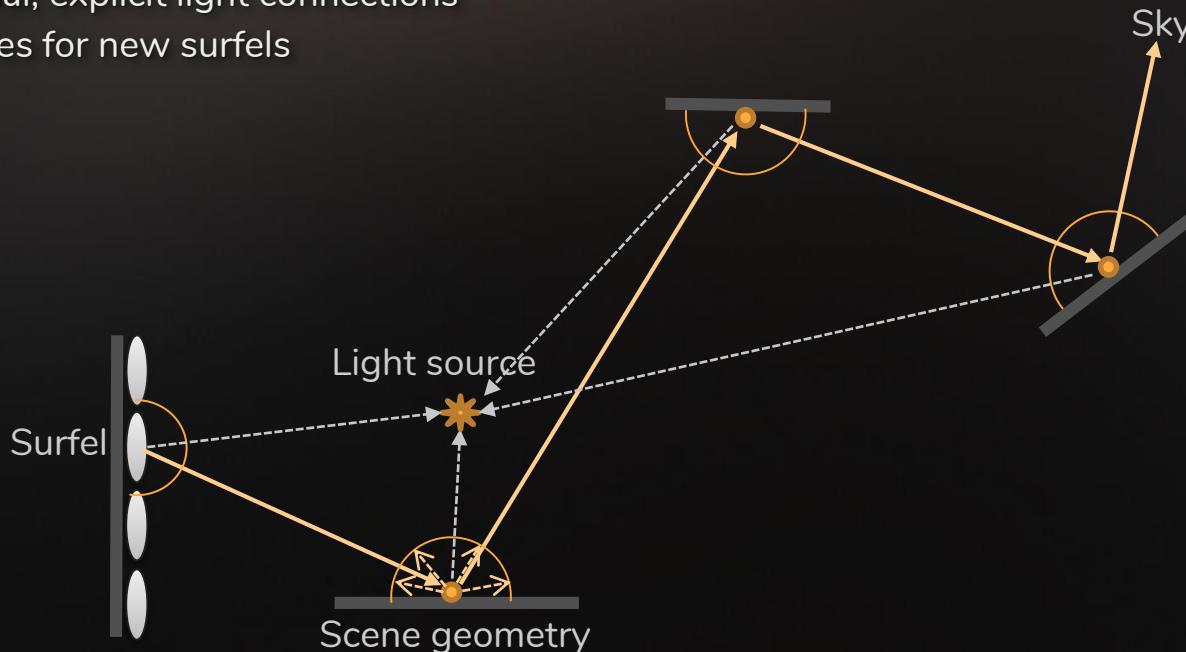
# Surfel Placement Algorithm

- Iterative screen-space hole filling
- Calculate surfel coverage for each pixel
- Find lowest coverage in tile (16x16 pixels)
- Probabilistically spawn surfel on pixel
  - Chance proportional to pixel's projected area
  - G-Buffer depth and normal
- Continue where coverage below threshold



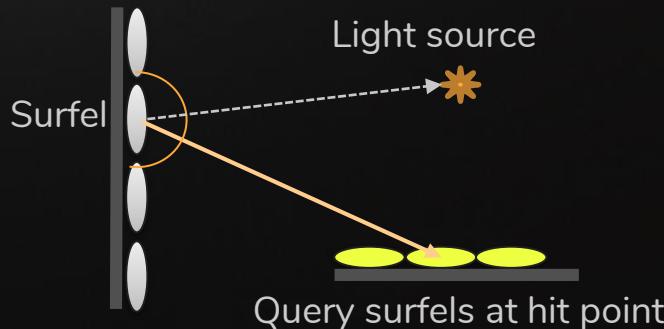
# Irradiance Calculation

- Path trace from surfels
  - Unidirectional, explicit light connections
  - More samples for new surfels



# Irradiance Calculation

- Path trace from surfels
  - Unidirectional, explicit light connections
  - More samples for new surfels
- Limit light path length
  - Sample previous GI at last bounce
- "Infinite" bounces amortized over frames



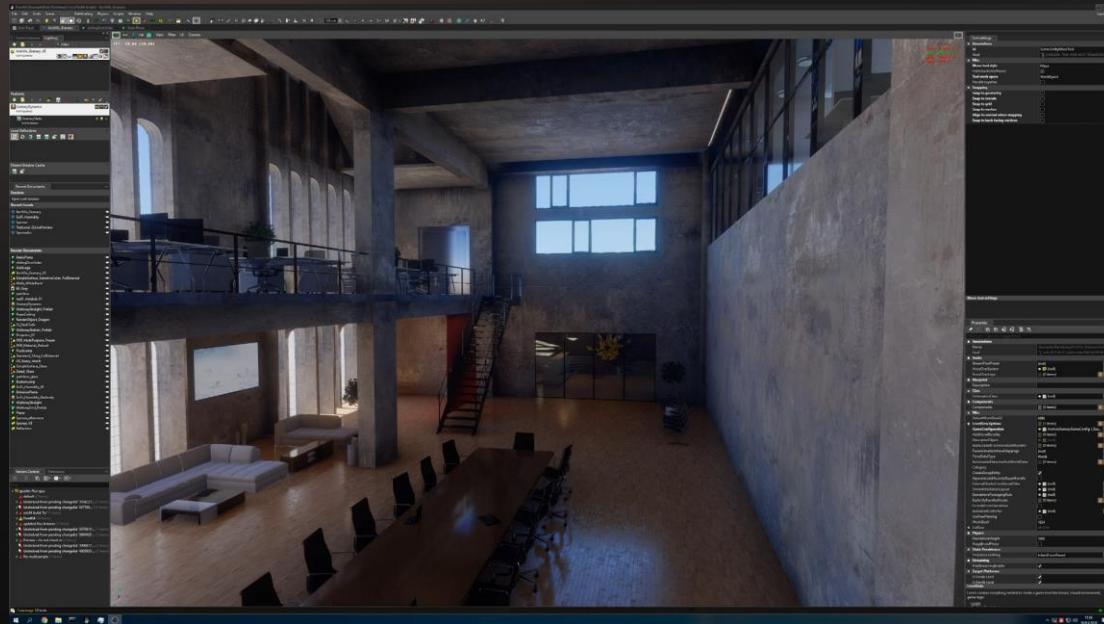
# Monte Carlo Integration

- Monte Carlo assumes immutable integrand
- Plain old mean estimation

$$\bar{x}_0 = 0$$

$$\bar{x}_{n+1} = \text{lerp}\left(\bar{x}_n, x_{n+1}, \frac{1}{n+1}\right)$$

- Can't use in dynamic GI
- Hard reset undesirable



Interactive GI in Frostbite [Hillaire 2018]

# Adaptive Integration

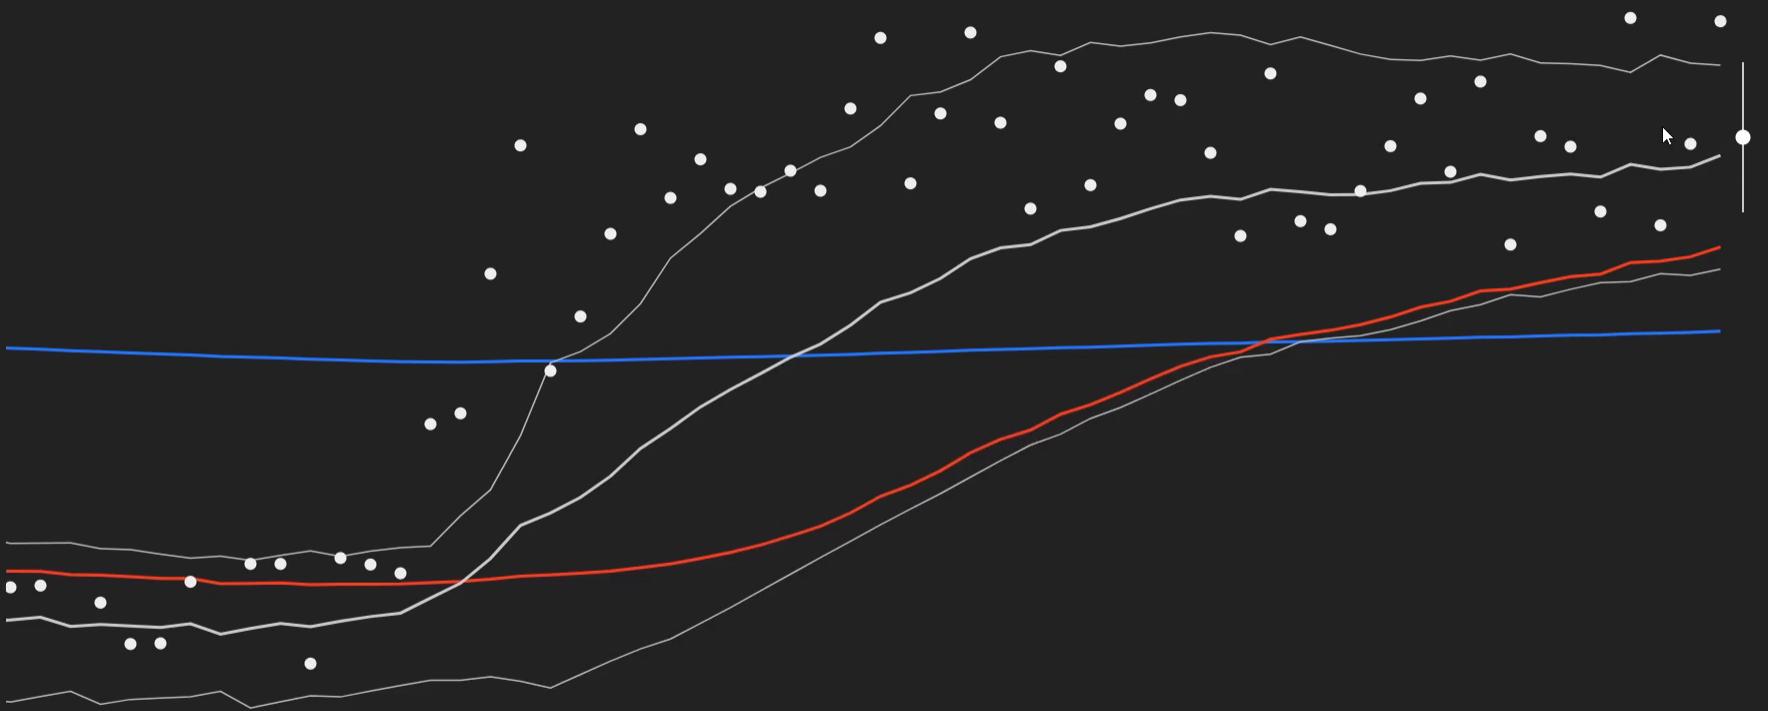
- Give up on fully converging
- Exponential moving average

$$\bar{x}_0 = 0$$

$$\bar{x}_{n+1} = \text{lerp}(\bar{x}_n, x_{n+1}, k)$$

- Estimate short-term statistics
  - Mean, variance
- Adapt blend factor  $k$ 
  - Reduce if variance high
  - Increase if mean far from short-term statistics





Adaptive integration in 1D

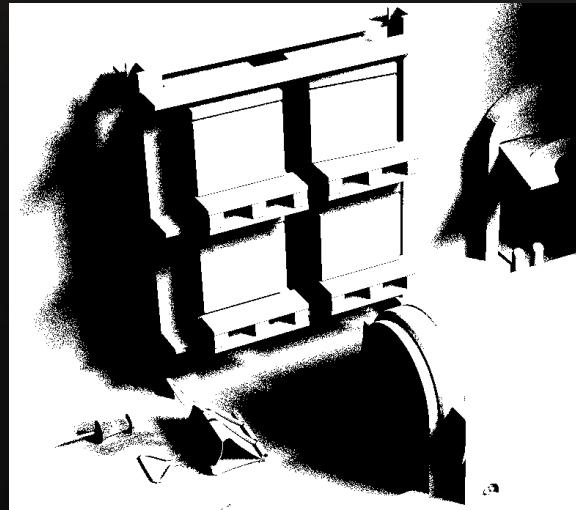
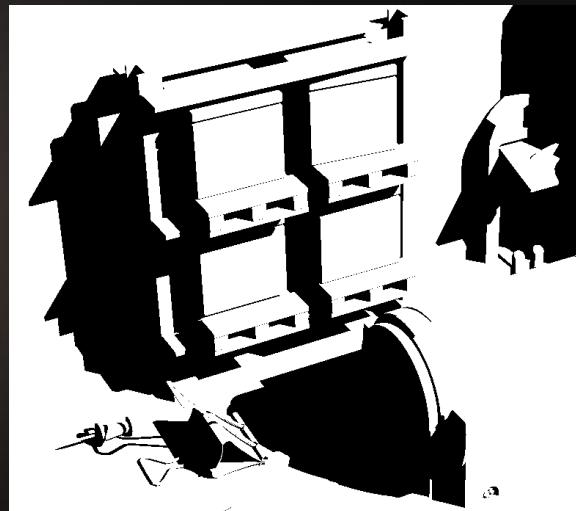
Adaptive mean

Exponential moving average

Short-term  
distribution bounds

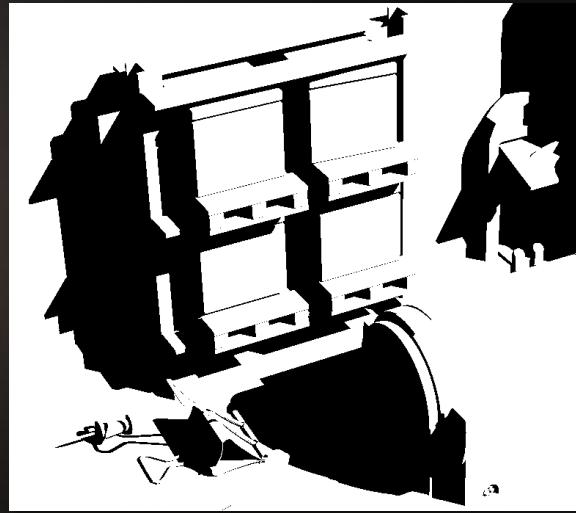
# Shadows

- Hard shadows trivial
  - Launch a ray towards light position
  - Check for any hit
- Soft shadows easy
  - Sample direction from uniform cone [PBRT]
  - Cone width determines penumbra



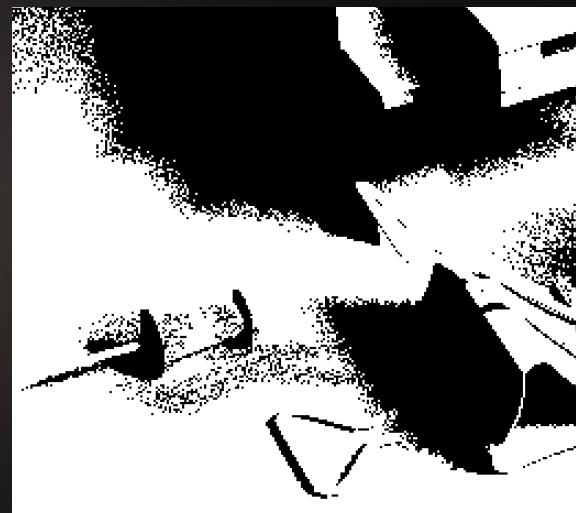
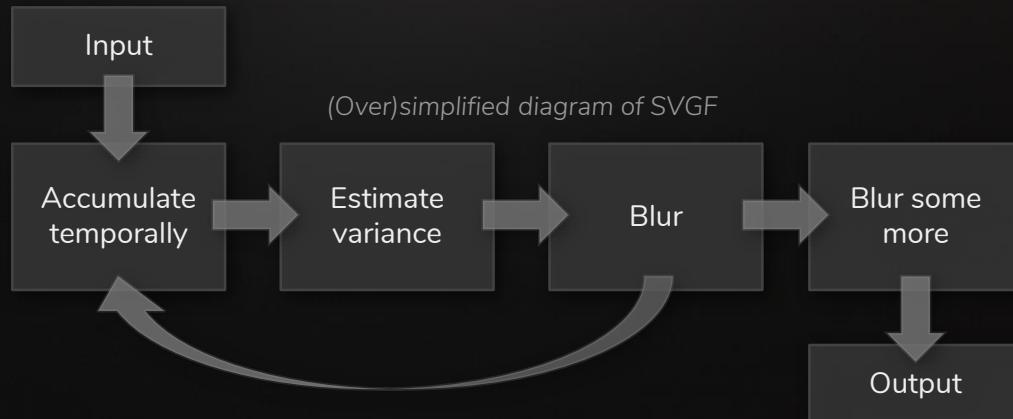
# Shadows

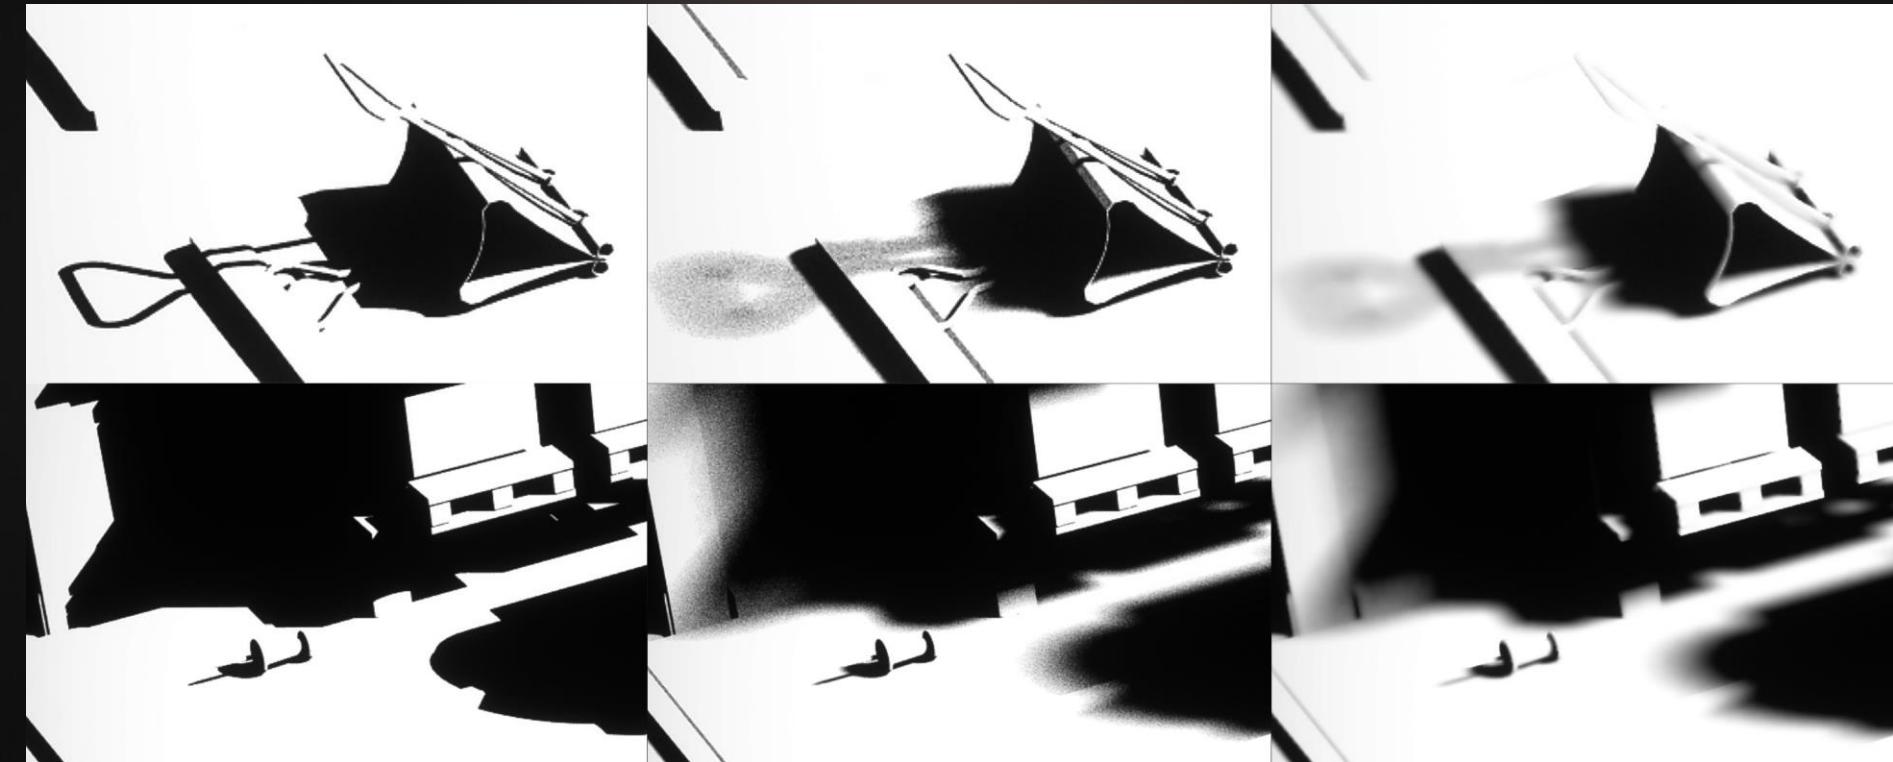
- Hard shadows trivial
  - Launch a ray towards light position
  - Check for any hit
- Soft shadows easy
  - Sample direction from uniform cone [PBRT]
  - Cone width determines penumbra
- ... Easy but noisy
  - 1 visibility sample per pixel not enough
  - Filter instead of shooting more



# Shadow Filtering

- Based on Nvidia's SVGF [Schied 2017]
  - Temporal accumulation
  - Multi-pass weighted blur
    - Kernel size driven by variance estimate
  - No shadow-specific heuristics





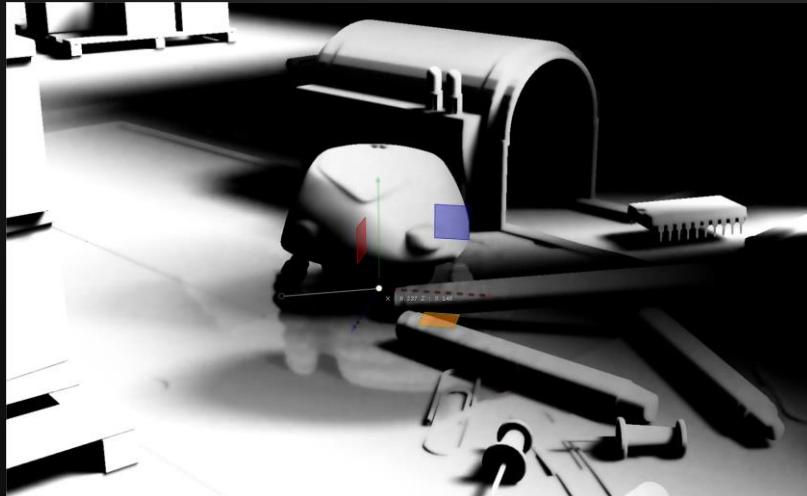
Hard Raytraced Shadows

Soft Raytraced Shadows (Unfiltered)

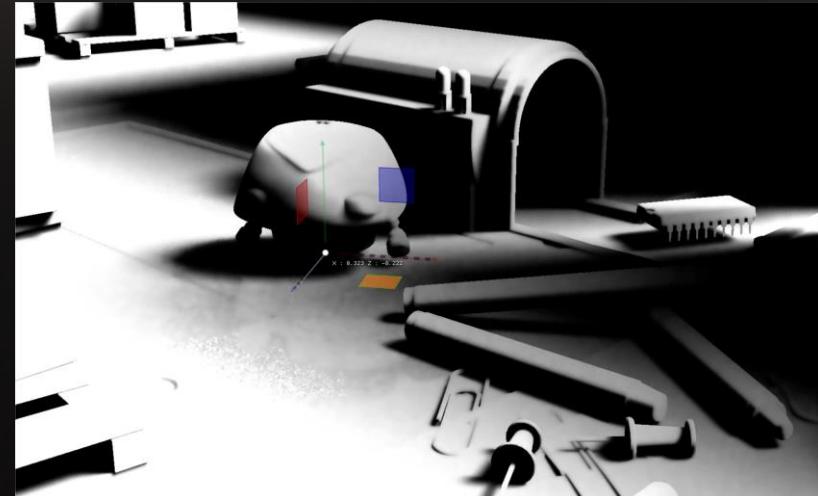
Soft Raytraced Shadows (Filtered)

# Shadow Filtering

- Modified to reduce ghosting
  - TAA-like color clip
    - Variance-based bounds (5x5 kernel) [Salvi 2016]



Baseline SVGF



+TAA-like clip

# Summary

- Replace fine-tuned hacks with unified approaches
  - Reconstruction and filtering instead
- New tool in the box
  - Solve sparse and incoherent problems
  - Not a silver bullet
- Use rays wisely
  - Pica Pica shoots ~2.25 rays per pixel
- Real-time visuals with (almost) path traced quality achievable today

# Thanks

- **SEED**

- Johan Andersson
- Colin Barré-Brisebois
- Jasper Bekkers
- Joakim Bergdahl
- Ken Brown
- Dean Calver
- Dirk de la Hunt
- Jenna Frisk
- Paul Greveson
- Henrik Halen
- Effeli Holst
- Andrew Lauritzen

- Magnus Nordin
- Niklas Nummelin
- Anastasia Opara
- Kristoffer Sjöö
- Ida Winterhaven
- Graham Wihlidal

- **Microsoft**

- Chas Boyd
- Ivan Nevraev
- Amar Patel
- Matt Sandy

- **NVIDIA**

- Tomas Akenine-Möller
- Nir Benty
- Jiho Choi
- Peter Harrison
- Alex Hyder
- Jon Jansen
- Aaron Lefohn
- Ignacio Llamas
- Henry Moreton
- Martin Stich

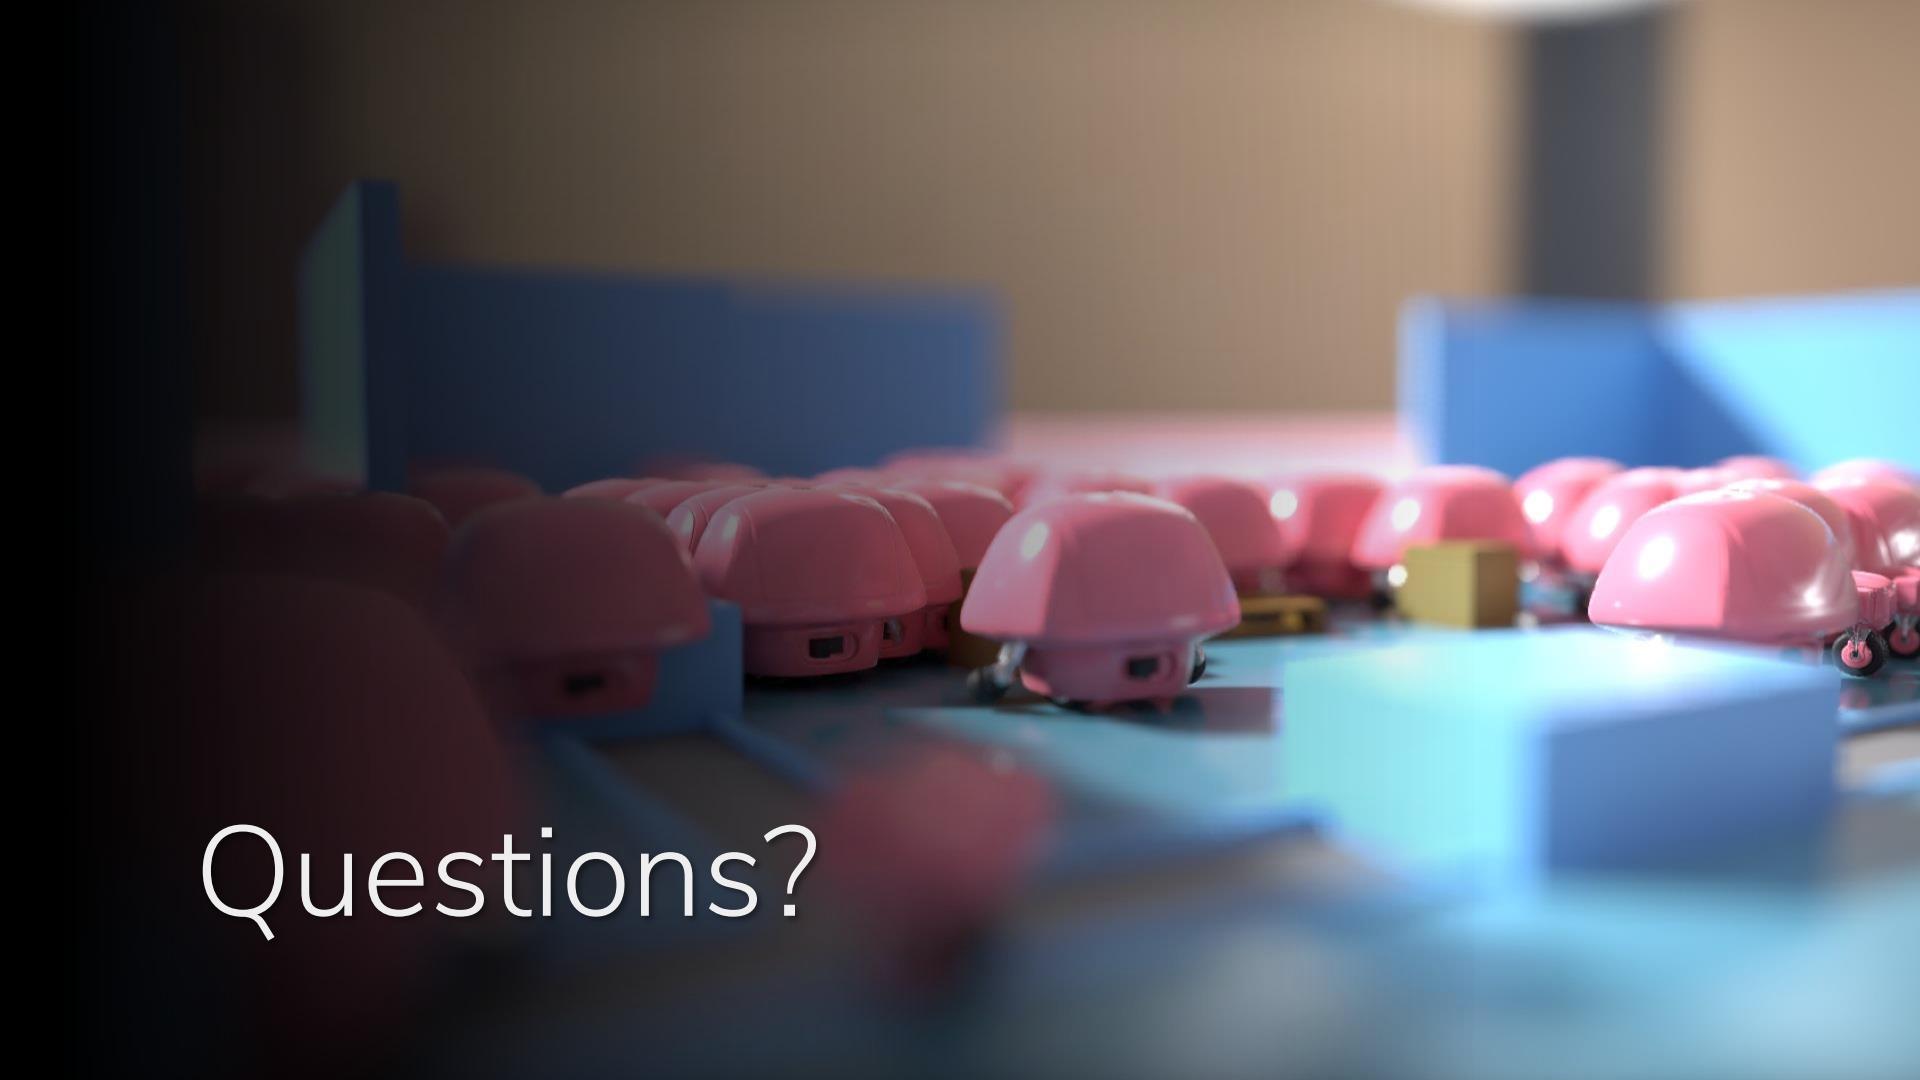


SEED // SEARCH FOR EXTRAORDINARY EXPERIENCES DIVISION

STOCKHOLM – LOS ANGELES – MONTRÉAL – REMOTE

[WWW.EA.COM/SEED](http://WWW.EA.COM/SEED)

WE'RE HIRING!

A close-up photograph of a row of small, pink, mushroom-shaped robots. Each robot has a white circular sensor on its cap and two black wheels at the base. They are lined up on a light blue surface. In the background, there are blurred blue and yellow structures, possibly parts of a playground or a room. The lighting is soft, creating a shallow depth of field.

Questions?

# References

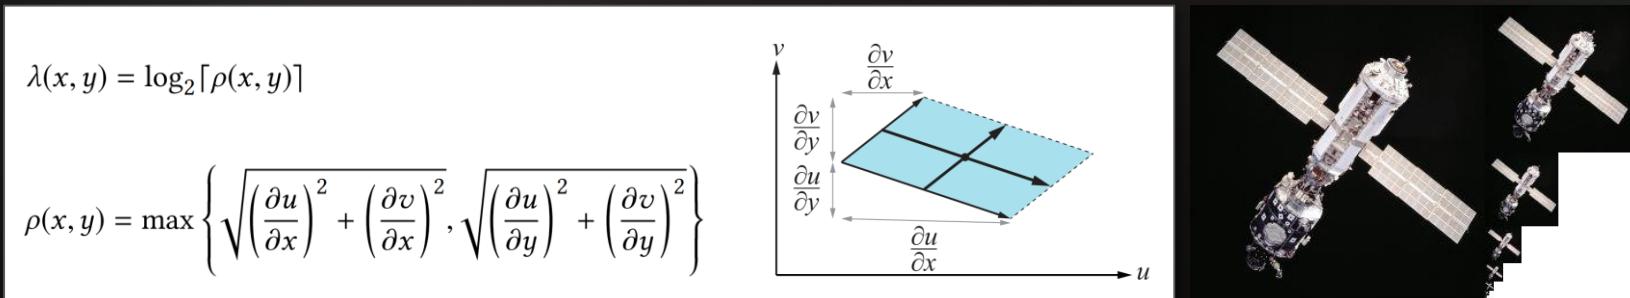
- **[Harmer 2018]** Jack Harmer, Linus Gisslén, Henrik Holst, Joakim Bergdahl, Tom Olsson, Kristoffer Sjöö and Magnus Nordin "Imitation Learning with Concurrent Actions in 3D Games". [available online](#).
- **[Barré-Brisebois 2011]** Barré-Brisebois, Colin and Bouchard, Marc. "Approximating Translucency for a Fast, Cheap and Convincing Subsurface Scattering Look", [available online](#).
- **[Barré-Brisebois 2017]** Barré-Brisebois, Colin. "A Certain Slant of Light: Past, Present and Future Challenges of Global Illumination in Games", [available online](#).
- **[Igehy 1999]** Igehy, Homan. "Tracing Ray Differentials", [available online](#).
- **[PBRT]** Pharr, Matt. Jakob, Wenzel and Humphreys, Greg. "Physically Based Rendering", Book, <http://www.pbrt.org/>.
- **[Schied 2017]** Schied, Christoph et. Al. "Spatiotemporal Variance-Guided Filtering: Real-Time Reconstruction for Path-Traced Global Illumination", NVIDIA Research, [available online](#).
- **[Stachowiak 2015]** Stachowiak, Tomasz. "Stochastic Screen-Space Reflections", [available online](#).
- **[Weidlich 2007]** Weidlich, Andrea and Wilkie, Alexander. "Arbitrarily Layered Micro-Facet Surfaces", [available online](#).
- **[Williams 1983]** Williams, Lance. "Pyramidal Parametrics", [available online](#).
- **[Robison 2009]** Austin Robison, Peter Shirley. "Image Space Gathering", [available online](#).
- **[Heitz 2018]** Eric Heitz, Stephen Hill, Morgan McGuire. "Combining Analytic Direct Illumination and Stochastic Shadows", [available online](#).
- **[Karis 2014]** Brian Karis. "High-Quality Temporal Supersampling", [available online](#).
- **[Lehtinen 2008]** Jaakko Lehtinen, Matthias Zwicker, Emmanuel Turquin, Janne Kontkanen, Frédéric Durand, François Sillion, Timo Aila. "A Meshless Hierarchical Representation for Light Transport", [available online](#).
- **[Jimenez 2016]** Jorge Jimenez, Xian-Chun Wu, Angelo Pesce, Adrian Jarabo. "Practical Realtime Strategies for Accurate Indirect Occlusion", [available online](#).
- **[Opara 2018]** Anastasia Opara. "Creativity of Rules and Patterns", [available online](#).
- **[Salvi 2016]** Marco Salvi. "An excursion in temporal super sampling", [available online](#).
- **[Hillaire 2018]** Sébastien Hillaire. "Real-time Raytracing for Interactive Global Illumination Workflows in Frostbite", [available online](#).

# Bonus

# Texture Level-of-Detail

What about texture level of detail?

- Mipmapping [Williams 1983] is the standard method to avoid texture aliasing:



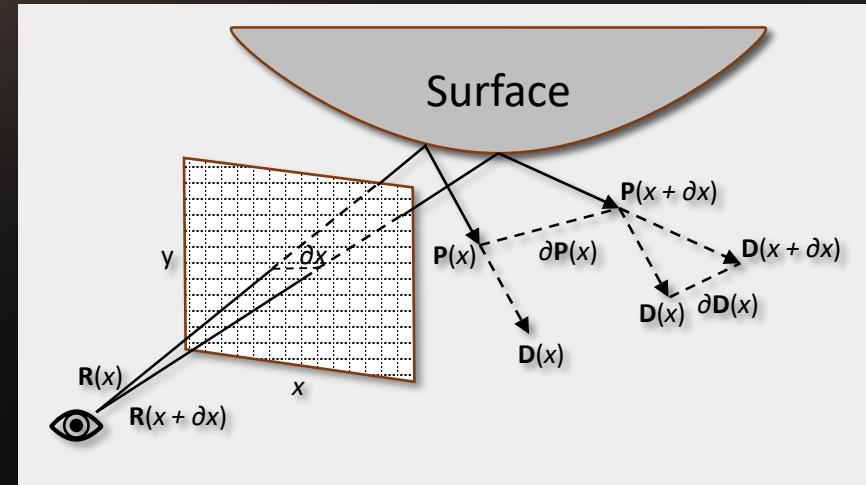
Left: level-of-detail ( $\lambda$ ), partial derivatives and the parallelogram-approximated texture-space footprint of a pixel. Right: mipmap chain

- Screen-space pixel maps to approximately one texel in the mipmap hierarchy
- Supported by all GPUs for rasterization via shading quad and derivatives

# Texture Level-of-Detail

No shading quads for ray tracing!

- Traditionally: *Ray Differentials*
  - Estimates the footprint of a pixel by computing world-space derivatives of the ray with respect to the image plane
  - Have to differentiate (virtual offset) rays
  - Heavier payload (12 floats) for subsequent rays (can) affect performance. Optimize!



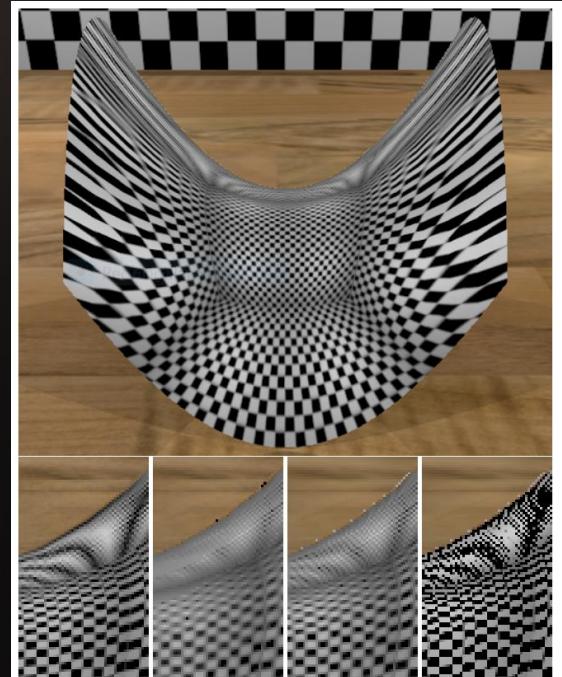
Ray Differentials [Igehy99]

- Alternative: always sample mip 0 with bilinear filtering (with extra samples)
  - Leads to aliasing and additional performance cost

# Texture Level-of-Detail

Together with  **NVIDIA**. Research, we developed a texture LOD technique for raytracing:

- Heuristic based on **triangle properties**, a **curvature estimate**, **distance**, and **incident angle**
  - Similar quality to ray differentials with single trilinear lookup
  - Single value stored in the payload for subsequent rays
- Upcoming publication by:
  - Tomas Akenine-Möller (NV), Magnus Andersson (NV), Colin Barré-Brisebois (EA), Jim Nilsson (NV), Robert Toth (NV)



Ground Truth, Ray Differentials, Ours, Mip0  
Digital Dragons

# mGPU

- Explicit Heterogenous Multi-GPU
- Parallel Fork-Join Style
- Resources copied through system memory using copy queue
- Minimize PCI-E transfers
- Approach
  - Run ray generation on primary GPU
  - Copy results in sub-regions to other GPUs
  - Run tracing phases on separate GPUs
  - Copy tracing results back to primary GPU
  - Run filtering on primary GPU

